



Concurrent Objects



Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

Concurrent Computation




Objectivism

- What is a concurrent object?
 - How do we **describe** one?
 - How do we **implement** one?
 - How do we **tell if we're right**?


Objectivism

- What is a concurrent object?
 - How do we **describe** one?
 - How do we **tell if we're right**?

FIFO Queue: Enqueue Method



FIFO Queue: Dequeue Method




A Lock-Based Queue

```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```

A Lock-Based Queue


```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```



Queue fields
protected by single
shared lock

A Lock-Based Queue


```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```



Initially head = tail


Implementation: Deq

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```



Implementation: Deq

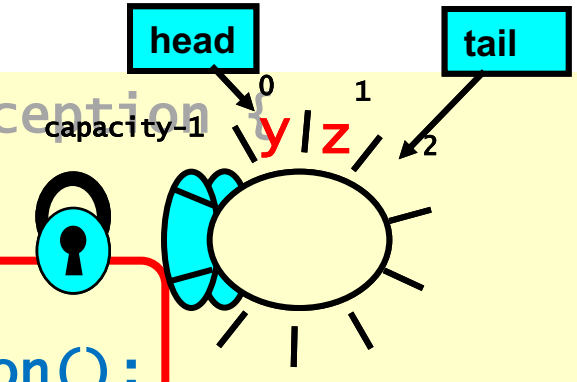
```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



Method calls
mutually exclusive

Implementation: Deq

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```




The diagram illustrates a queue implementation. It features a circular array with slots. A blue padlock icon represents a lock. A red box highlights the condition `if (tail == head)` in the code, with a red arrow pointing from it to the text "If queue empty throw exception".

Diagram details: A circular array with slots. A blue padlock icon represents a lock. A red box highlights the condition `if (tail == head)` in the code, with a red arrow pointing from it to the text "If queue empty throw exception".

Labels in the diagram: `head` (pointing to slot 0), `tail` (pointing to slot 2), `capacity-1` (pointing to slot 1), `y` (in slot 1), `z` (in slot 2).

Implementation: Deq


```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```



Queue not empty:
remove item and update
head

Implementation: Deq


```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```



Return result

Implementation: Deq

```
public T deq() throws EmptyException {
    lock.lock();
    try {
        if (tail == head)
            throw new EmptyException();
        T x = items[head % items.length];
        head++;
        return x;
    } finally {
        lock.unlock();
    }
}
```



Release lock no
matter what!

Implementation: Deq

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

**Should be correct because
modifications are mutually
exclusive...**

Now consider the following implementation

- The same thing without mutual exclusion
- For simplicity, only two threads
 - One thread **enq only**
 - The other **deq only**

Wait-free 2-Thread Queue

```
public class WaitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```

Wait-free 2-Thread Queue


```
public class waitFreeQueue {
```

```
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];
```

```
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }
```


```
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }
```

```
}}
```



Wait-free 2-Thread Queue

```
public class waitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Queue is updated  
        return item;  
    }  
}
```



How do we define “correct”
when modifications are not
mutually exclusive?

Defining concurrent queue implementations

- Need a way to specify a concurrent queue object
- Need a way to prove that an algorithm implements the object's specification
- Lets talk about object specifications ...

Correctness and Progress

- In a concurrent setting, we need to specify both the safety and the liveness properties of an object
- Need a way to define
 - when an implementation is correct
 - the conditions under which it guarantees progress

Let's begin with correctness

Sequential Objects

- Each object has a ***state***
 - Usually given by a set of ***fields***
 - Queue example: sequence of items
- Each object has a set of ***methods***
 - Only way to manipulate state
 - Queue example: **enq** and **deq** methods

Sequential Specifications

- If (precondition)
 - the object is in such-and-such a state
 - before you call the method,
- Then (postcondition)
 - the method will return a particular value
 - or throw a particular exception.
- and (postcondition, con't)
 - the object will be in some other state
 - when the method returns,

Pre and PostConditions for Dequeue

- **Precondition:**
 - Queue is non-empty
- **Postcondition:**
 - Returns first item in queue
- **Postcondition:**
 - Removes first item in queue

Pre and PostConditions for Dequeue

- **Precondition:**
 - Queue is empty
- **Postcondition:**
 - Throws Empty exception
- **Postcondition:**
 - Queue state unchanged


Why Sequential Specifications Totally Rock

- Interactions among methods captured by side-effects on object state
 - State meaningful between method calls
- Documentation size linear in number of methods
 - Each method described in isolation
- Can add new methods
 - Without changing descriptions of old methods


What About Concurrent Specifications ?

- Methods?
- Documentation?
- Adding new methods?


Methods Take Time




Methods Take Time



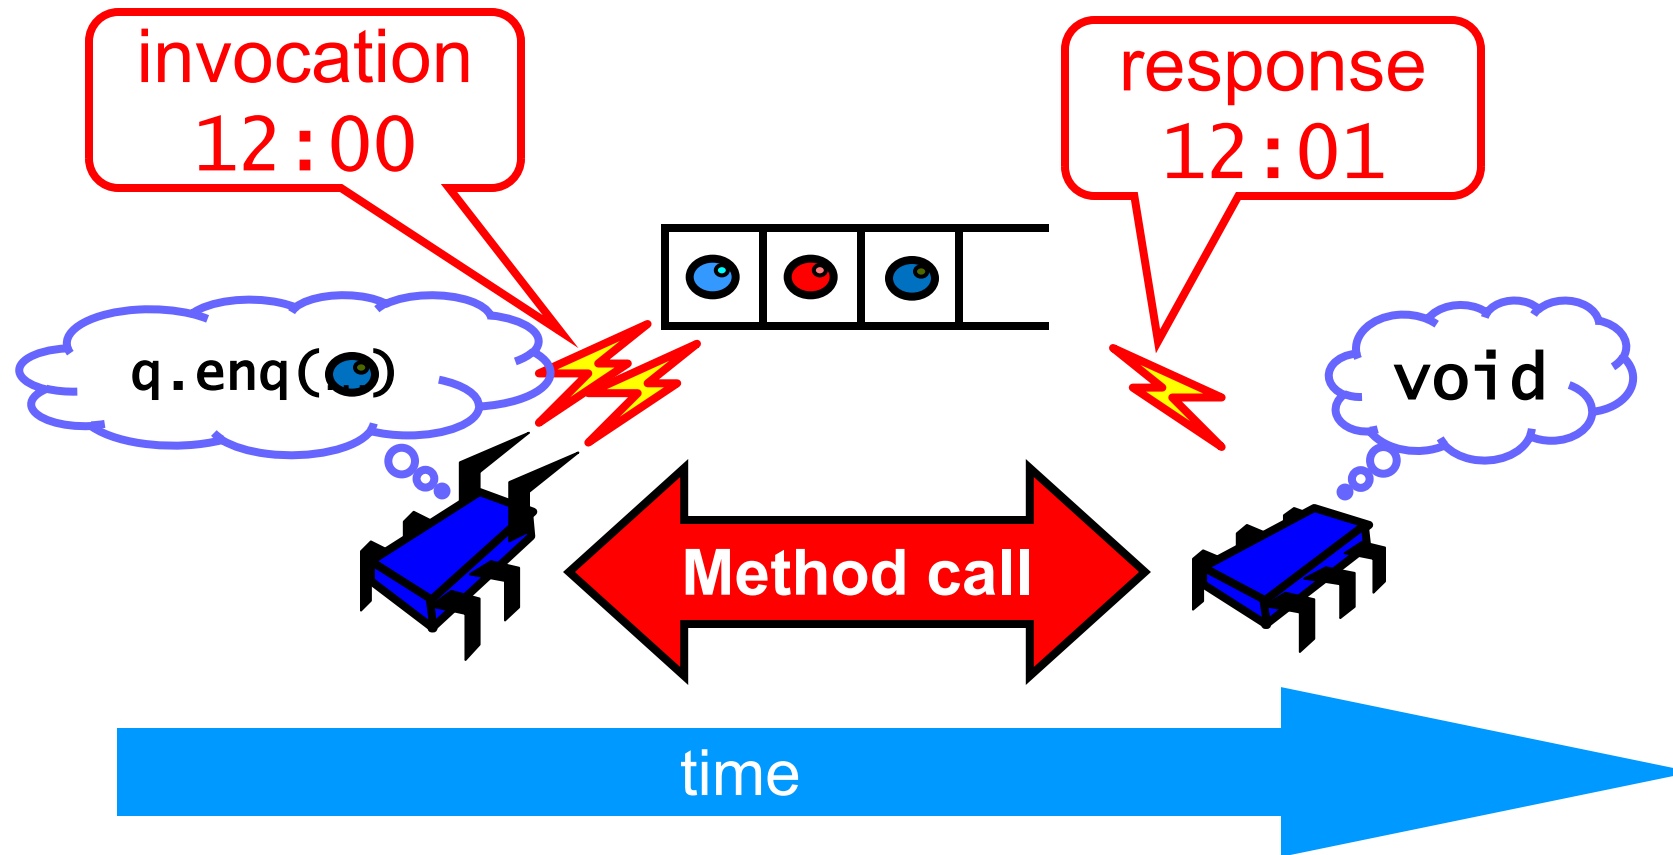
Methods Take Time



Methods Take Time




Methods Take Time




Sequential vs Concurrent

- Sequential
 - Methods take time? Who knew?
- Concurrent
 - Method call is not an event
 - Method call is an interval.

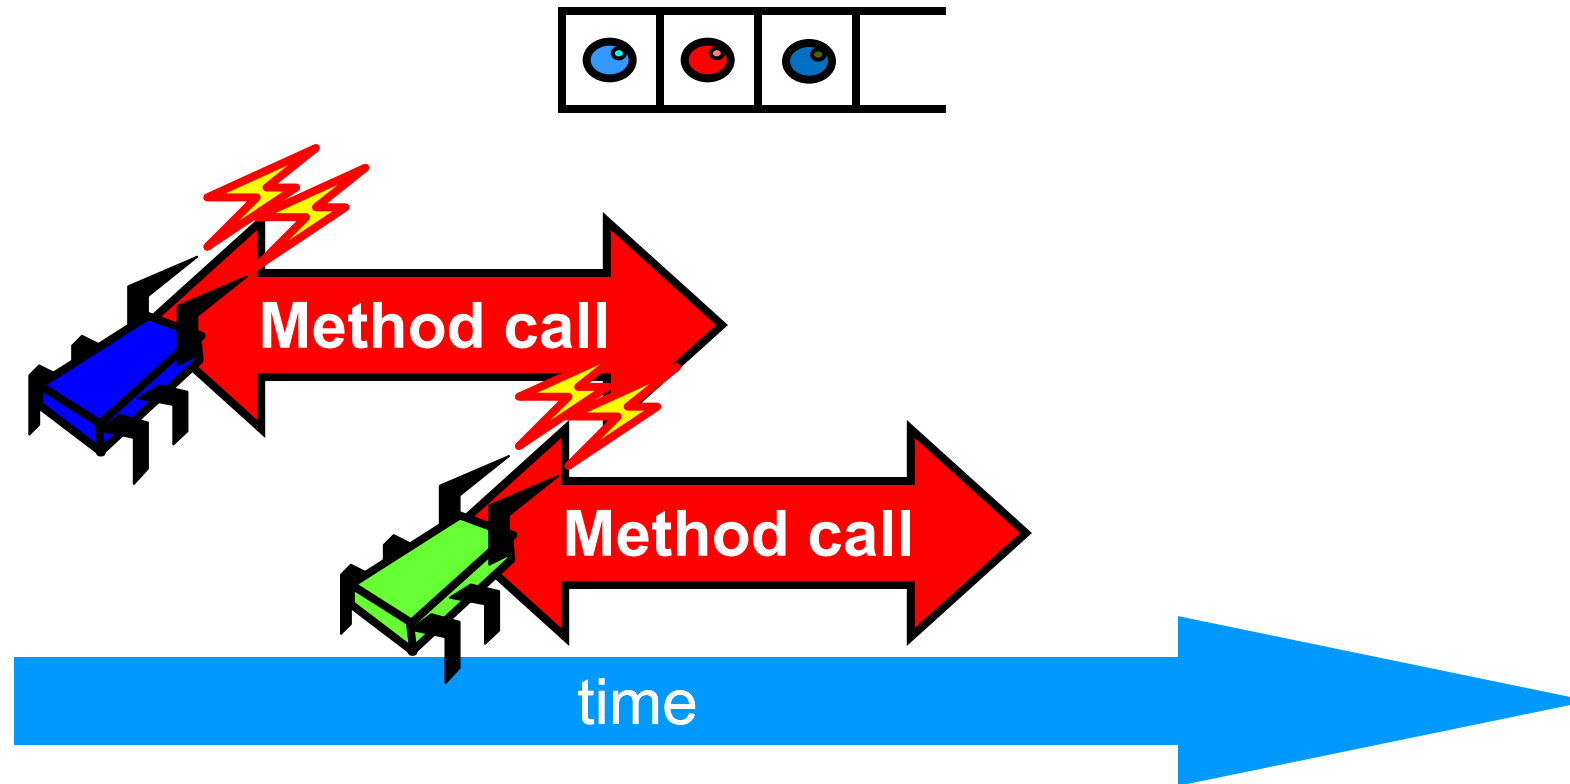
Concurrent Methods Take **Overlapping** Time




Concurrent Methods Take **Overlapping** Time



Concurrent Methods Take **Overlapping** Time



Concurrent Methods Take **Overlapping** Time



Sequential vs Concurrent

- Sequential:
 - Object needs meaningful state only **between** method calls
- Concurrent
 - Because method calls overlap, object might **never** be between method calls

Sequential vs Concurrent

- Sequential:
 - Each method described in isolation
- Concurrent
 - Must characterize ***all*** possible interactions with concurrent calls
 - What if two enqs overlap?
 - Two deqs? enq and deq? ...

Sequential vs Concurrent

- Sequential:
 - Can add new methods without affecting older methods
- Concurrent:
 - Everything can potentially interact with everything else

Sequential vs Concurrent

- Sequential:
 - Can add new methods without affecting older methods
- Concurrent:
 - Everything can potentially interact with everything else

Panic!

The Big Question

- What does it **mean** for a *concurrent* object to be correct?
 - What *is* a concurrent FIFO queue?
 - FIFO means strict temporal order
 - Concurrent means ambiguous temporal order

Intuitively...


```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Intuitively...

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

All modifications
of queue are done
mutually exclusive

Let's capture the idea of describing
the concurrent via the sequential




Linearizability

- Each method should
 - “take effect”
 - Instantaneously
 - Between invocation and response events
- Object is correct if this “sequential” behavior is correct
- Any such concurrent object is
 - **Linearizable**TM


Is it really about the object?

- Each method should
 - “take effect”
 - Instantaneously
 - Between invocation and response events
- Sounds like a property of an execution...
- A linearizable object: one all of whose possible executions are linearizable


Example




Example



Example

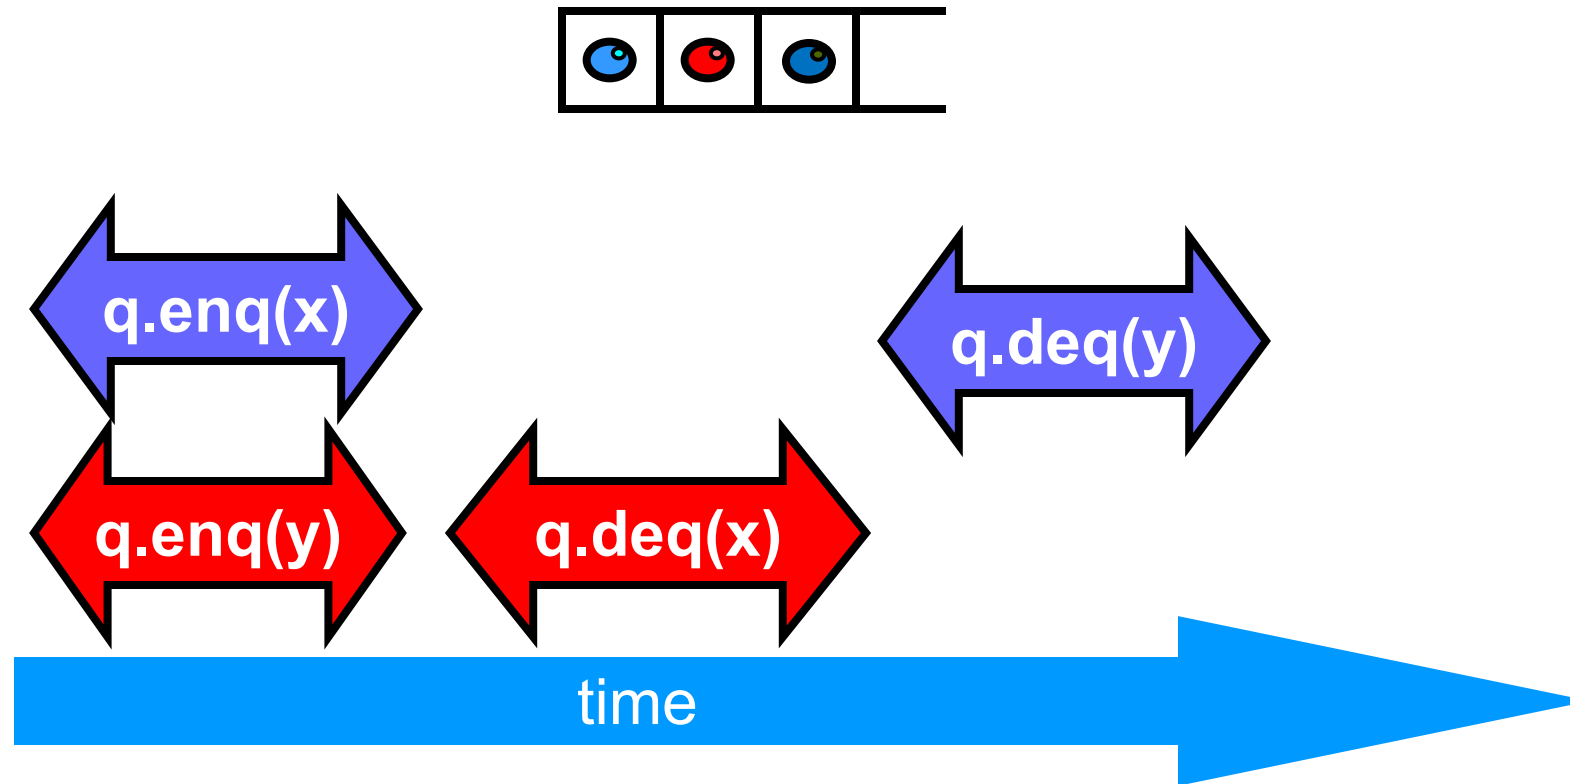


Example

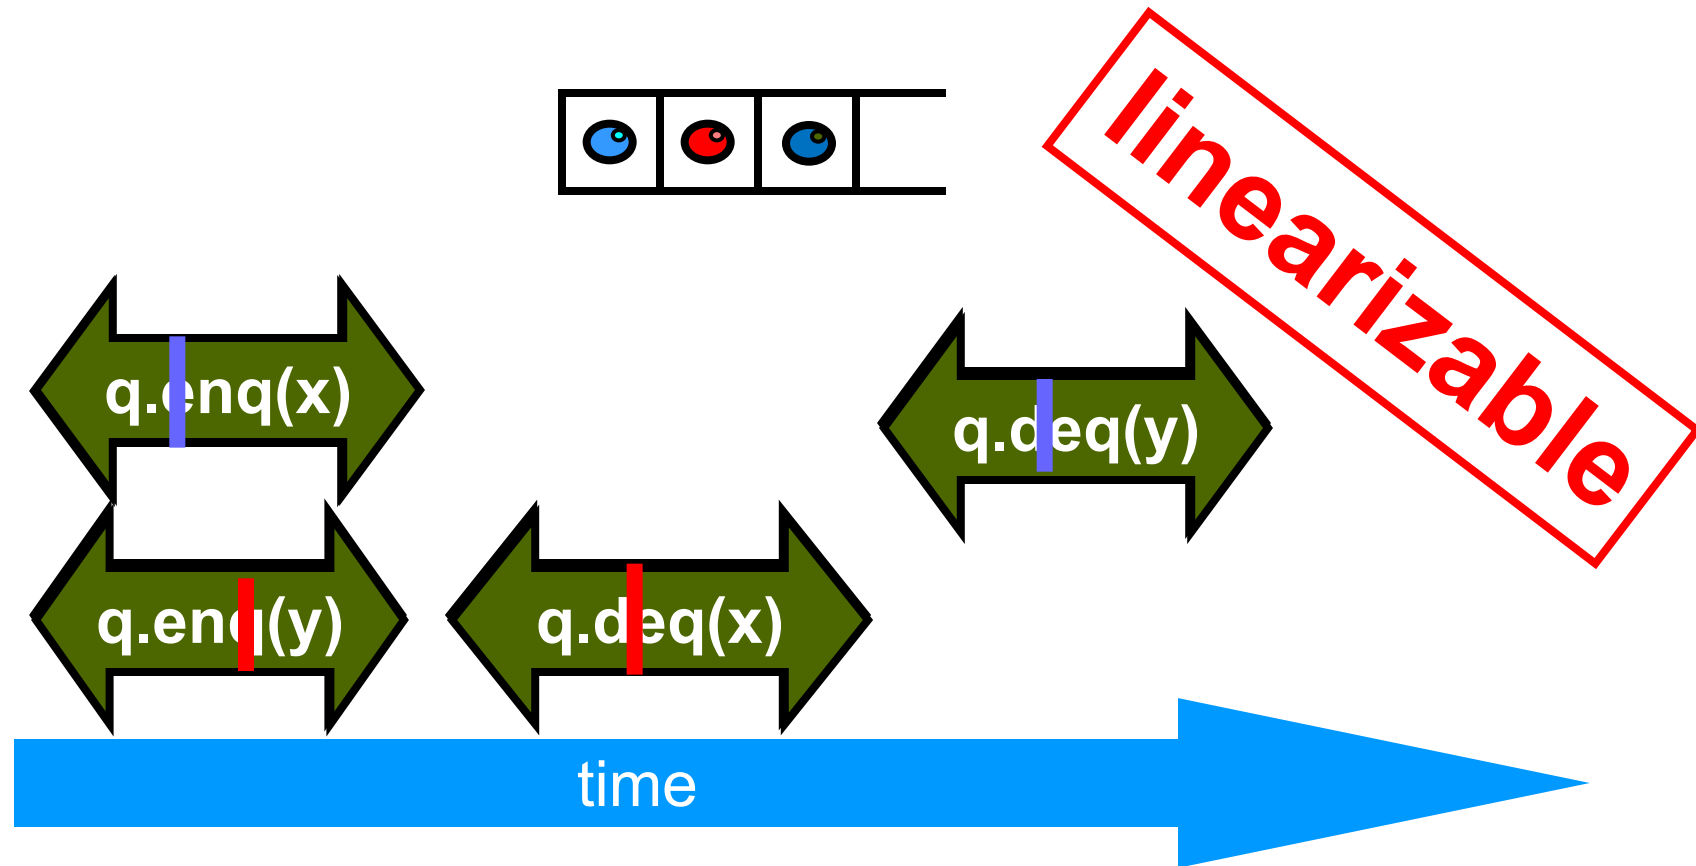




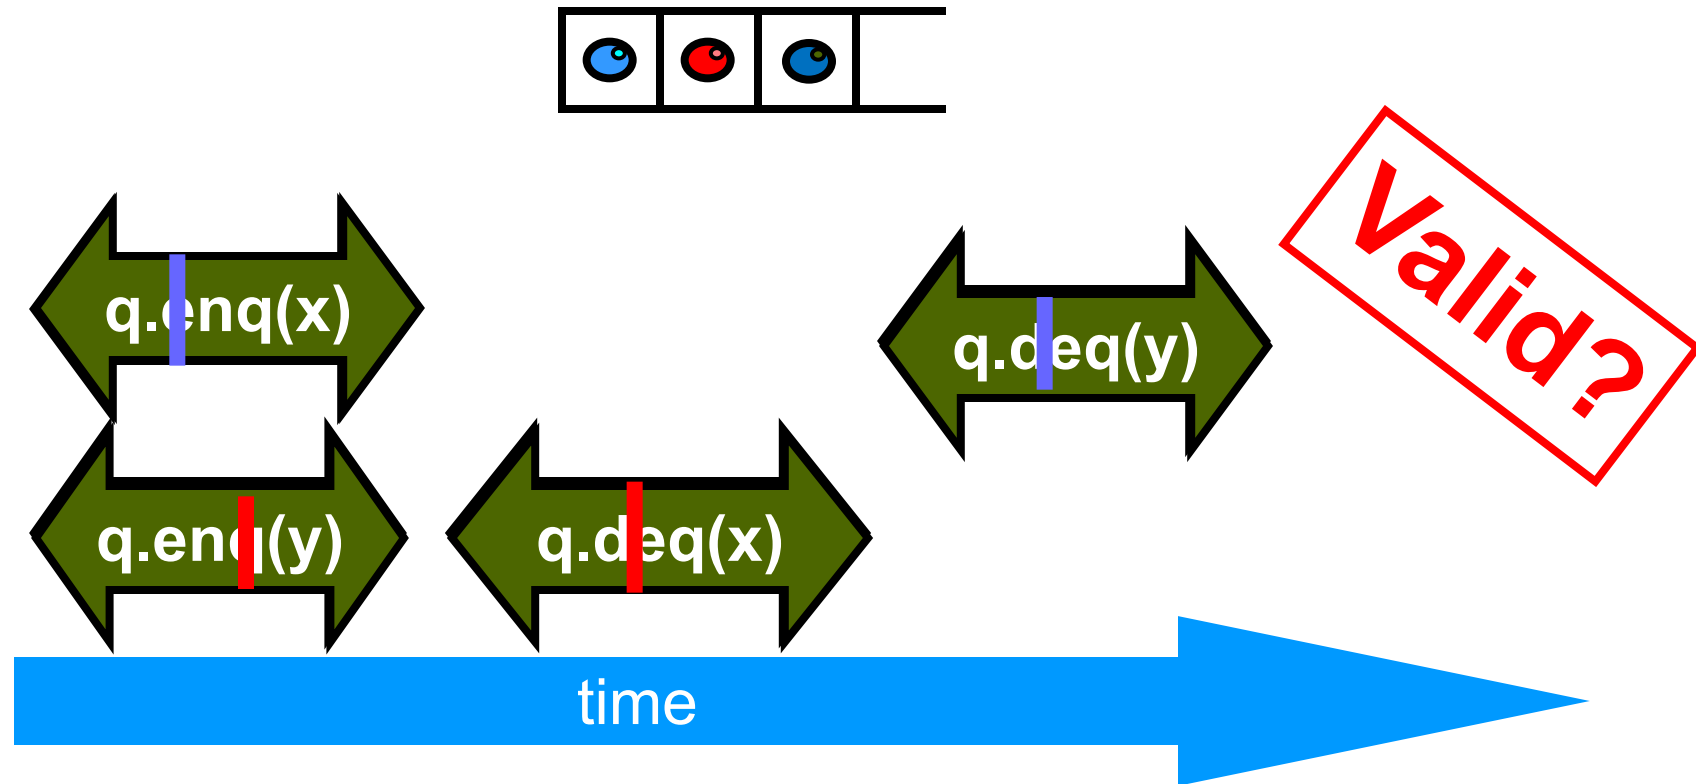
Example




Example




Example



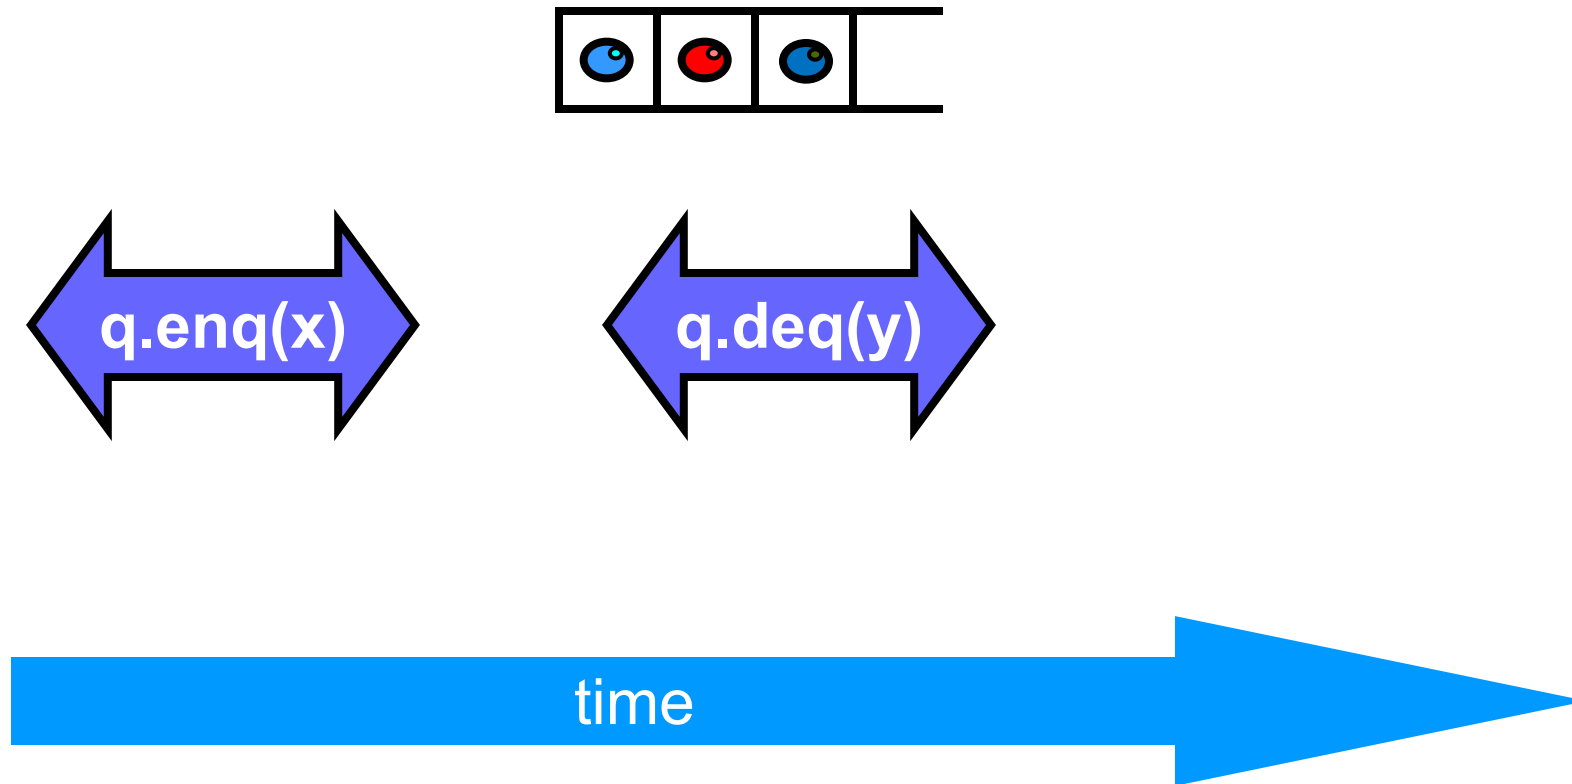
Example



Example

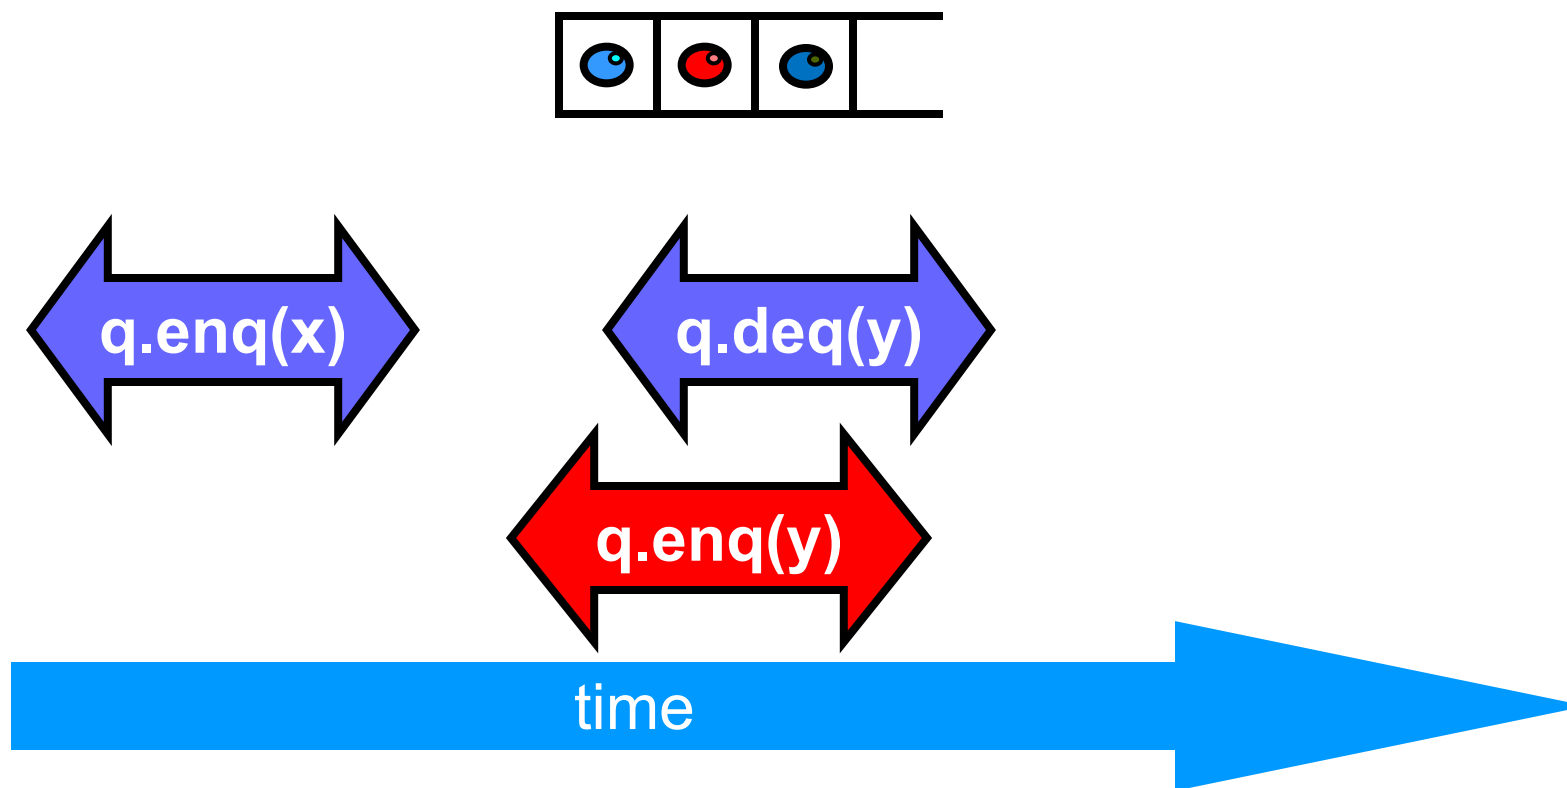


Example



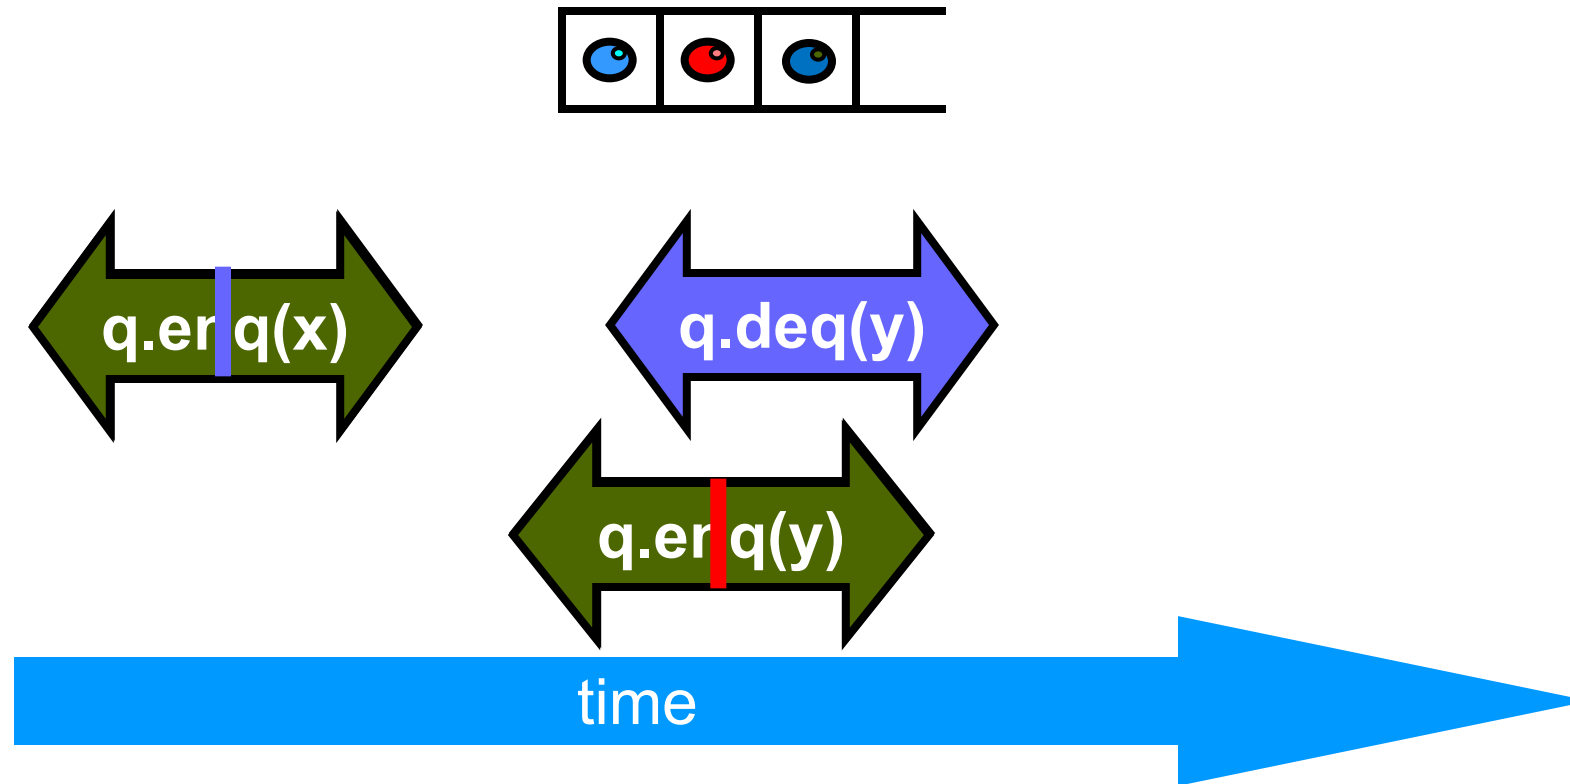


Example



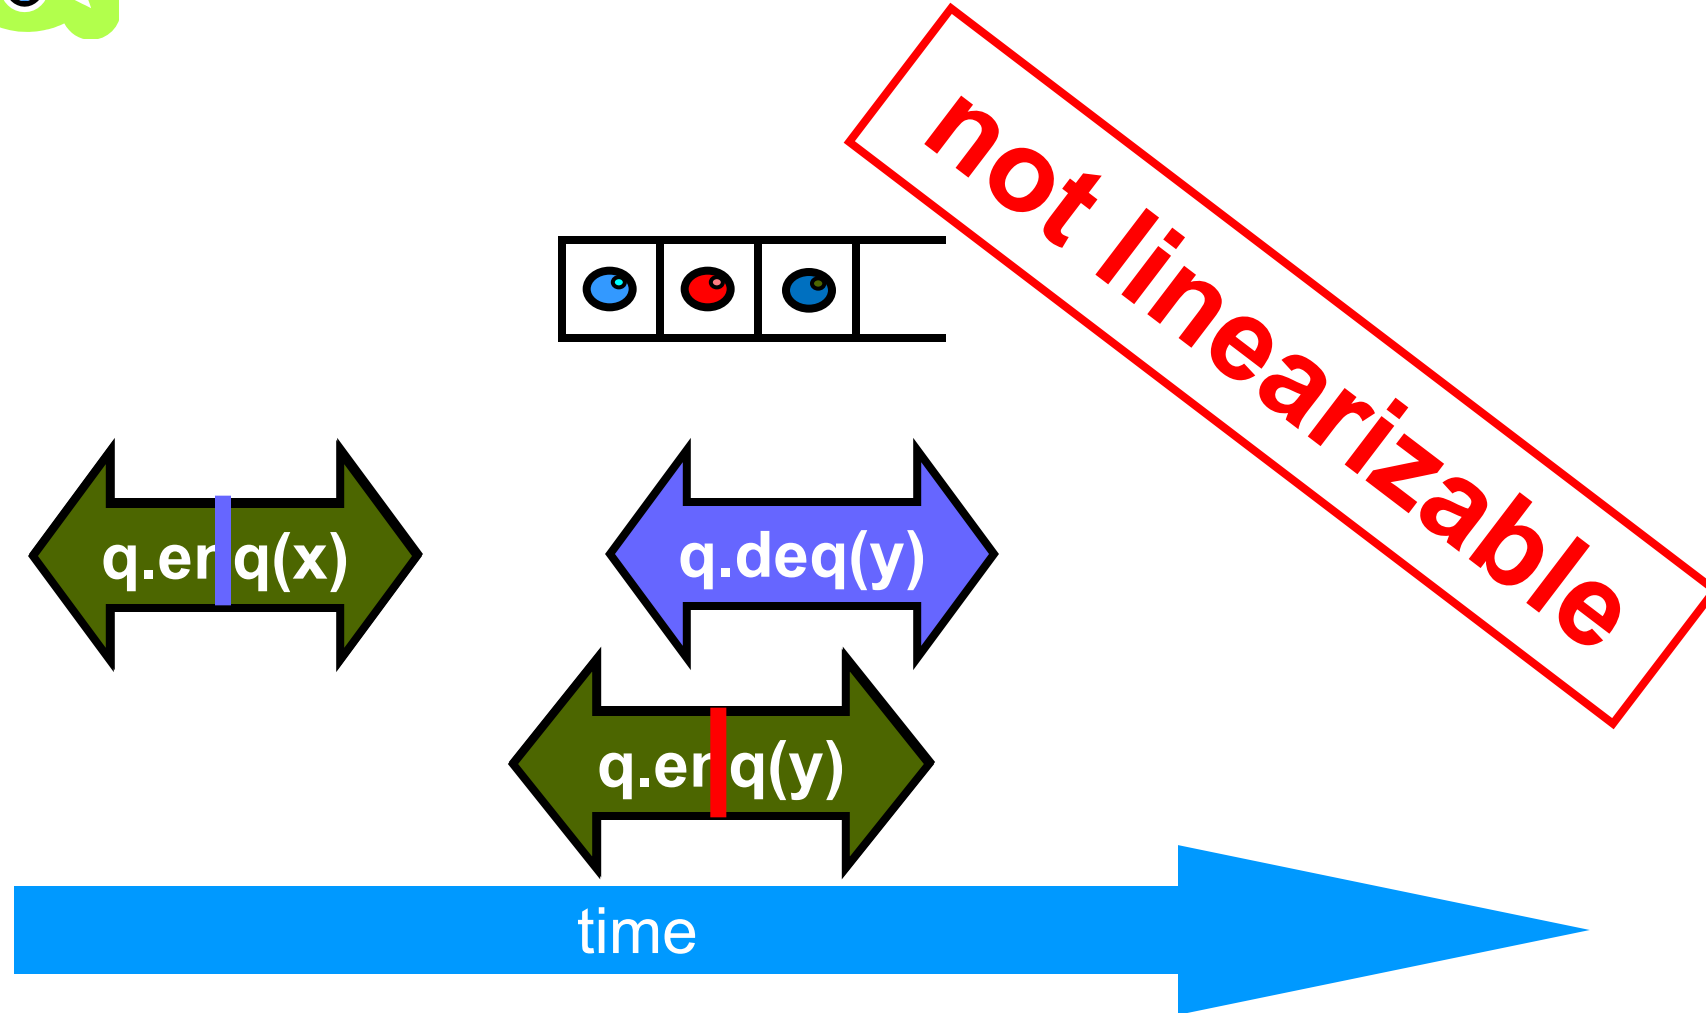


Example







Example



Example

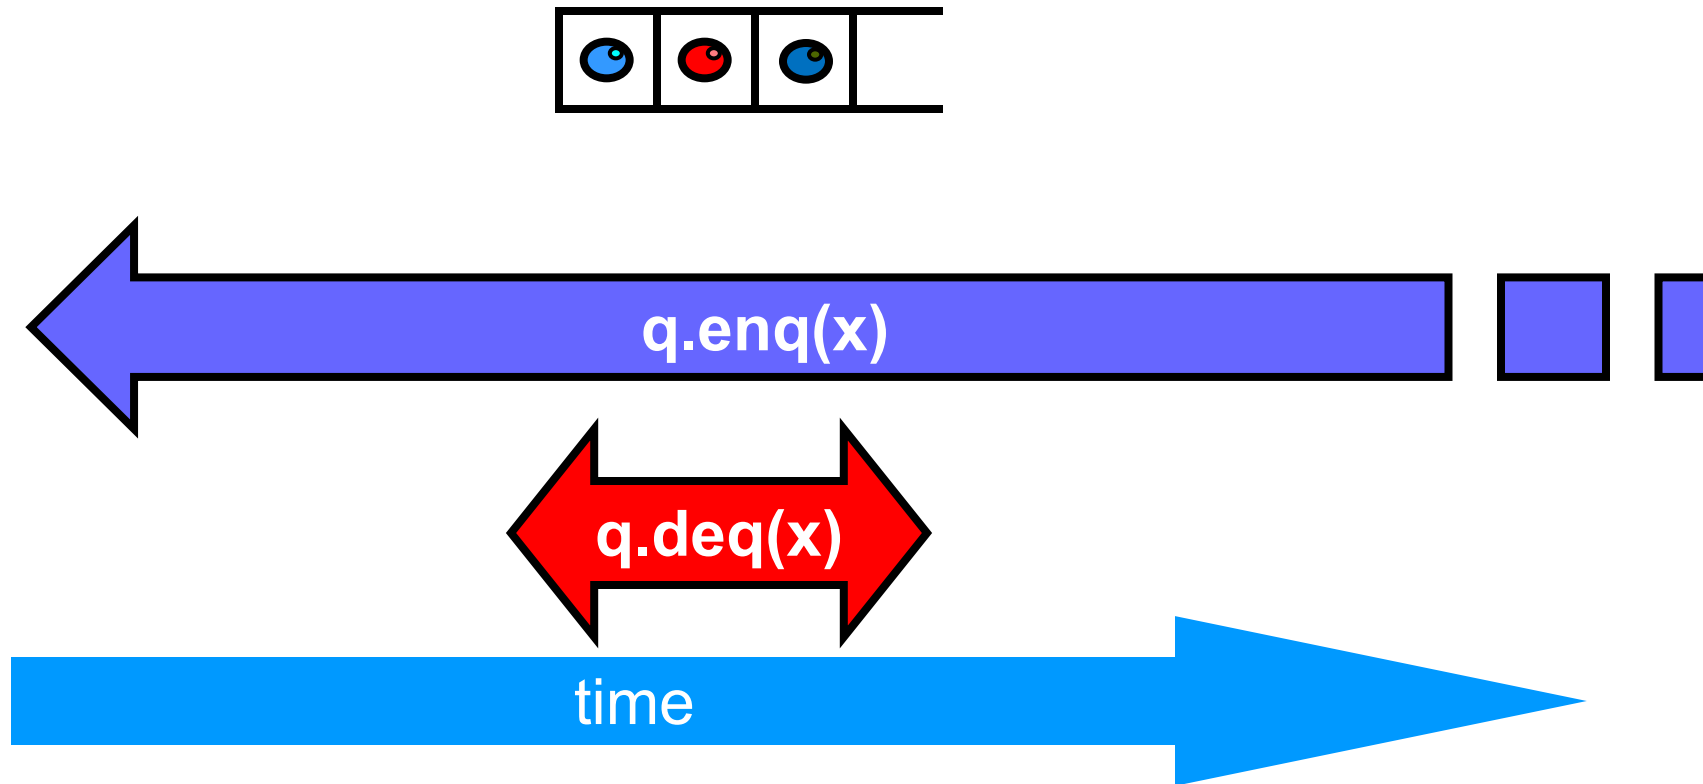


Example



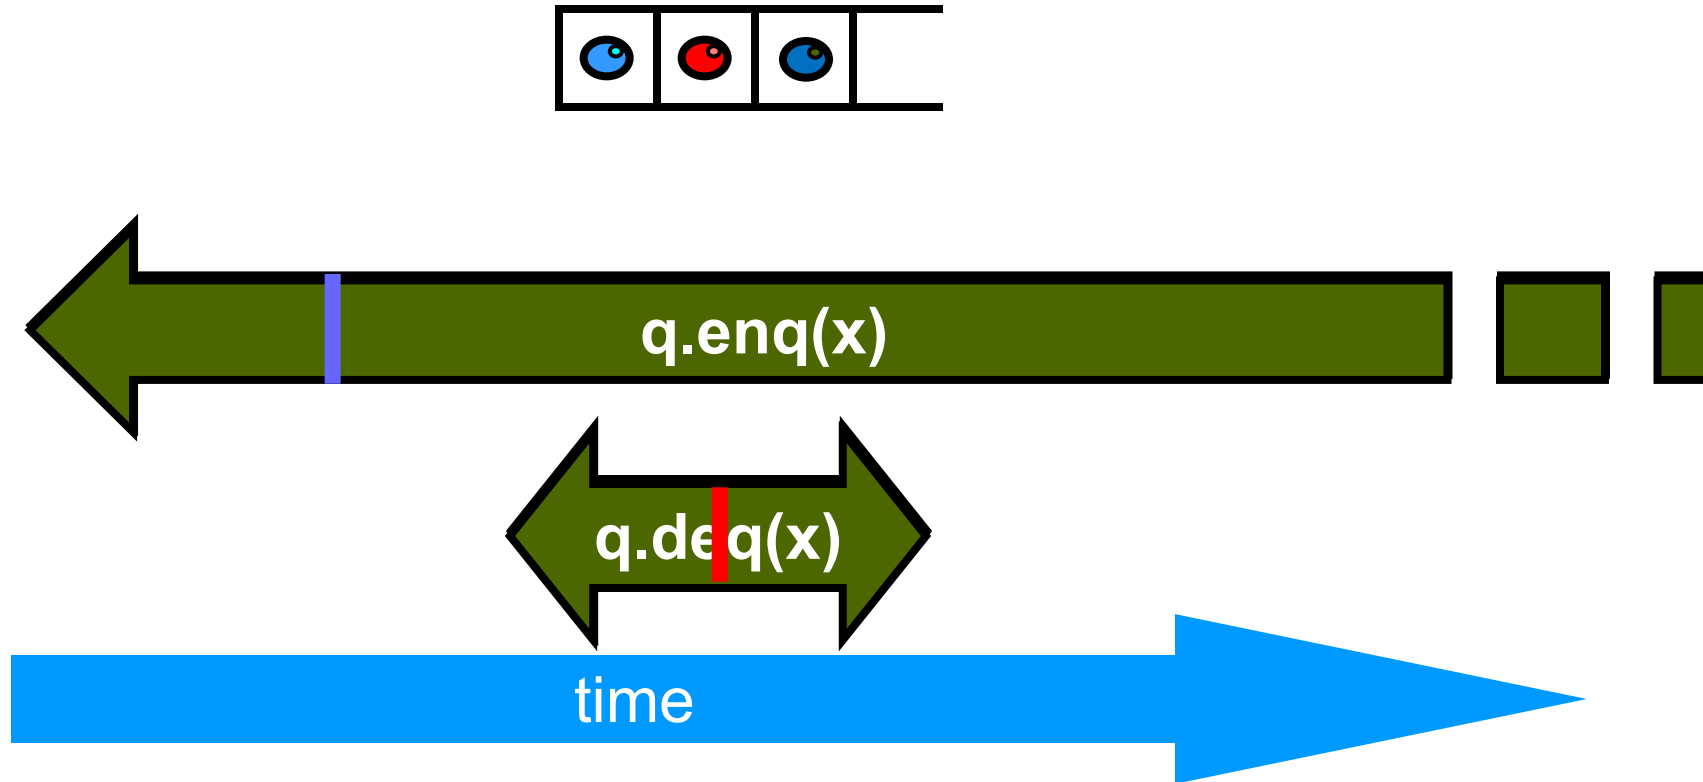


Example



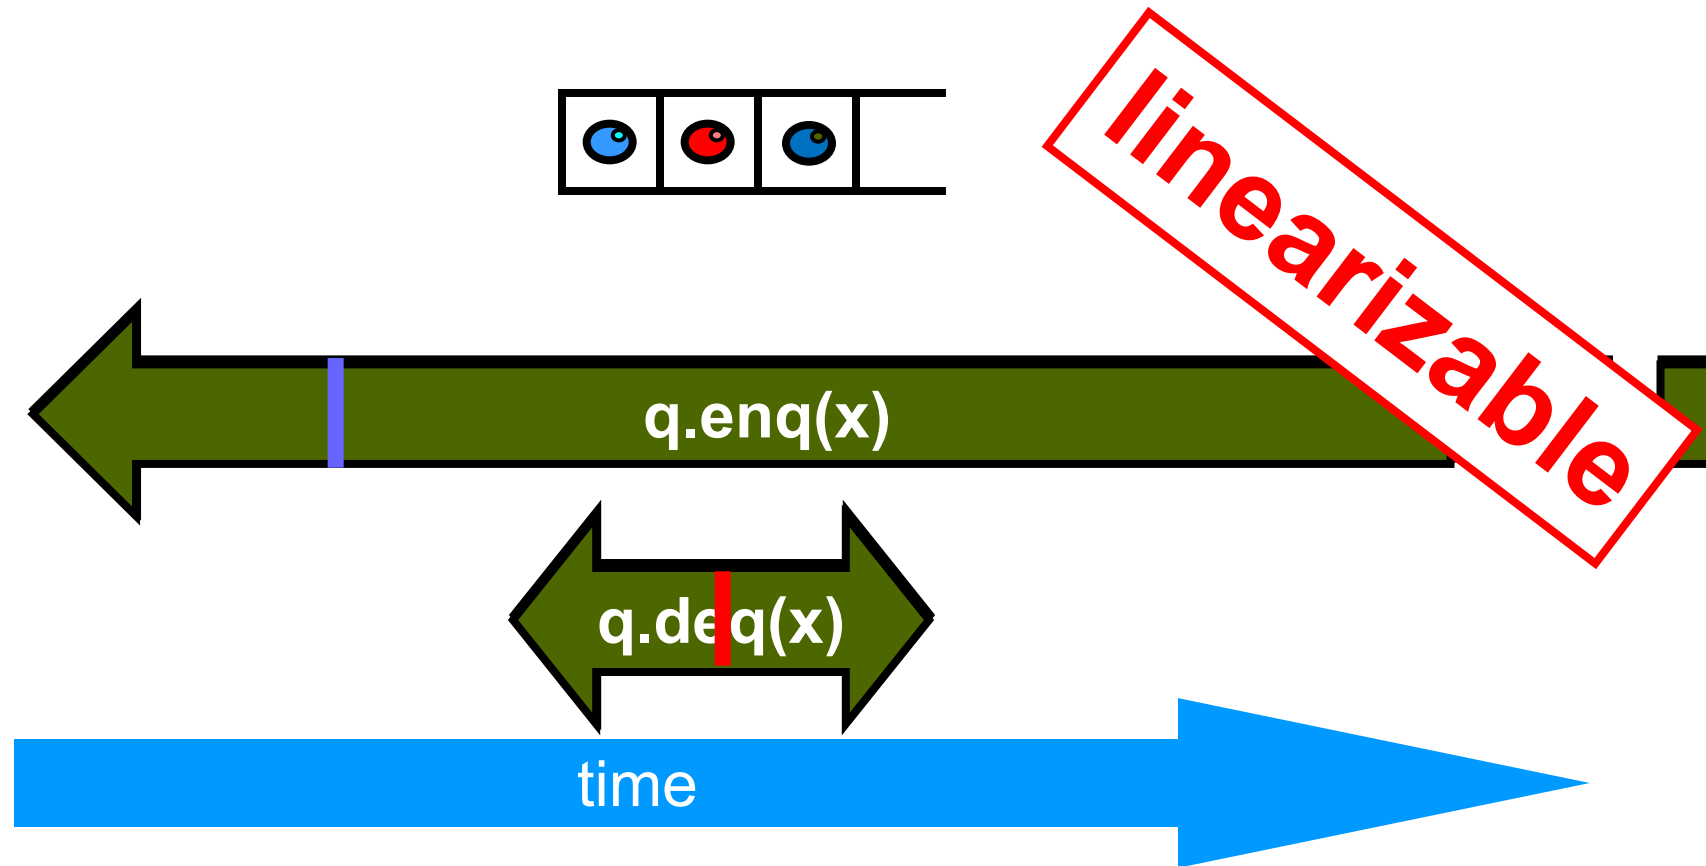


Example







Example



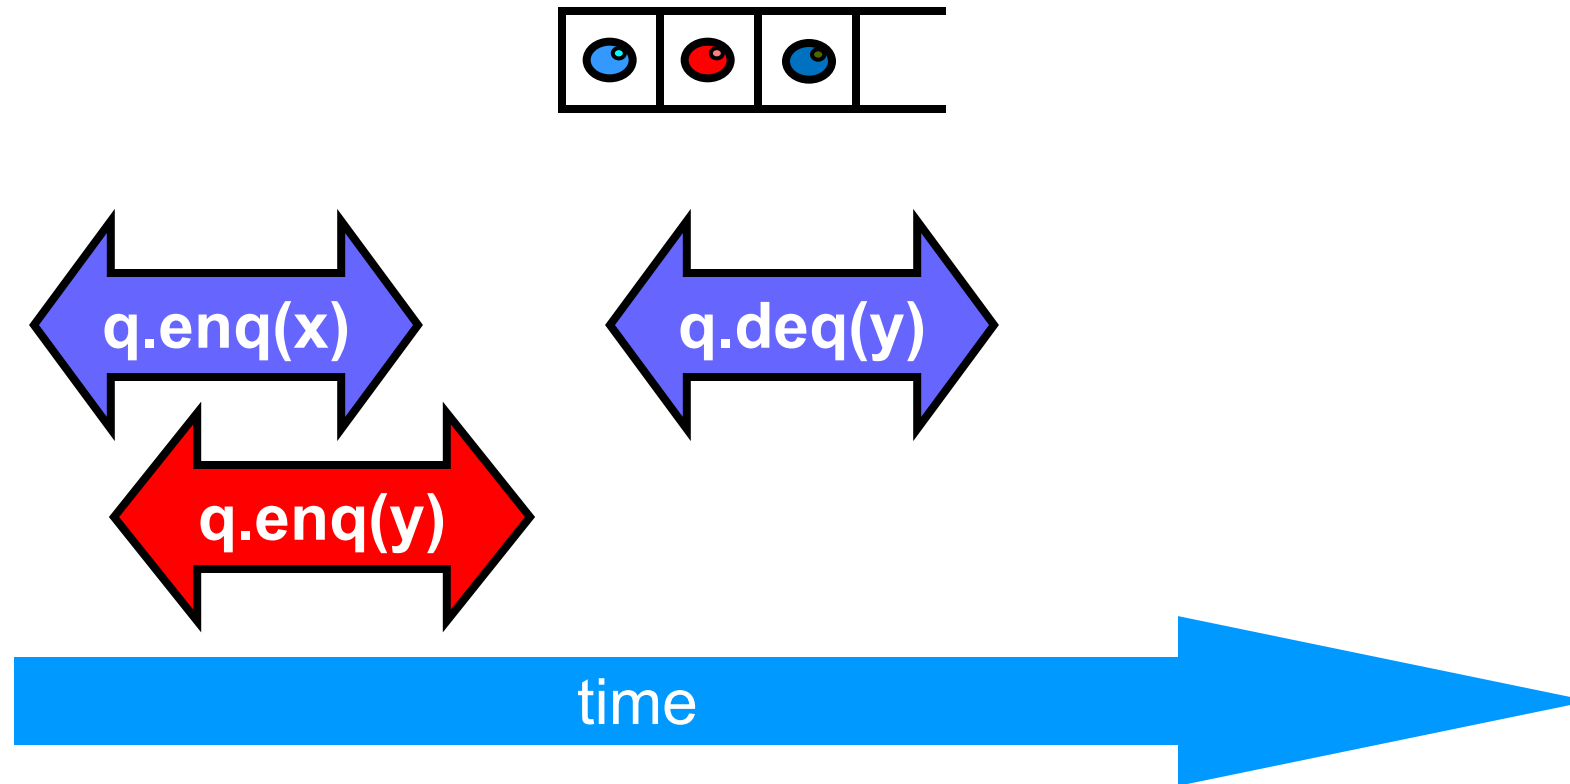
Example



Example

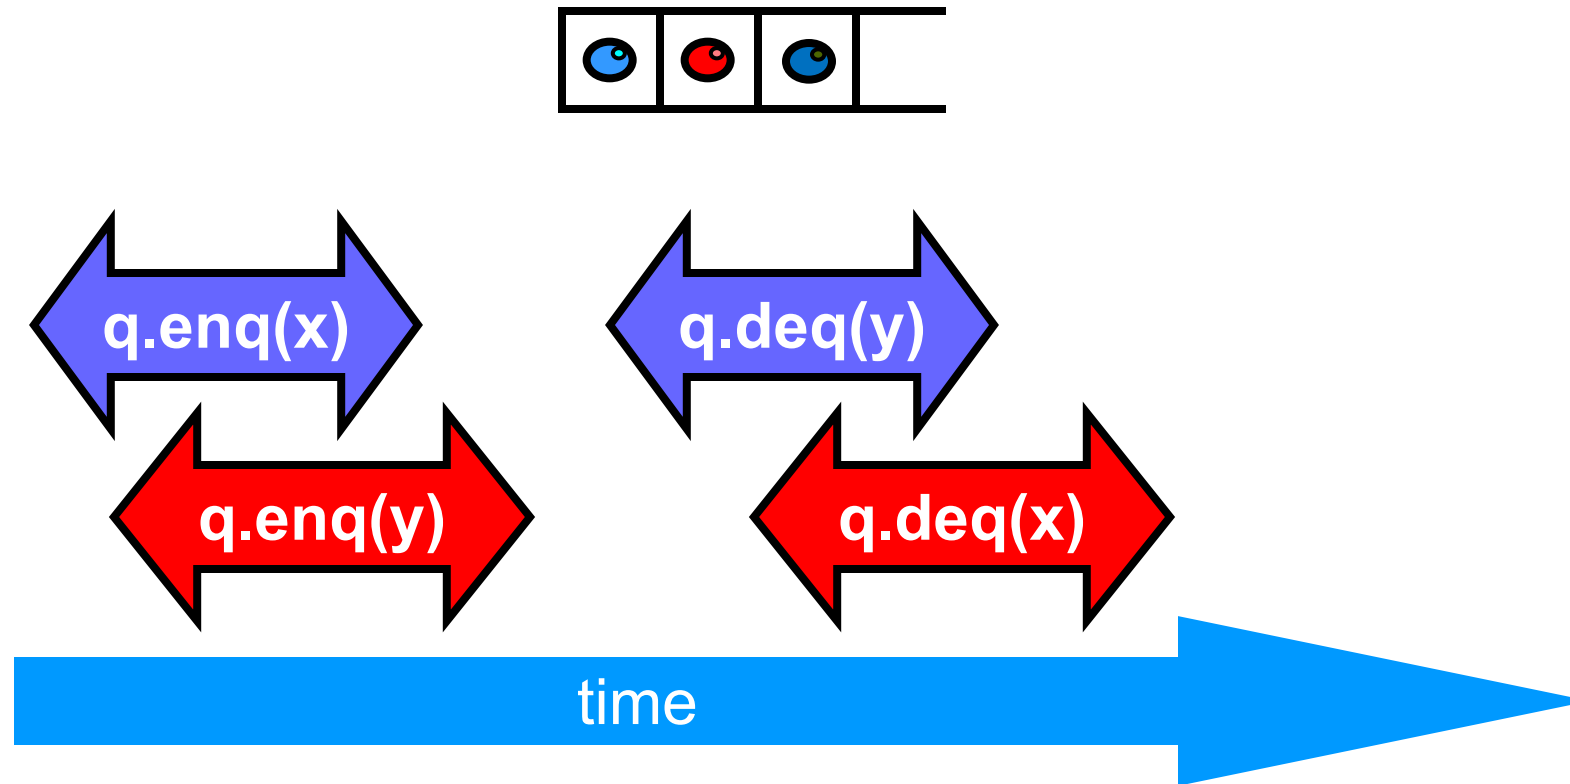


Example



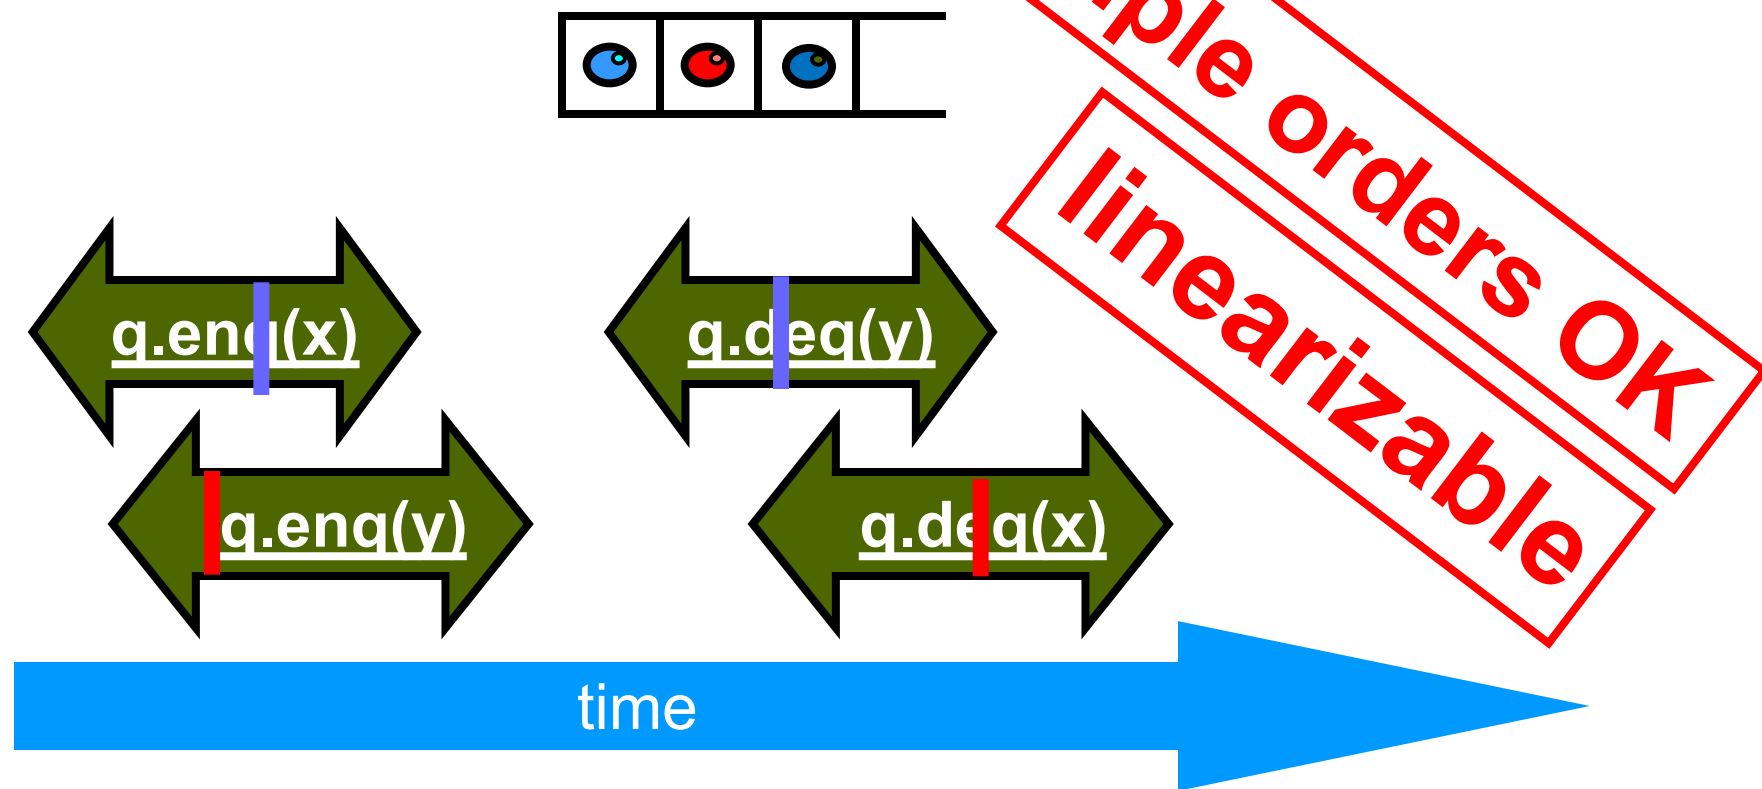


Example

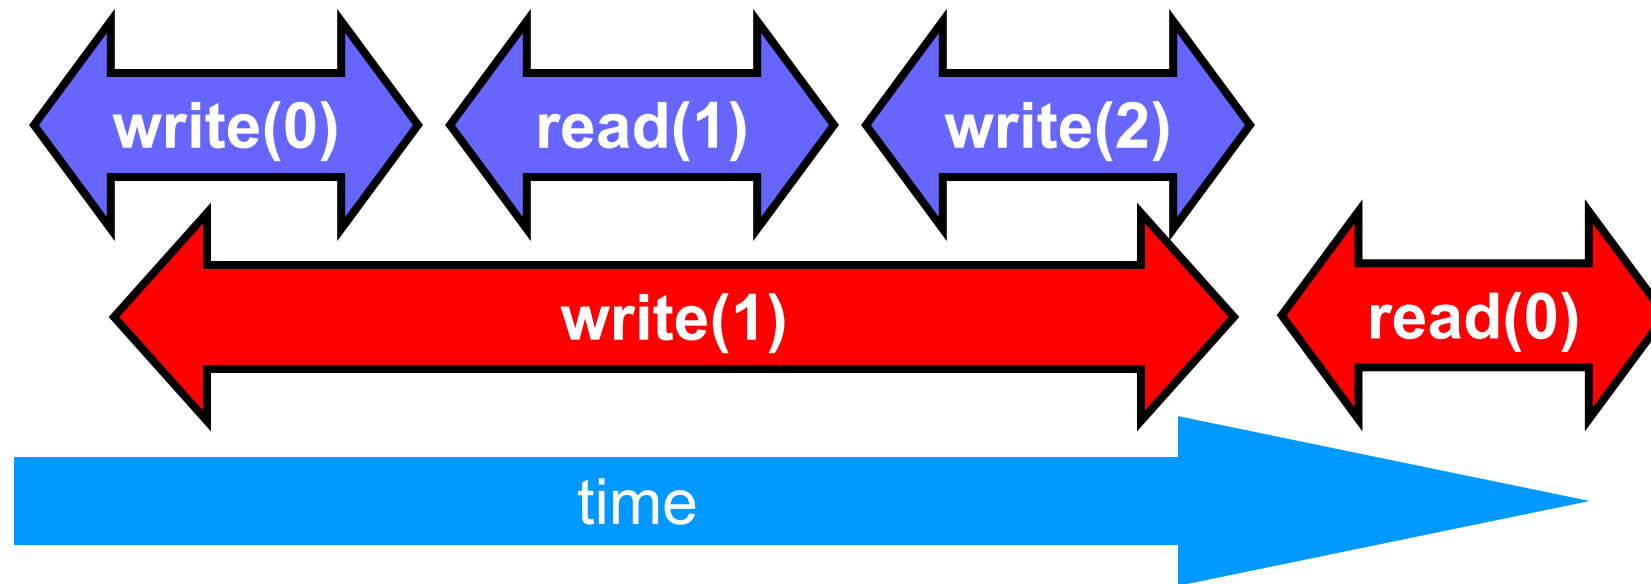


Comme ci
Comme ça

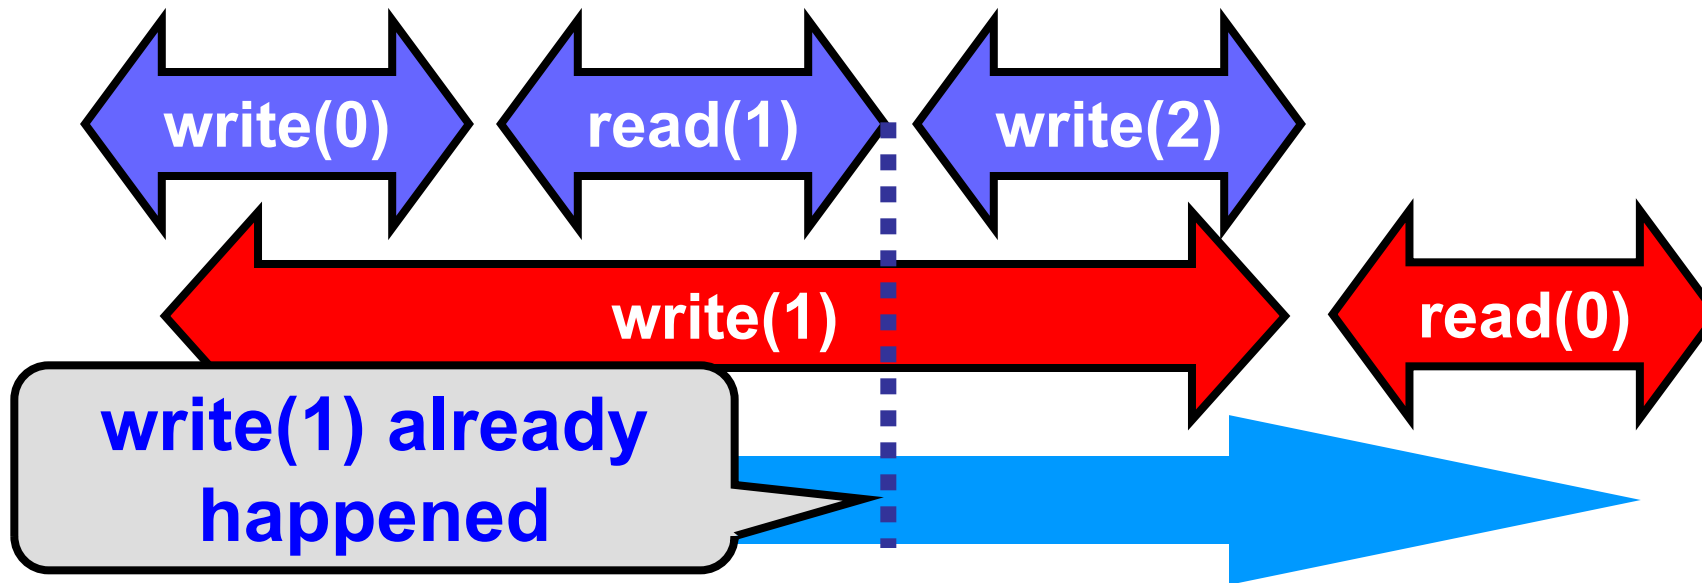
Example



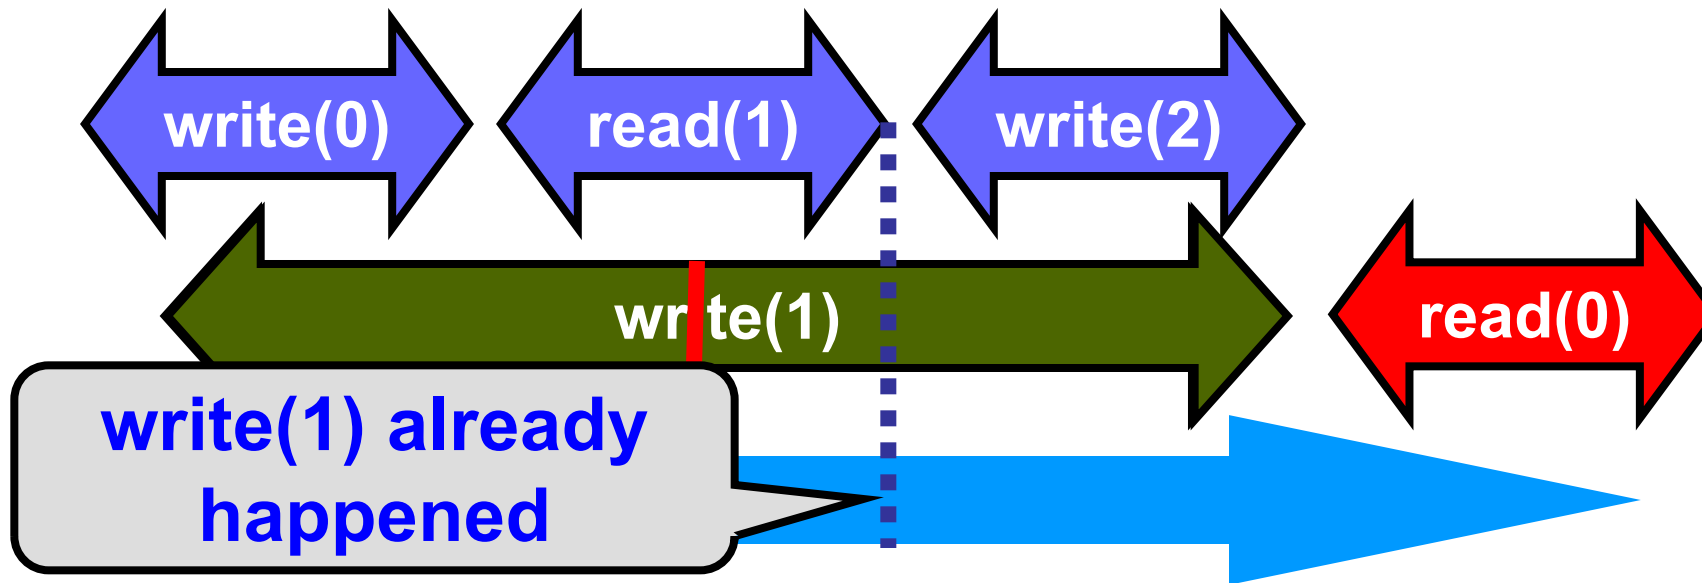
Read/Write Register Example



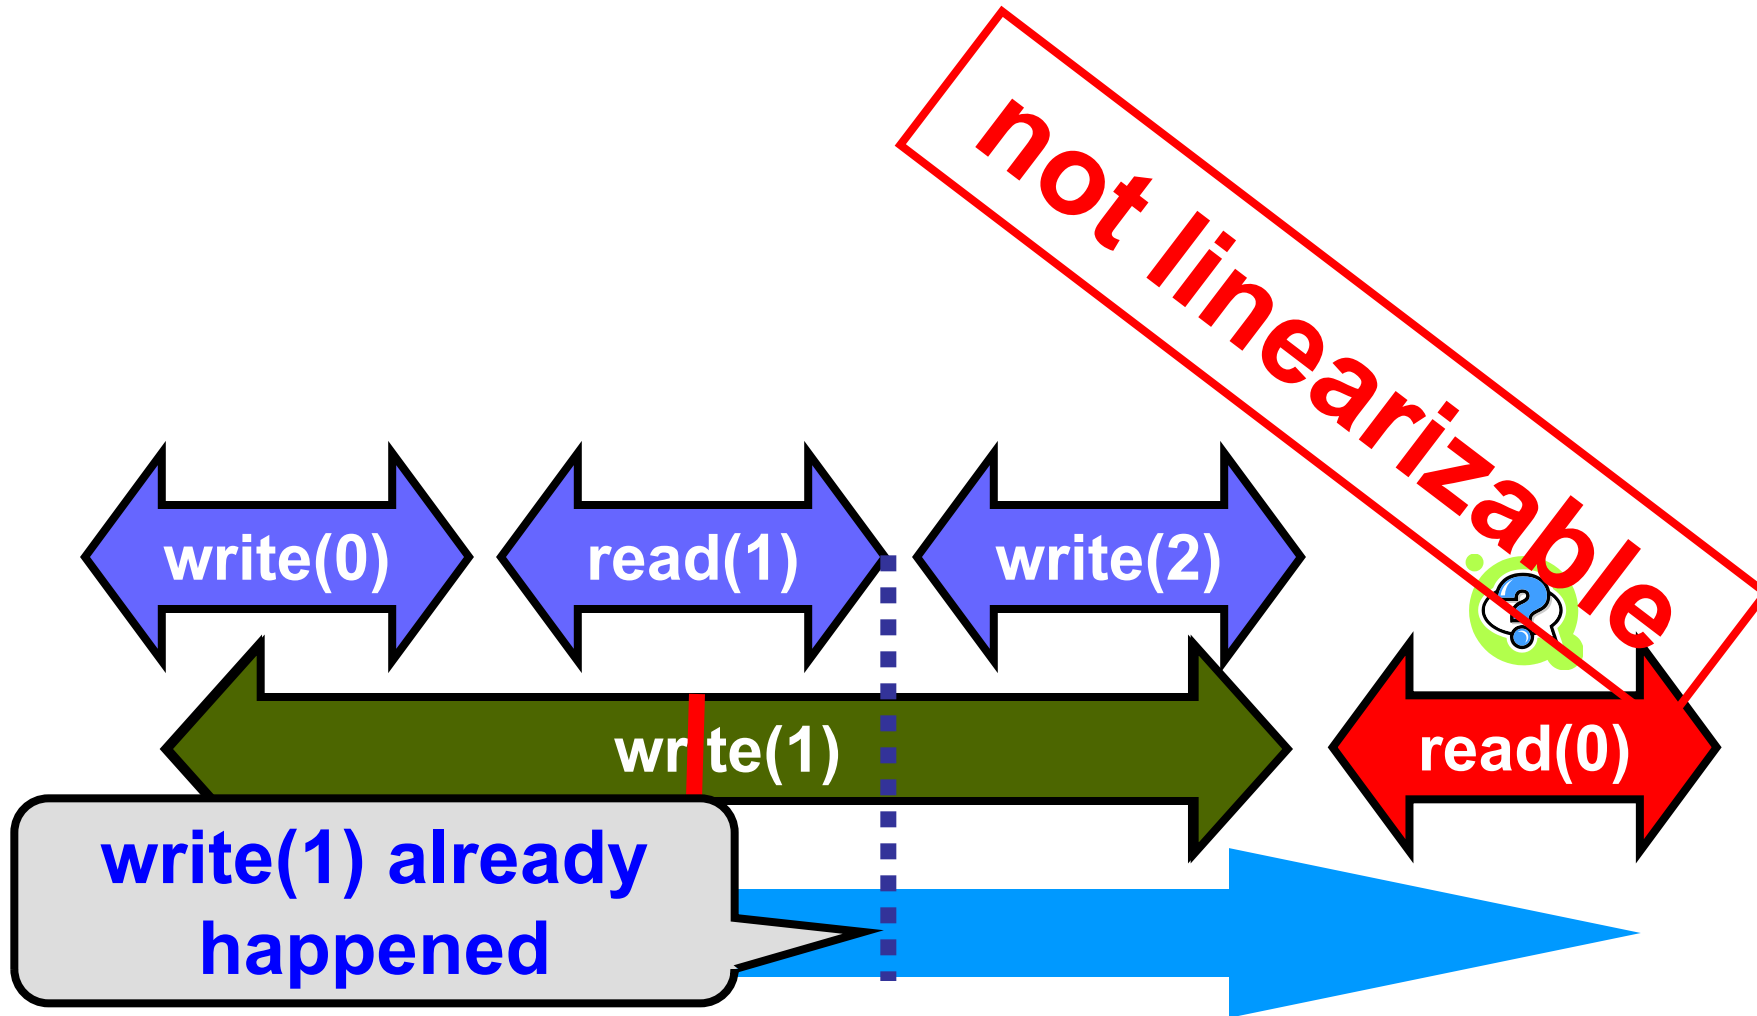
Read/Write Register Example



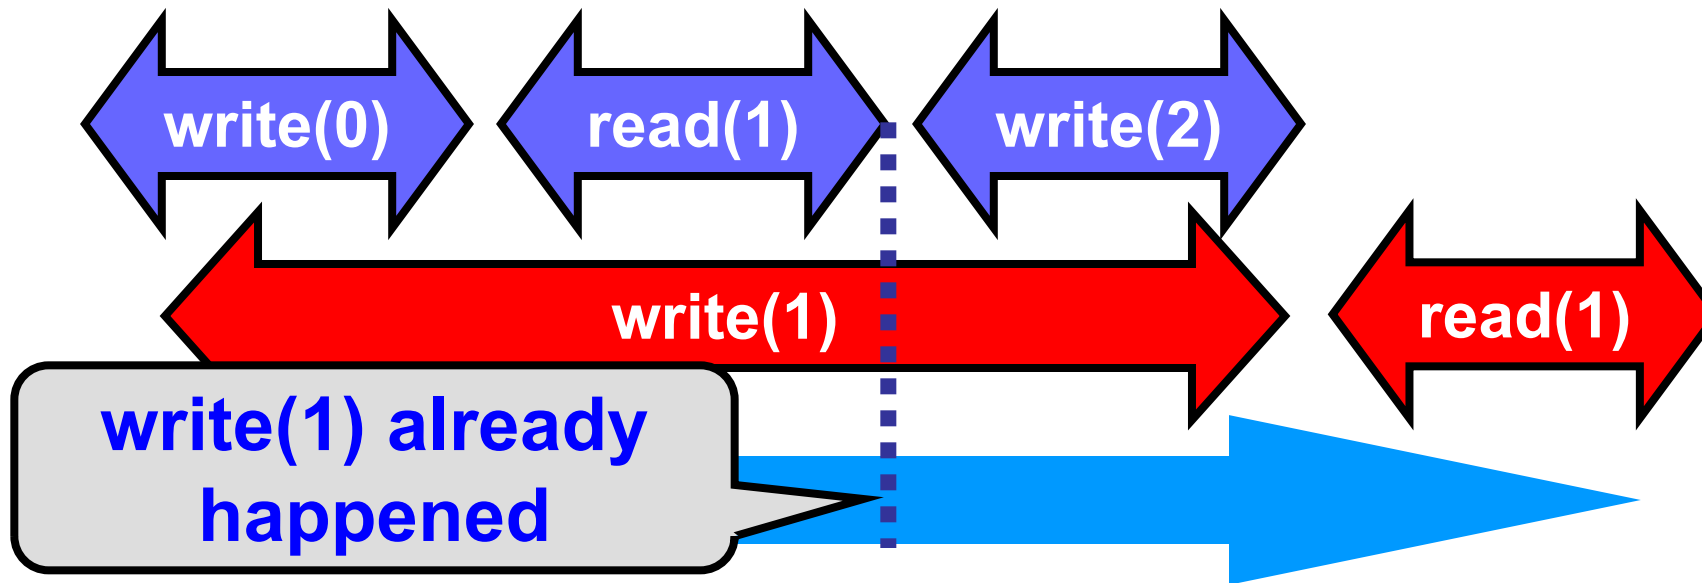
Read/Write Register Example



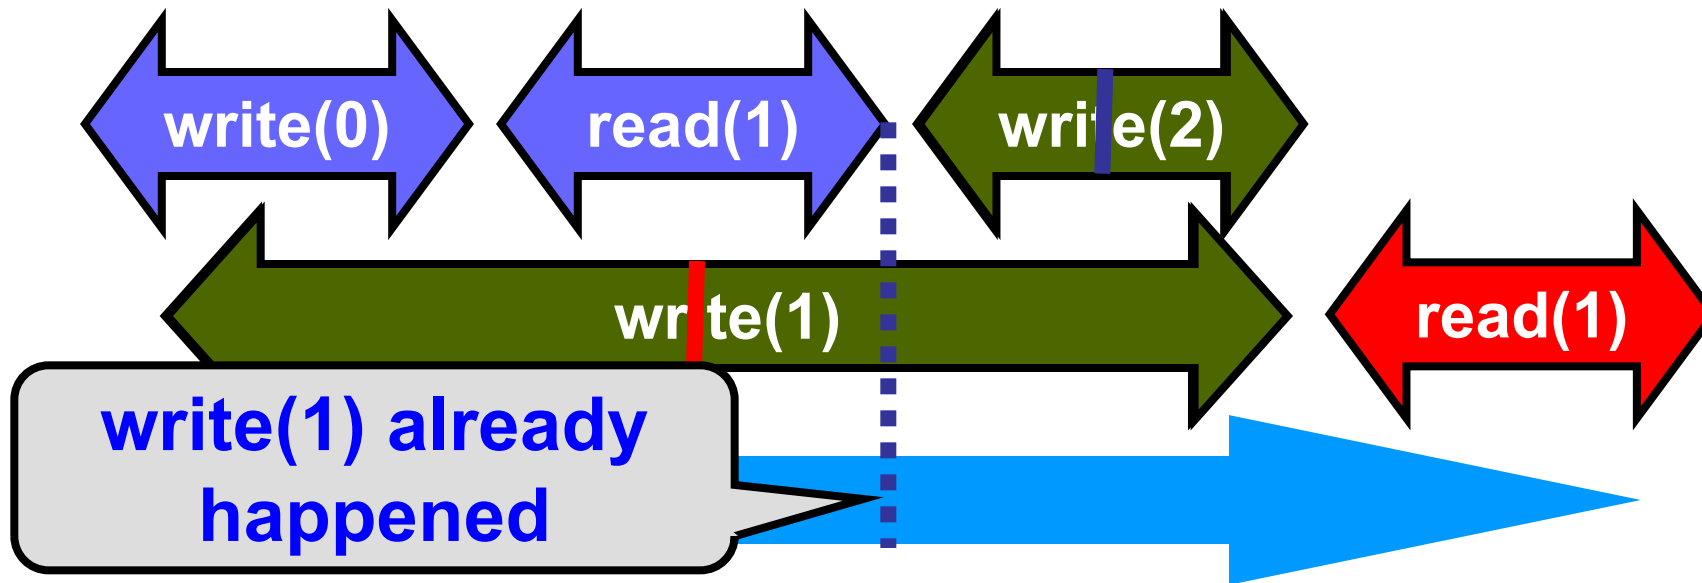
Read/Write Register Example



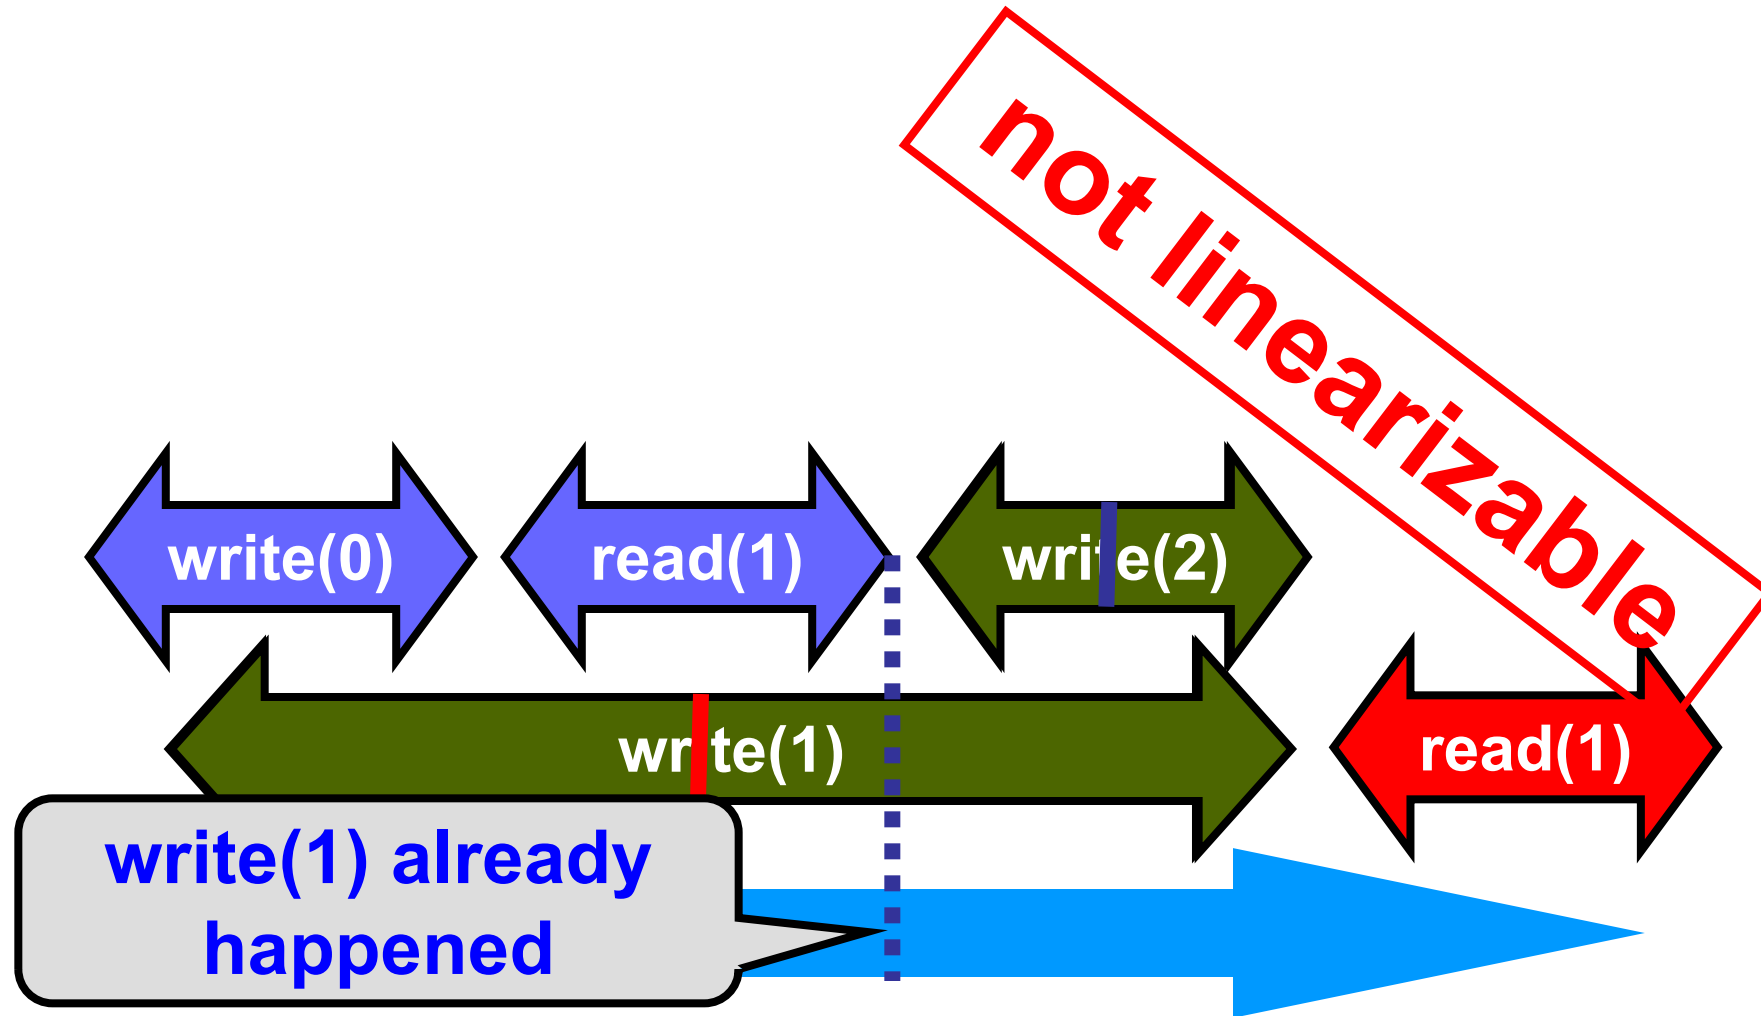
Read/Write Register Example



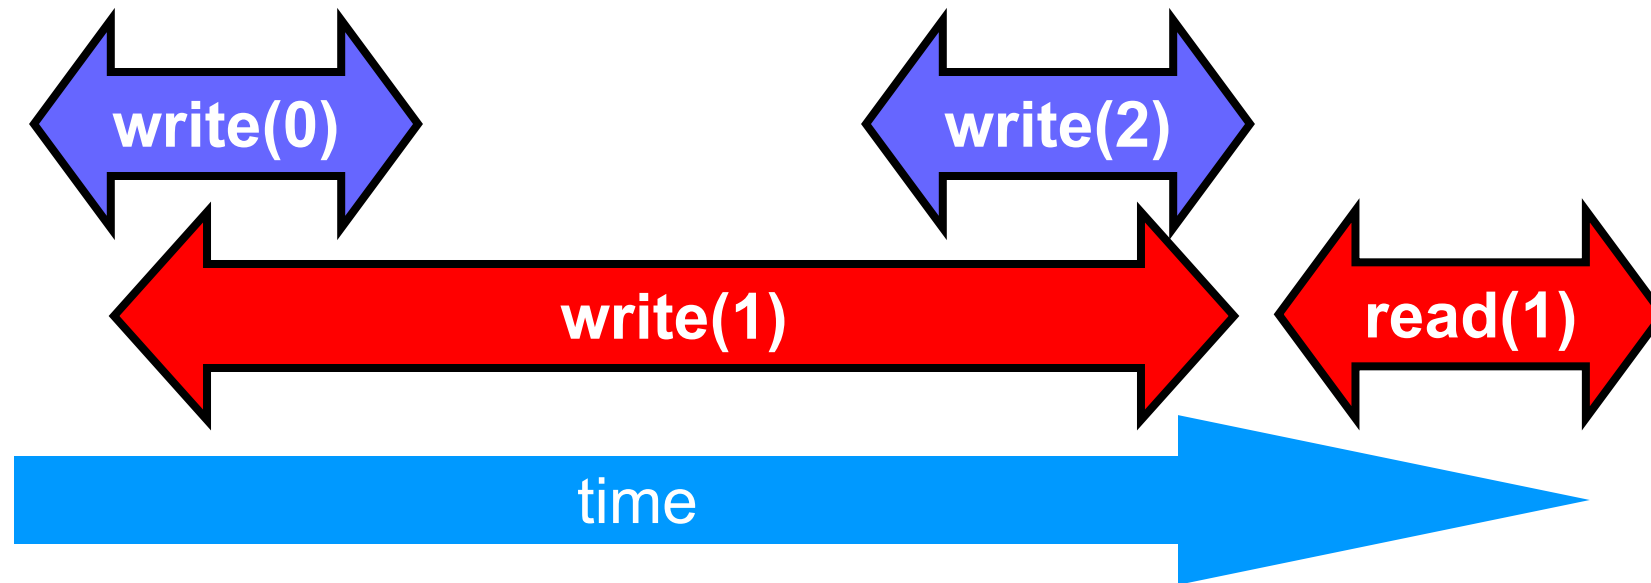
Read/Write Register Example



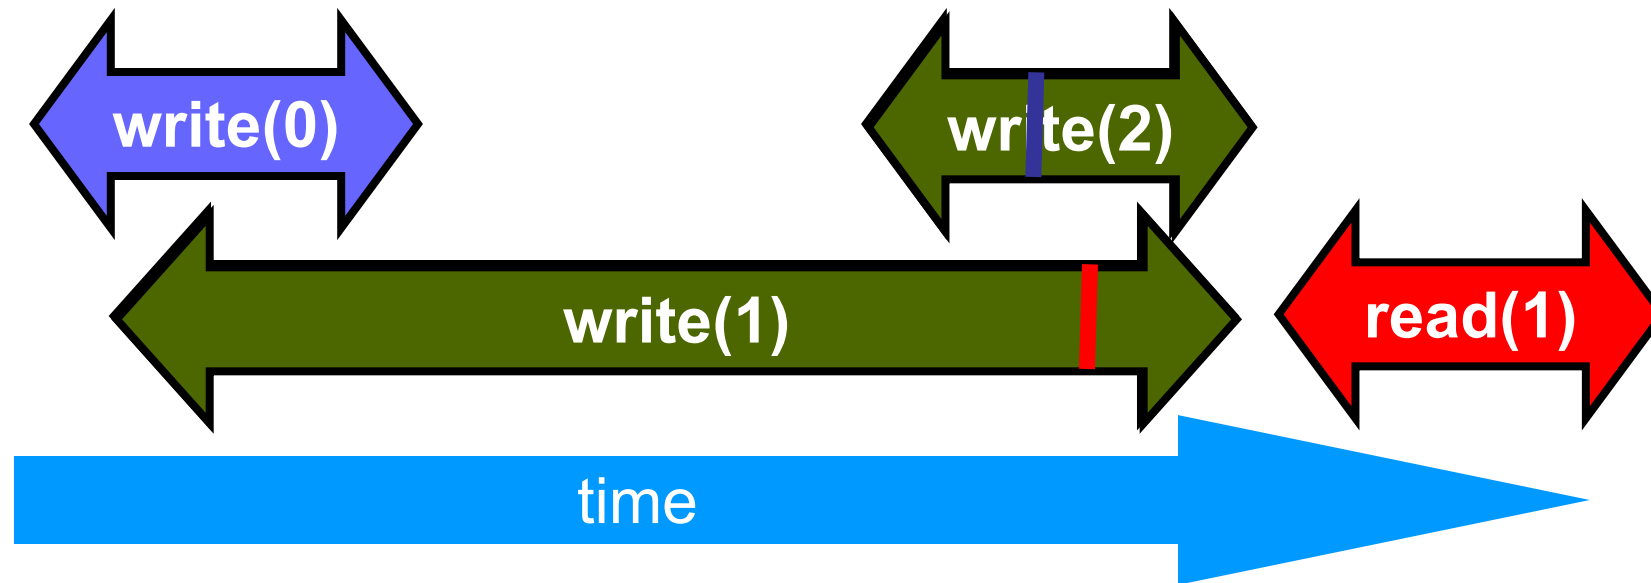
Read/Write Register Example



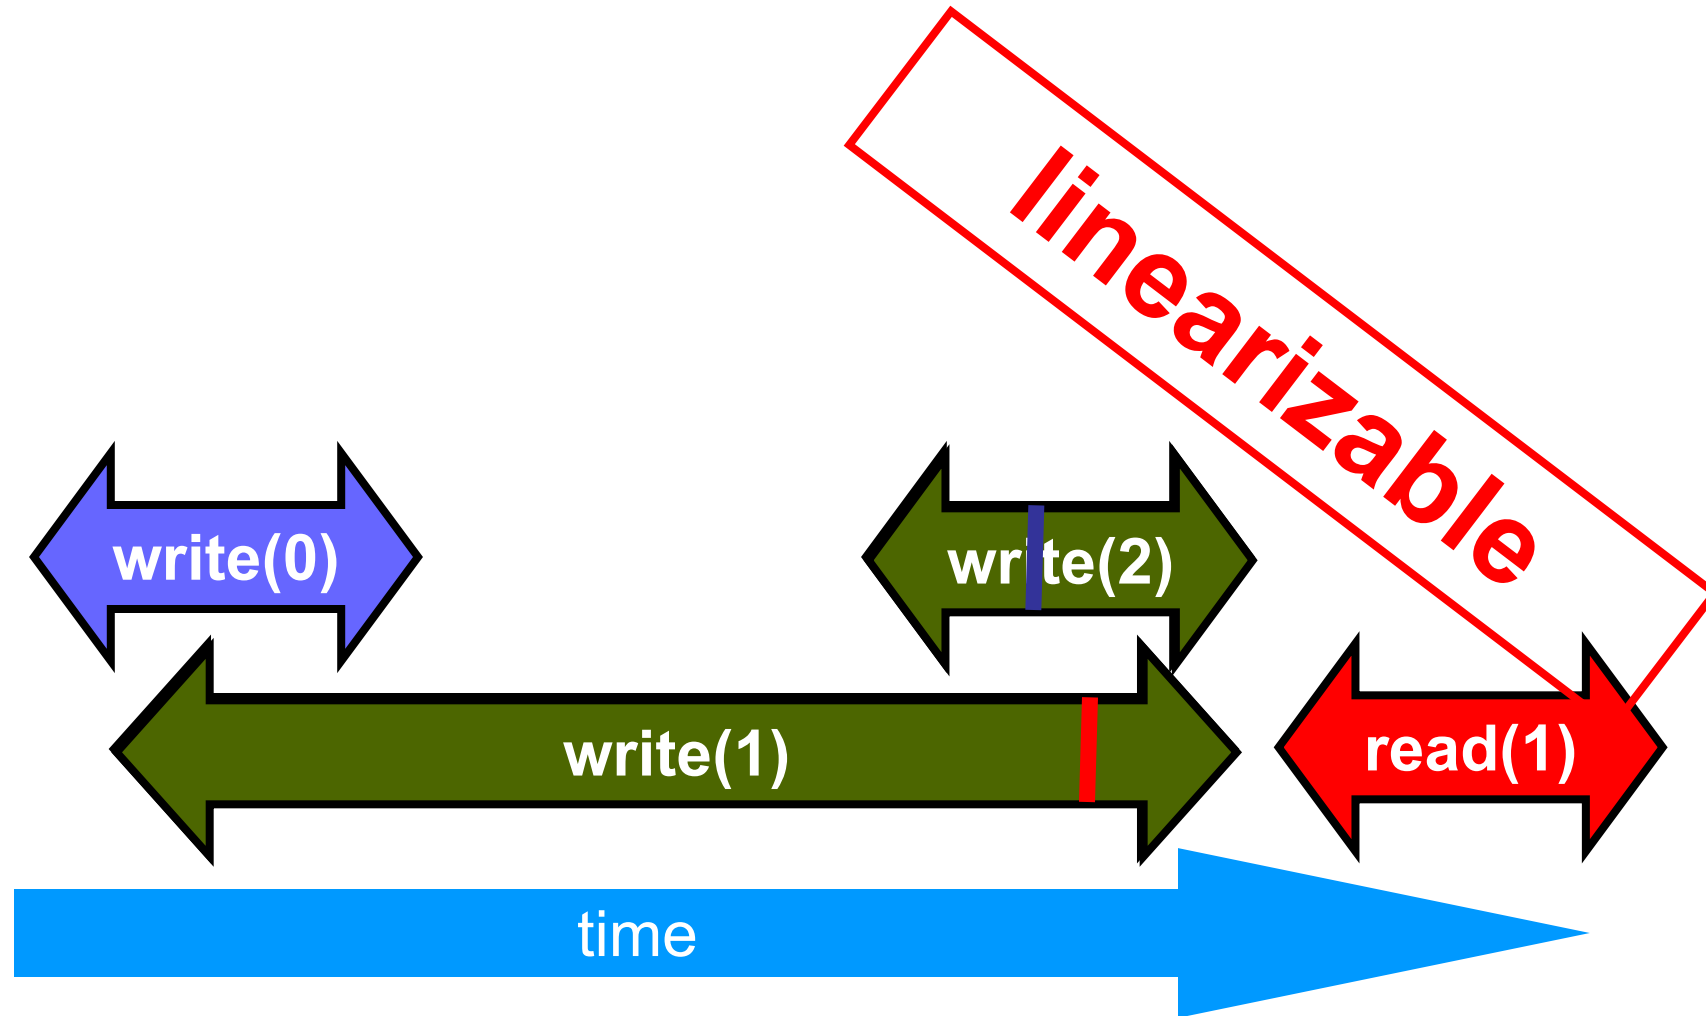
Read/Write Register Example



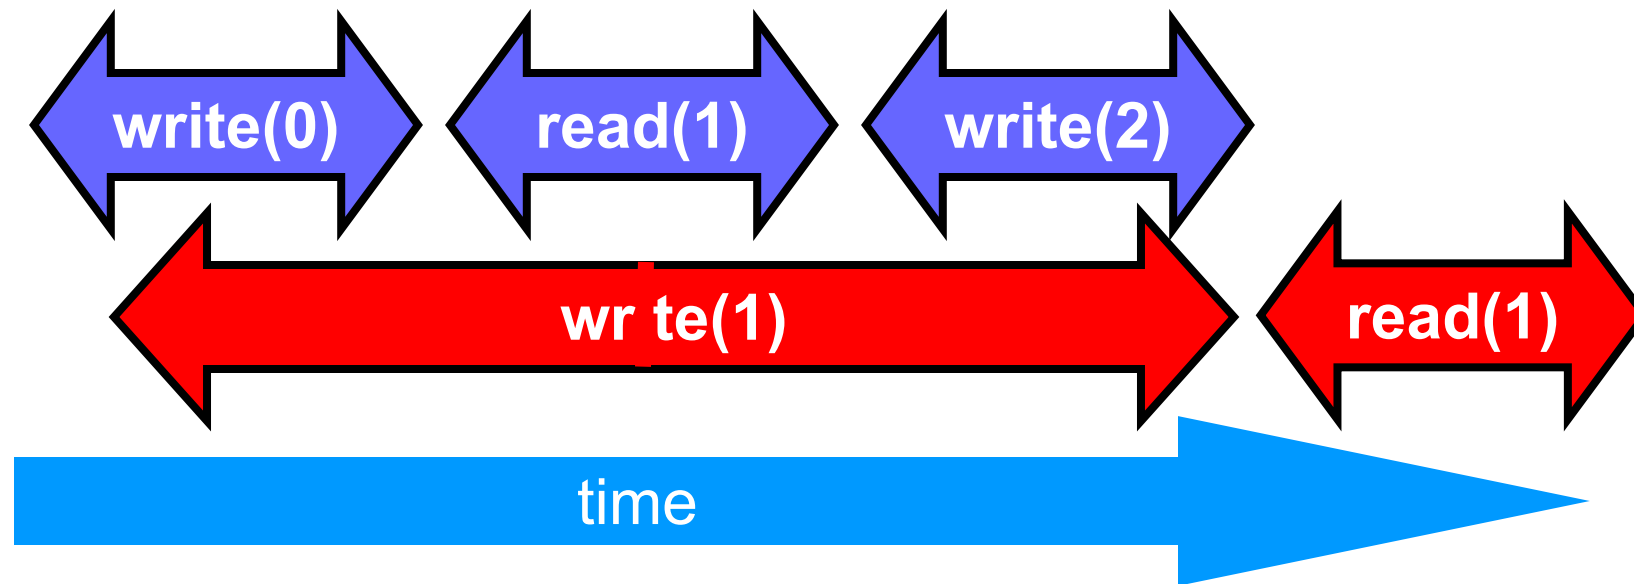
Read/Write Register Example



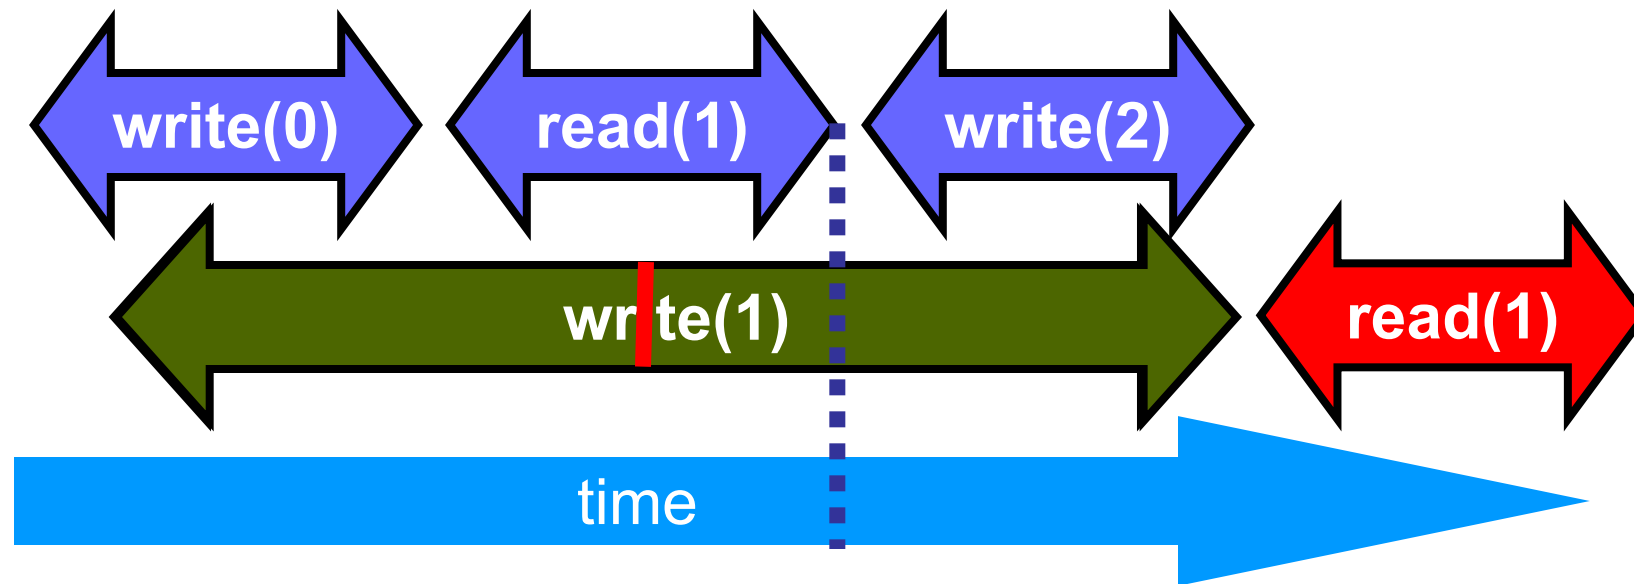
Read/Write Register Example



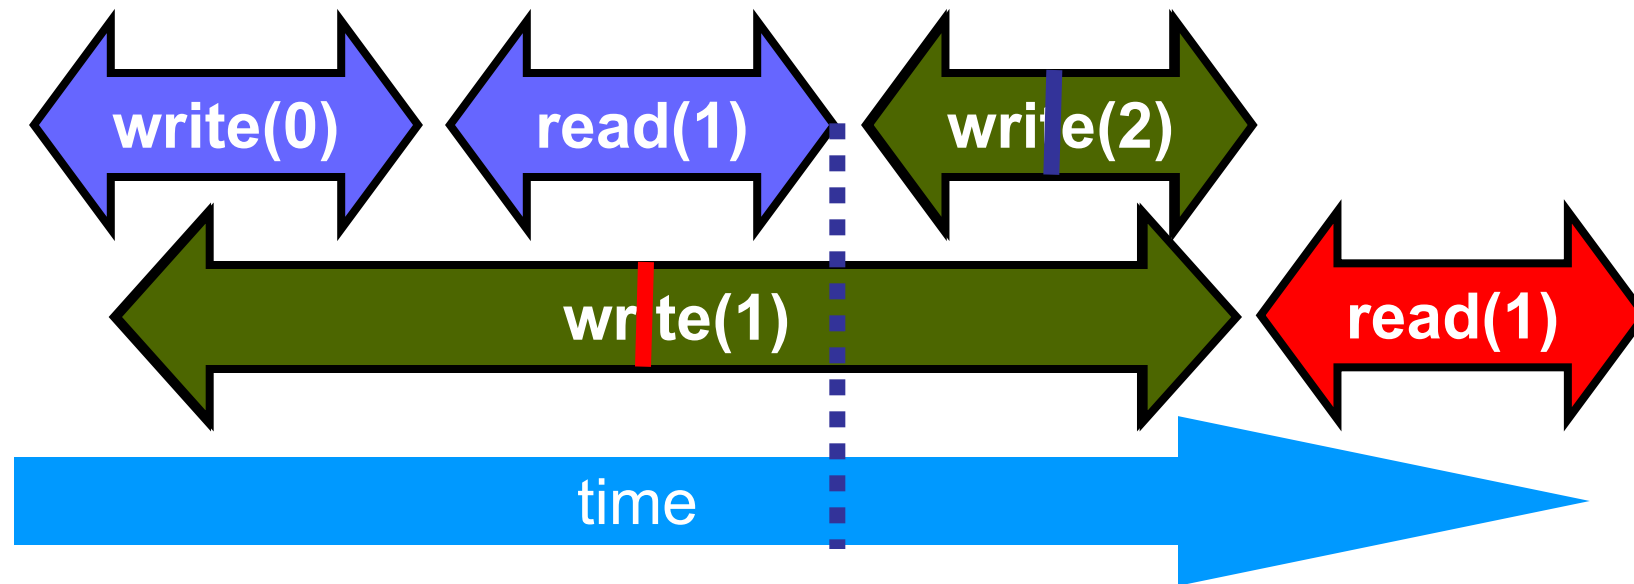
Read/Write Register Example



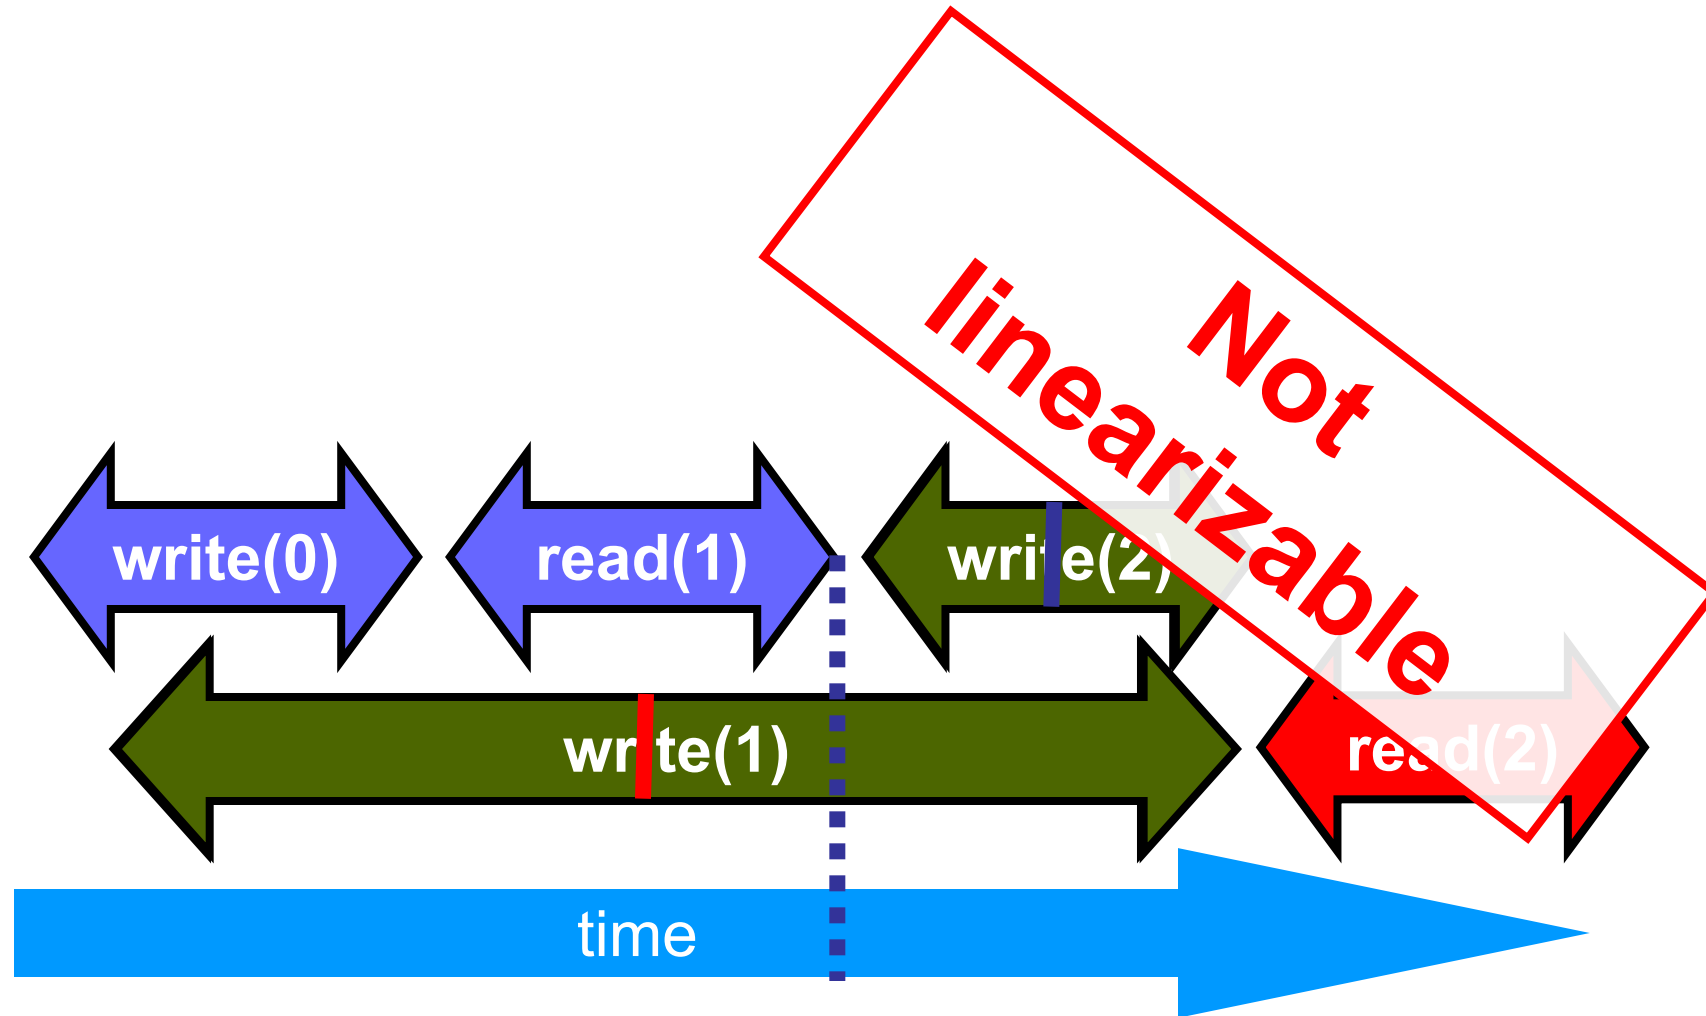
Read/Write Register Example



Read/Write Register Example



Read/Write Register Example



Talking About Executions

- Why?
 - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
 - In some cases, linearization point depends on the execution

Formal Model of Executions

- Define precisely what we mean
 - Ambiguity is bad when intuition is weak
- Allow reasoning
 - Formal
 - But mostly informal
 - In the long run, actually more important

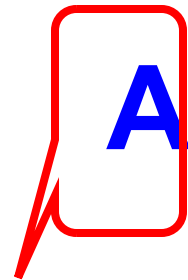
Split Method Calls into Two Events

- Invocation
 - method name & args
 - `q.enq(x)`
- Response
 - result or exception
 - `q.enq(x)` returns `void`
 - `q.deq()` returns `x`
 - `q.deq()` throws `empty`

Invocation Notation


A q.enq(x)

Invocation Notation


 **A q.enq(x)**

thread


Invocation Notation



Invocation Notation



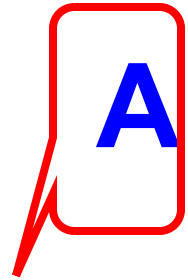
Invocation Notation



Response Notation


A q: void

Response Notation


 **A q: void**

thread


Response Notation




Response Notation



Response Notation



Response Notation



History - Describing an Execution


H =

- A q.enq(3)
- A q:void
- A q.enq(5)
- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

**Sequence of
invocations and
responses**

Definition

- Invocation & response *match* if



Object Projections

H =

A	q.enq(3)
A	q:void
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

Object Projections

A q.enq(3)

A q:void

H|q =

B q.deq()

B q:3

Thread Projections

$H =$

- A `q.enq(3)`
- A `q:void`
- B `p.enq(4)`
- B `p:void`
- B `q.deq()`
- B `q:3`

Thread Projections

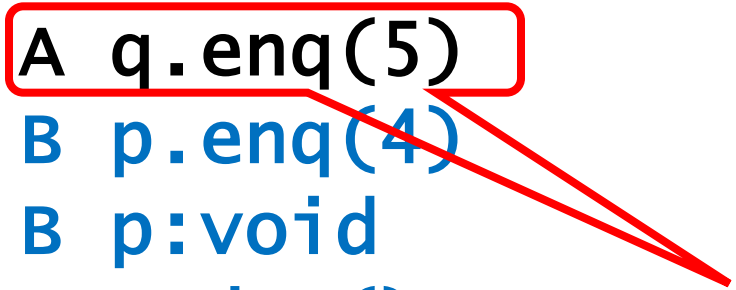
$H|B =$
B p.enq(4)
B p:void
B q.deq()
B q:3

Complete Subhistory

H =

A	q.enq(3)
A	q:void
A	q.enq(5)
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

An invocation is *pending* if it has no matching response



Complete Subhistory

H =

A	q.enq(3)
A	q:void
A	q.enq(5)
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

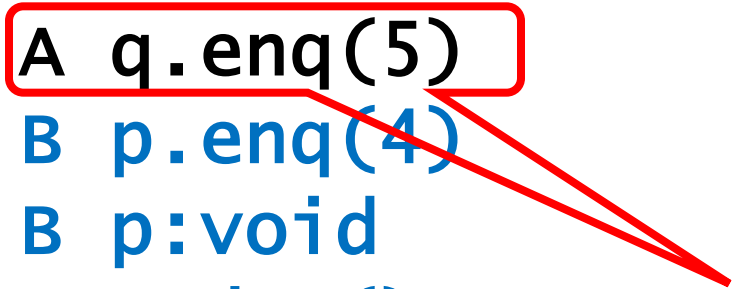
**May or may not
have taken effect**

Complete Subhistory

H =

A	q.enq(3)
A	q:void
A	q.enq(5)
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

**discard pending
invocations**



Complete Subhistory

A q.enq(3)
A q:void

Complete(H) = B p.enq(4)
B p:void
B q.deq()
B q:3

Sequential Histories

A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
A q:enq(5)

Sequential Histories

A q.enq(3)

A q:void

match

B p.enq(4)

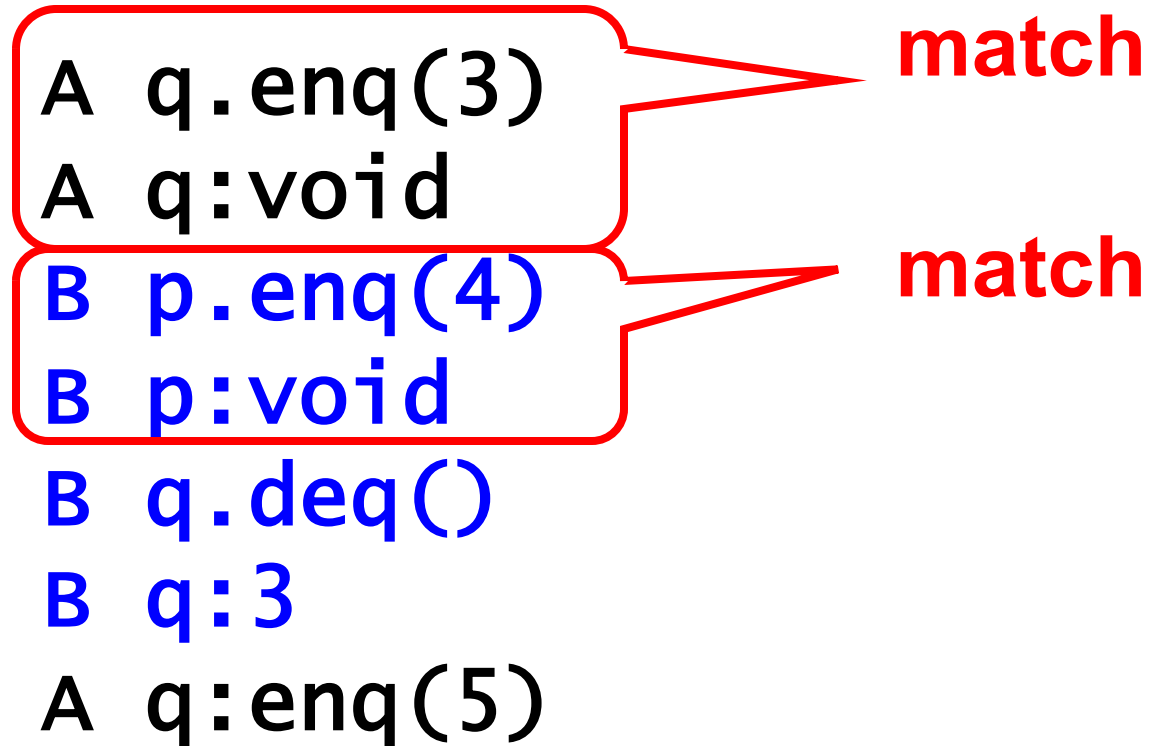
B p:void

B q.deq()

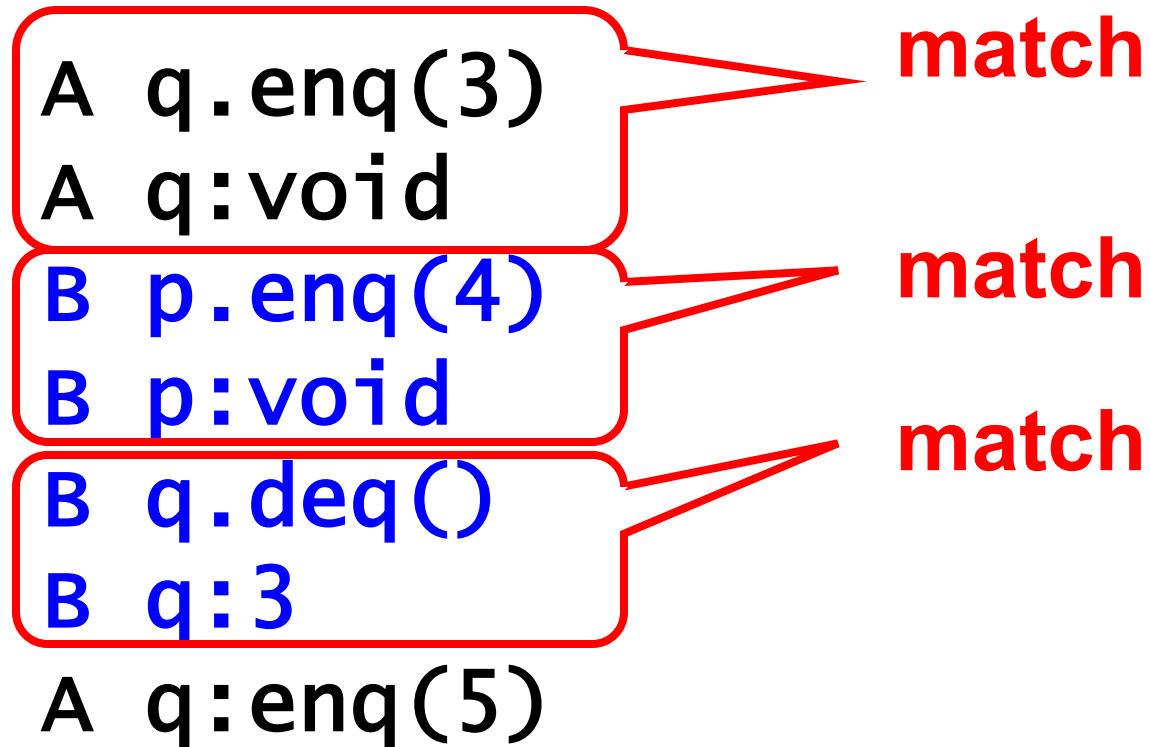
B q:3

A q:enq(5)


Sequential Histories



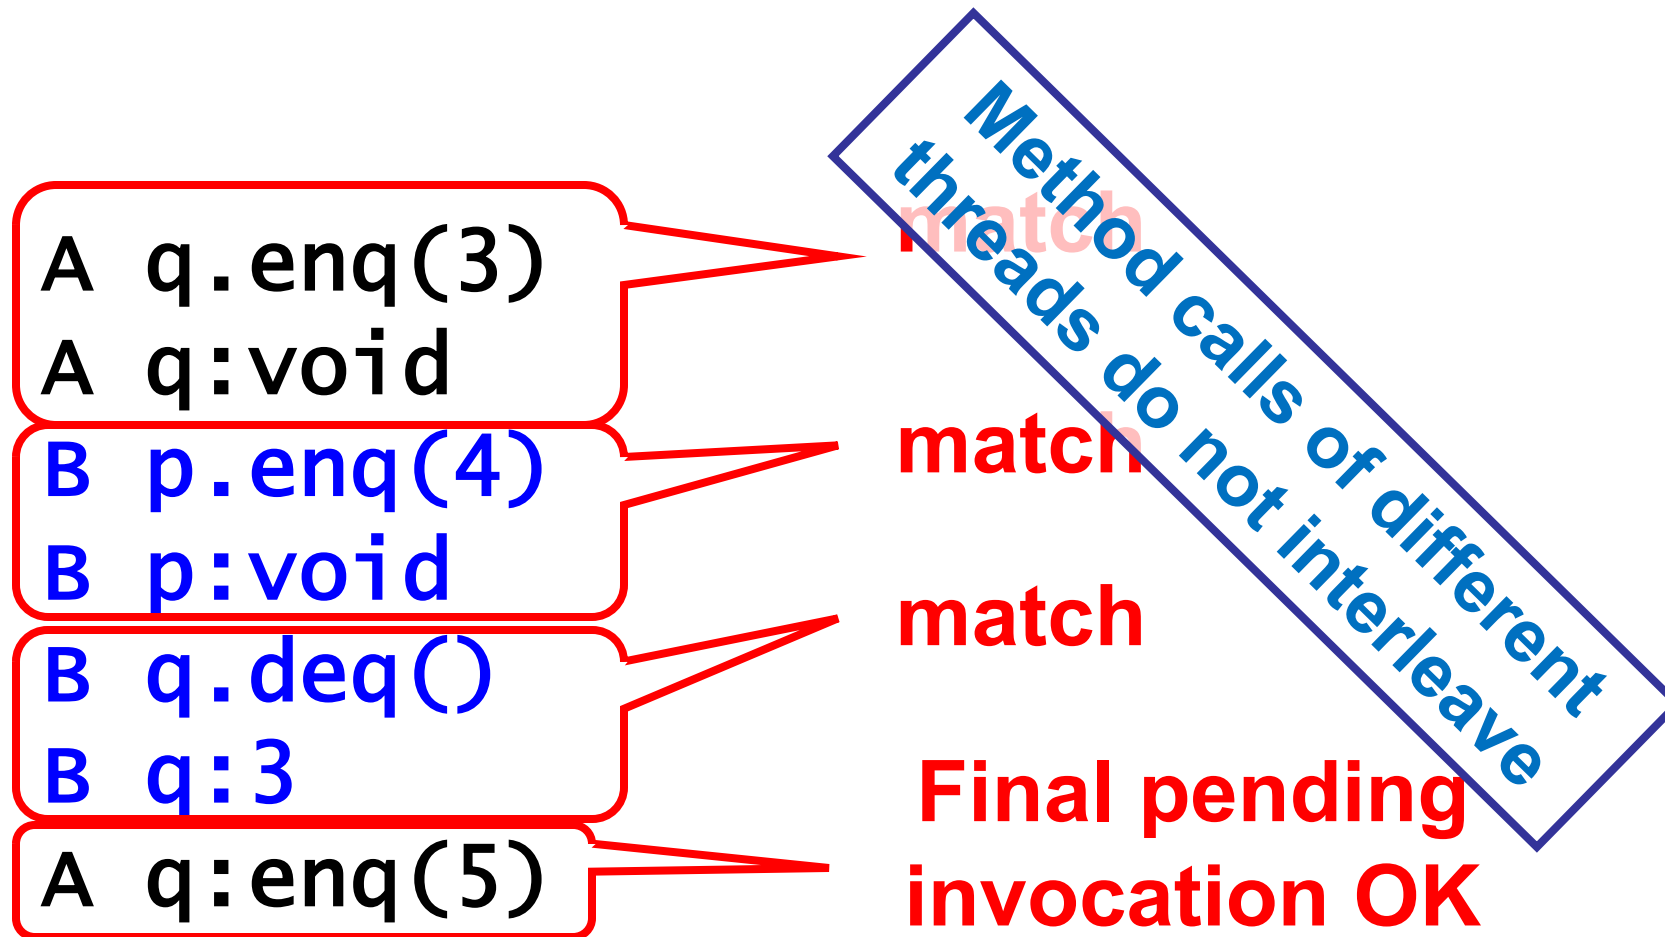
Sequential Histories



Sequential Histories



Sequential Histories



Well-Formed Histories

H=

A	q.enq(3)
B	p.enq(4)
B	p:void
B	q.deq()
A	q:void
B	q:3

Well-Formed Histories

Per-thread projections
sequential

$H =$	A q.enq(3)	$H \mid B =$	B p.enq(4)
	B p.enq(4)		B p:void
	B p:void		B q.deq()
	B q.deq()		B q:3
	A q:void		
	B q:3		

Well-Formed Histories

Per-thread projections
sequential

$H =$

A	q.enq(3)
B	p.enq(4)
B	p:void
B	q.deq()
A	q:void
B	q:3

$H | B =$

B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

$H | A =$

A	q.enq(3)
A	q:void

Equivalent Histories

Threads see the same
thing in both

$$\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$$

H=

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```

G=

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```

Sequential Specifications

- A sequential specification is some way of telling whether a
 - Single-thread, single-object history
 - Is legal
- For example:
 - Pre and post-conditions
 - But plenty of other techniques exist ...

Legal Histories

- A sequential (multi-object) history H is legal if
 - For every object x
 - $H|x$ is in the sequential spec for x

Precedence

A q.enq(3)

B p.enq(4)


B p.void

A q:void

B q.deq()


B q:3

A method call **precedes**
another if response
event precedes
invocation event




Non-Precedence

```
A q.enq(3)
B p.enq(4)
B p.void
B q.deq()
A q:void
B q:3
```




Some method calls
overlap one another



Notation

- Given
 - History H
 - method executions m_0 and m_1 in H
- We say $m_0 \rightarrow_H m_1$, if
 - m_0 precedes m_1
- Relation $m_0 \rightarrow_H m_1$ is a
 - Partial order
 - Total order if H is sequential



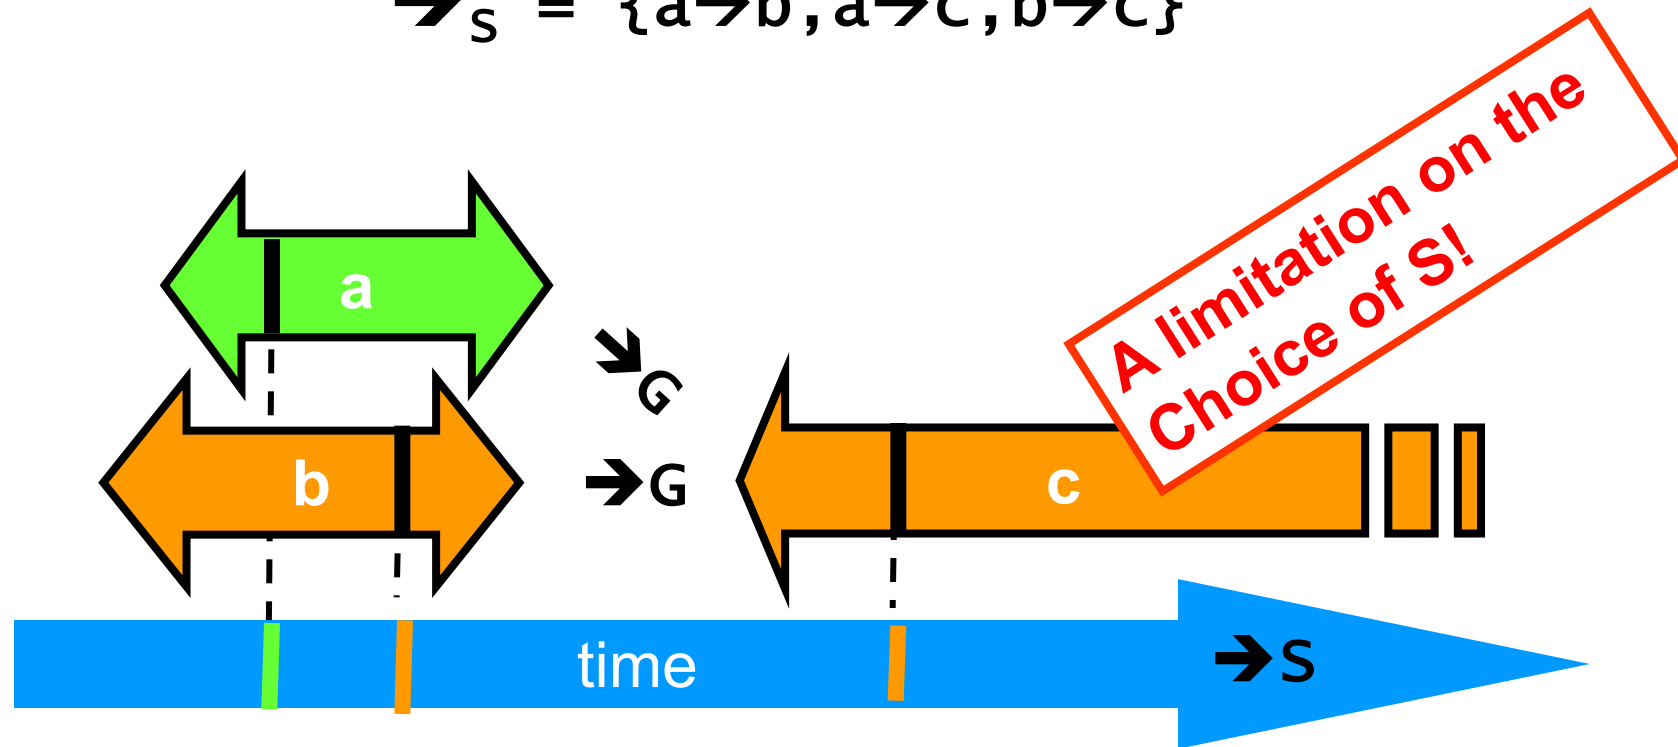
Linearizability

- History H is *linearizable* if it can be extended to G by
 - Appending zero or more responses to pending invocations
 - Discarding other pending invocations
- So that G is equivalent to
 - Legal sequential history S
 - where $\rightarrow_G \subset \rightarrow_S$

Ensuring $\rightarrow_G \subset \rightarrow_S$

$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$




Remarks

- Some pending invocations
 - Took effect, so keep them
 - Discard the rest
- Condition $\rightarrow_{\mathbf{G}} \subset \rightarrow_{\mathbf{S}}$
 - Means that **S** respects “real-time order” of **G**

Example

A q.enq(3)
B q.enq(4)
B q:void
B q.deq()
B q:4
B q:enq(6)



Example

A q.enq(3)

B q.enq(4)

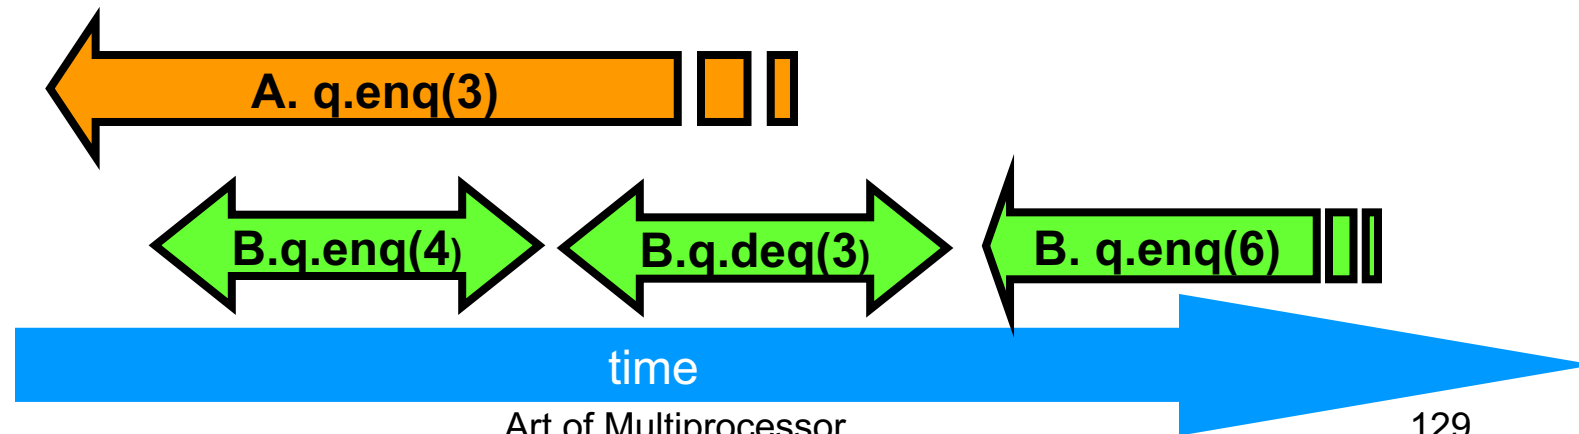
B q:void

B q.deq()

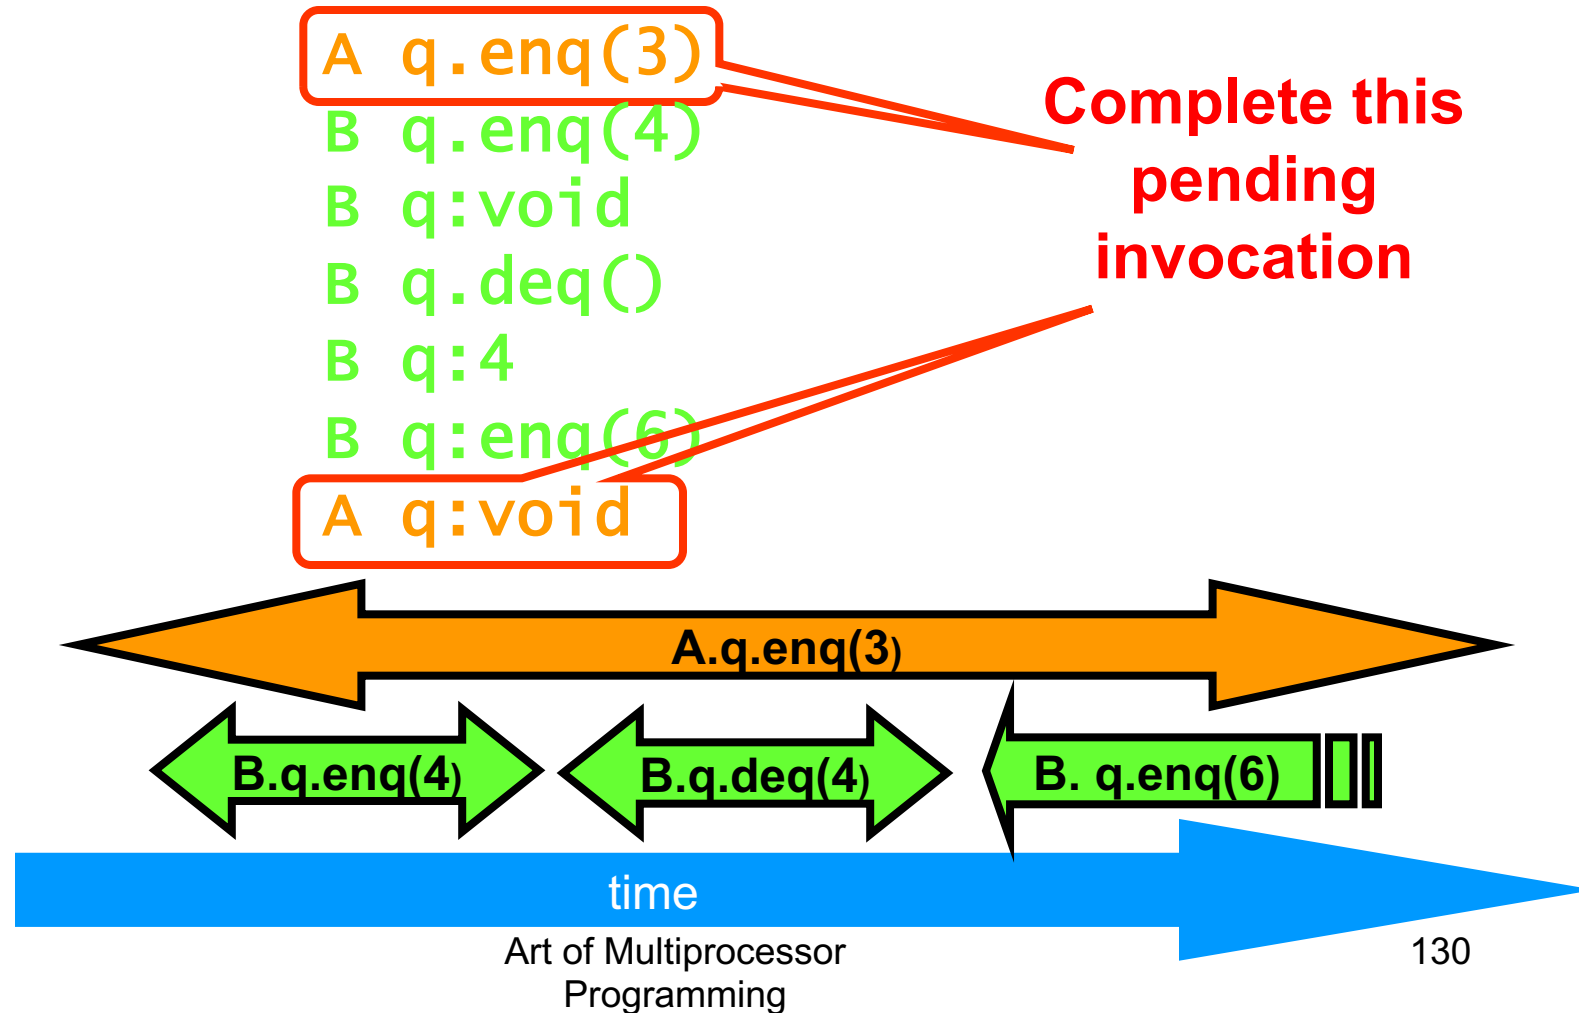
B q:4

B q:enq(6)

Complete this
pending
invocation



Example



Example

discard this one

A q.enq(3)

B q.enq(4)

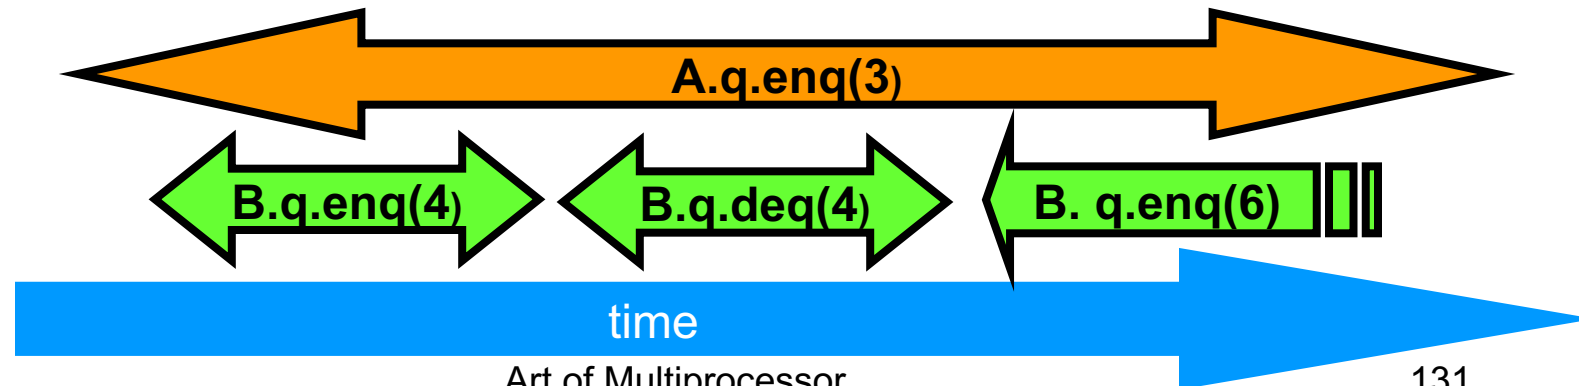
B q:void

B q.deq()

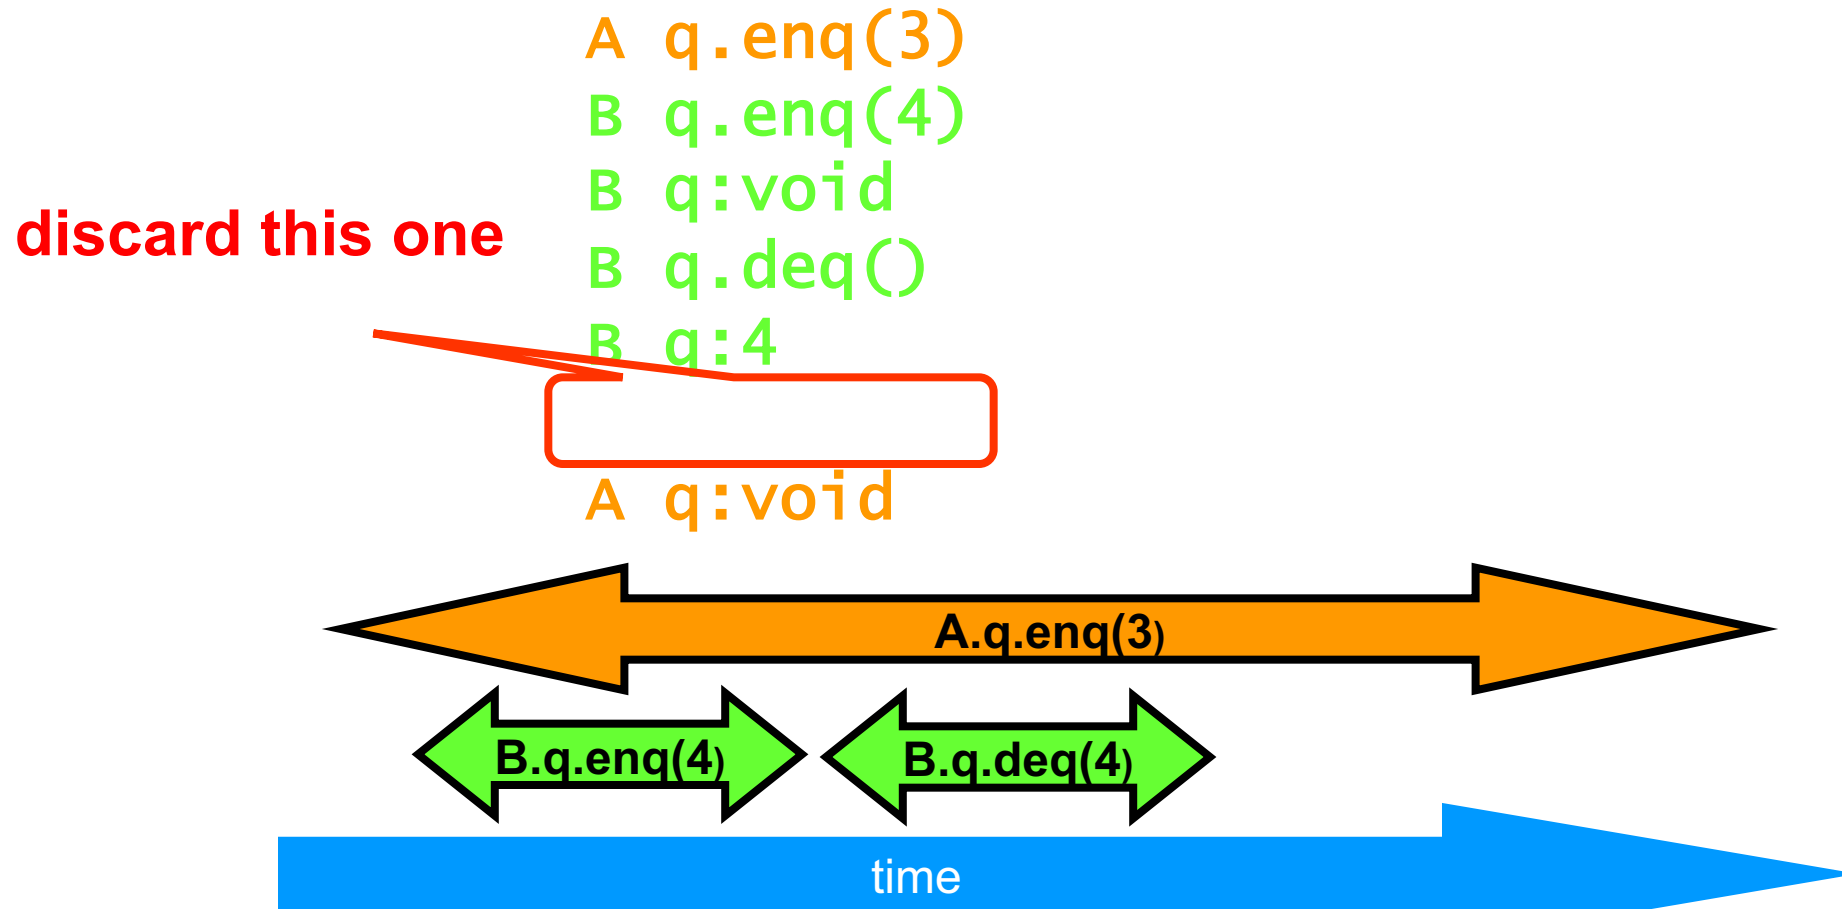
B q:4

B q:enq(6)

A q:void

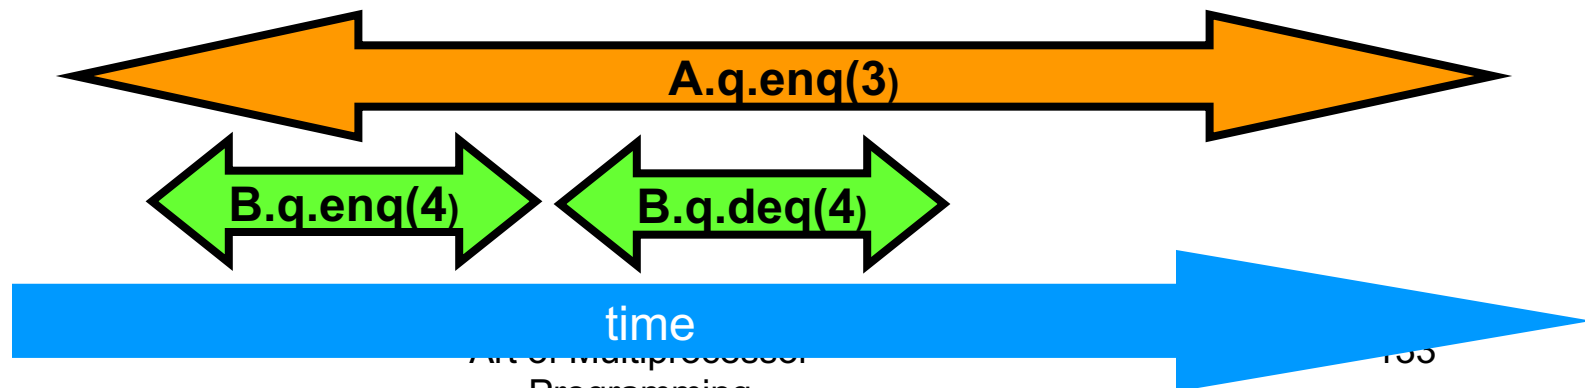


Example



Example

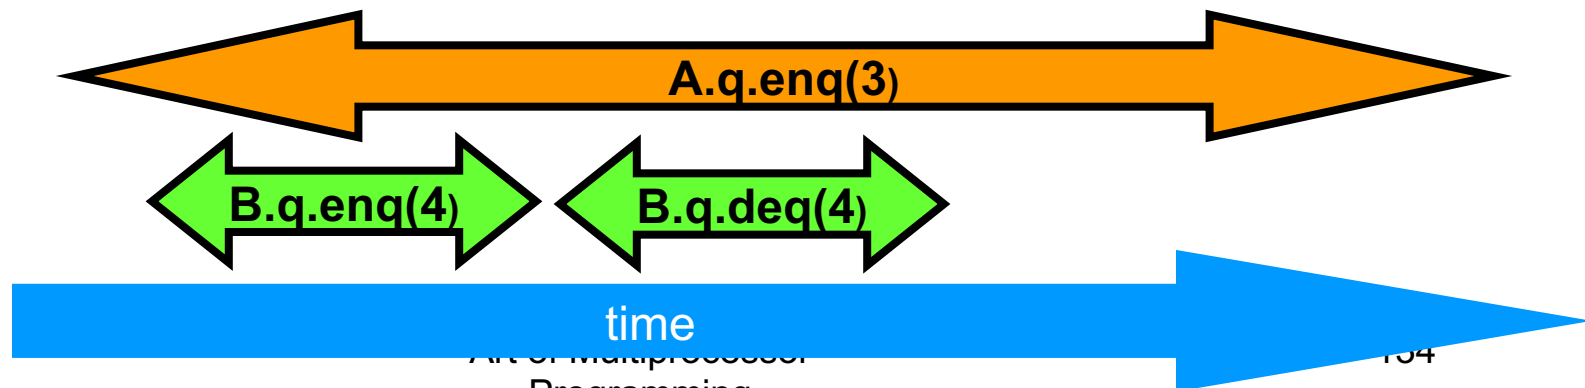
A q.enq(3)
B q.enq(4)
B q:void
B q.deq()
B q:4
A q:void



Example

A q.enq(3)
B q.enq(4)
B q:void
B q.deq()
B q:4
A q:void

B q.enq(4)
B q:void
A q.enq(3)
A q:void
B q.deq()
B q:4

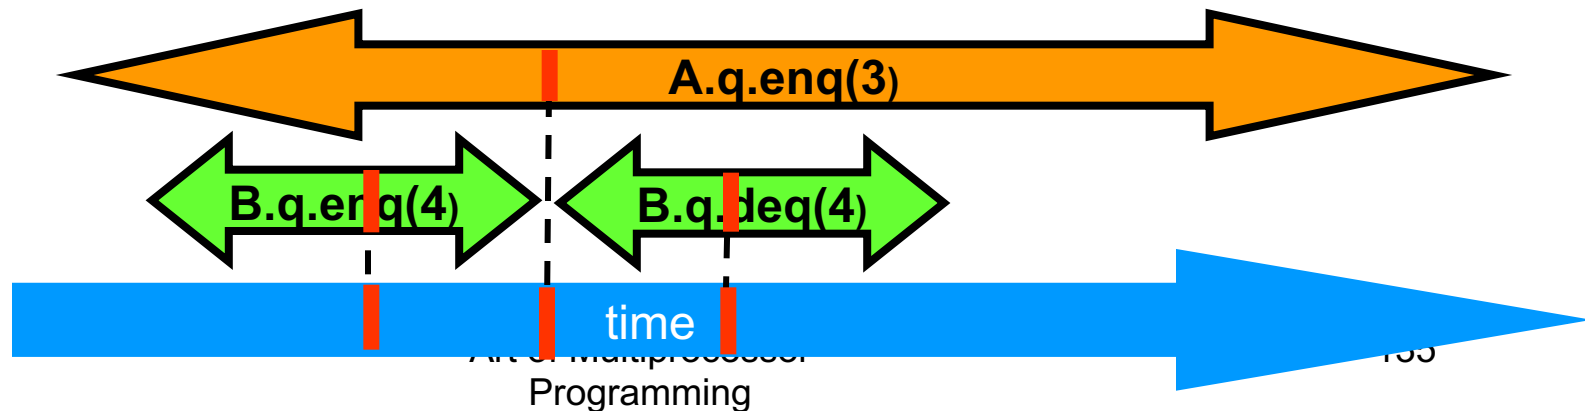


Example

Equivalent sequential history

A q.enq(3)
B q.enq(4)
B q:void
B q.deq()
B q:4
A q:void

B q.enq(4)
B q:void
A q.enq(3)
A q:void
B q.deq()
B q:4



Composability Theorem


- History H is linearizable if and only if
 - For every object x
 - $H|x$ is linearizable
- We care about objects only!
 - (Materialism?)

Why Does Composability Matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects

Reasoning About Linearizability: Locking

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```




Reasoning About Linearizability: Locking

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Linearization points
are when locks are
released

More Reasoning: Wait-free

```
public class waitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```



More Reasoning: Wait-free

**Remember that there
is only one enqueuer
and only one dequeuer**

```
public class WaitFreeQueue {
```

```
    int head = 0;
    int tail = 0;
    Item[] items = new Item[capacity];
```

```
    public void enq(Item x) {
        if (tail == capacity) throw
            new FullException();
        items[tail % capacity] = x;
    }
```


```
    public Item deq() {
        if (tail == head) throw
            new EmptyException();
        Item item = items[head % capacity];
        return item;
    }
```

```
}}
```

Linearization order is
order head and tail
fields modified

tail++;

head++;



Strategy

- Identify one atomic step where method “happens”
 - Critical section
 - Machine instruction
- Doesn't always work
 - Might need to define several different steps for a given method

Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don't leave home without it