
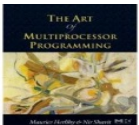


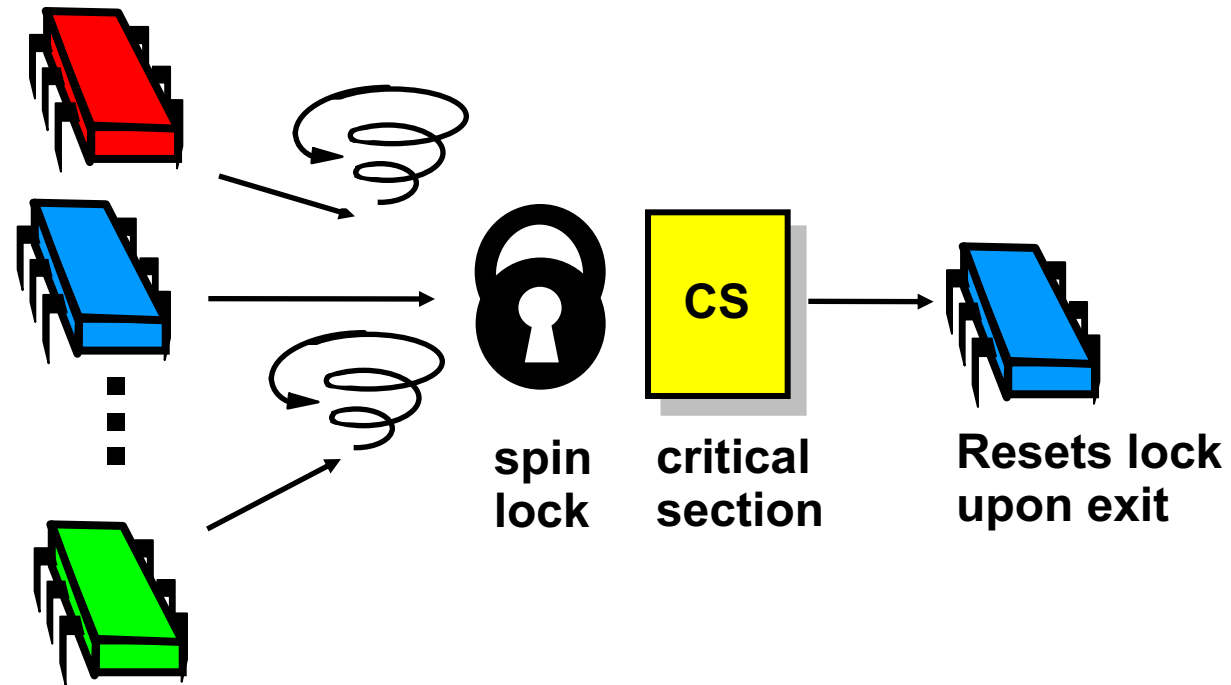
Linked Lists: Locking, Lock-Free, and Beyond ...



Hyungsoo Jung



Last Lecture: Spin-Locks



Today: Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Mostly fixed by Queue locks



Today: Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Mostly fixed by Queue locks
- Should increase throughput
 - Not possible if inherently sequential
 - Surprising things are parallelizable



Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using queue locks



Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using queue locks
 - Easy to reason about
 - In simple cases



Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using queue locks
 - Easy to reason about
 - In simple cases
- So, are we done?



Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”



Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse



Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse
- So why even use a multiprocessor?
 - Well, some apps inherently parallel ...



This Lecture

- Introduce four “patterns”
 - Bag of tricks ...
 - Methods that work more than once ...



This Lecture

- Introduce four “patterns”
 - Bag of tricks ...
 - Methods that work more than once ...
- For highly-concurrent objects
 - Concurrent access
 - More threads, more throughput



First:

Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components
- Methods conflict when they access
 - The same component ...
 - At the same time



Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over
- Evaluation
 - Usually cheaper than locking, but
 - Mistakes are expensive



Third:

Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted
 - Physical removal
 - Do what needs to be done



Fourth:

Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support
- Disadvantages
 - Complex
 - Sometimes high overhead



Linked List

- Illustrate these patterns ...
- Using a list-based Set
 - Common application
 - Building block for other apps



Set Interface

- Unordered collection of items
- No duplicates
- Methods
 - `add(x)` put `x` in set
 - `remove(x)` take `x` out of set
 - `contains(x)` tests if `x` in set



List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```



List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

Add item to set

List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
public boolean remove(T x);  
    public boolean contains(T x);  
}
```

Remove item from set

List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

Is item in set?

List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

List Node

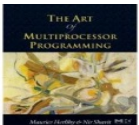
```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

item of interest

List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

Usually hash code




List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

Reference to next node

The List-Based Set



Sorted with Sentinel nodes
(min & max possible keys)

Reasoning about Concurrent Objects

- Invariant
 - Property that always holds
- Established because
 - True when object is **created**
 - Truth **preserved** by each method
 - Each **step** of each method



Specifically ...

- Invariants preserved by
 - `add()`
 - `remove()`
 - `contains()`
- Most steps are trivial
 - Usually one step tricky
 - Often linearization point



Interference

- Invariants make sense only if
 - methods considered
 - are the only modifiers



Interference

- Invariants make sense only if
 - methods considered
 - are the only modifiers
- Language encapsulation helps
 - List nodes not visible outside class




Interference

- Freedom from interference needed even for removed nodes
 - Some algorithms traverse removed nodes
 - Careful with **malloc()** & **free()** !
- Garbage collection helps here




Abstract Data Types

- Concrete representation:



- Abstract Type:
 - $\{a, b\}$




Abstract Data Types

- Meaning of rep given by abstraction map

– S(

) = {a, b}



Rep Invariant

- Which concrete values meaningful?
 - Sorted?
 - Duplicates?
- Rep invariant
 - Characterizes legal concrete reps
 - Preserved by methods
 - Relied on by methods



Blame Game

- Rep invariant is a **contract**
- Suppose
 - **add()** leaves behind 2 copies of x
 - **remove()** removes only 1
- Which is incorrect?



Blame Game

- Suppose
 - **add()** leaves behind 2 copies of x
 - **remove()** removes only 1
- Which is incorrect?
 - If rep invariant says *no duplicates*
 - **add()** is incorrect
 - Otherwise
 - **remove()** is incorrect



Rep Invariant (partly)

- Sentinel nodes
 - tail reachable from head
- Sorted
- No duplicates



Abstraction Map

- $S(\text{head}) =$
 - $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$
 - $\}$



Sequential List Based Set

Add()




Remove()




Sequential List Based Set


Add()




Remove()



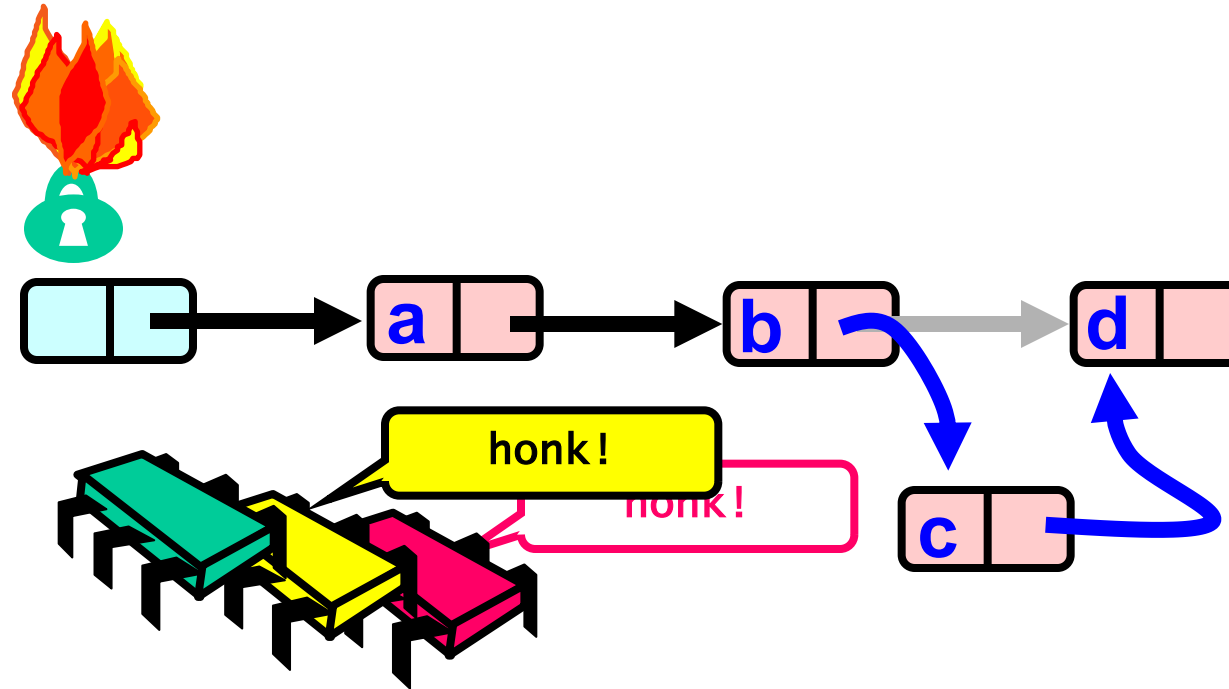
Coarse-Grained Locking



Coarse-Grained Locking



Coarse-Grained Locking



Simple but hotspot + bottleneck

Coarse-Grained Locking

- Easy, same as synchronized methods
 - “One lock to rule them all ...”



Coarse-Grained Locking

- Easy, same as synchronized methods
 - “One lock to rule them all ...”
- Simple, clearly correct
 - Deserves respect!
- Works poorly with contention
 - Queue locks help
 - But bottleneck still an issue



Fine-grained Locking

- Requires **careful** thought
 - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”





Fine-grained Locking


- Requires **careful** thought
 - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
- Split object into pieces
 - Each piece has own lock
 - Methods that work on disjoint pieces need not exclude each other




Hand-over-Hand locking




Art of Multiprocessor Programming




Hand-over-Hand locking




Hand-over-Hand locking





Hand-over-Hand locking




Hand-over-Hand locking




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes



Concurrent Removes




Uh, Oh





Uh, Oh

Bad news, c not removed



Problem

- To delete node **c**
 - Swing node **b**'s next field to **d**
- Problem is,
 - Someone deleting **b** concurrently could direct a pointer to **c**

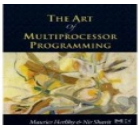




Insight


- If a node is locked
 - No one can delete node's *successor*
- If a thread locks
 - Node to be deleted
 - And its predecessor
 - Then it works




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node



Removing a Node



Removing a Node



Remove method

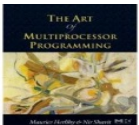
```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```



Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

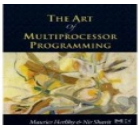
Key used to order node



Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        currNode.unlock();  
        predNode.unlock();  
    }  
}
```

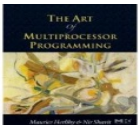
Predecessor and current nodes



Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Make sure
locks released**



Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

Everything else


Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


lock pred == head



Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


Lock current



Remove method


```
try {
    pred = this.head;
    pred.lock();
    curr = pred.next;
    curr.lock();
    ...
} finally { ... }
```

Traversing list



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



Remove: searching


```
while (curr.key <= key) {
```

```
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }
```

```
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

```
return false;
```

Search key range



Remove: searching

```
while (curr.key <= key) {
```

```
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }
```

```
    pred.unlock();
```

```
    pred = curr;
```


```
    curr = curr.next;
```

```
    curr.lock();
```

```
}
```

```
return false;
```


At start of each loop:
curr and pred locked



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item)  
        pred.next = curr.next;  
    return true;  
}  
pred.unlock();  
pred = curr;  
curr = curr.next;  
curr.lock();  
}
```


If item found, remove node



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

If node found, remove it




Remove: searching

Unlock predecessor

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

pred.unlock();




Remove: searching

Only one node locked!

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```


pred.unlock();



Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

demote current



Remove: searching


```
while (curr.key <= key) {  
    Find and lock new current  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

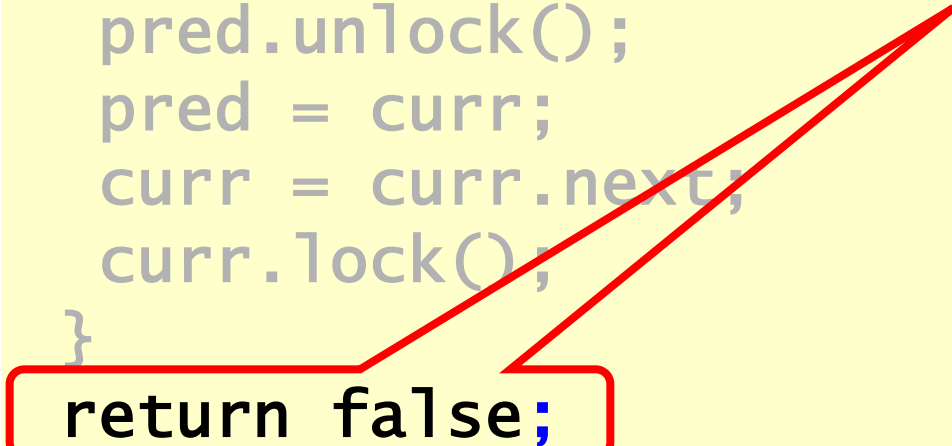
Lock invariant restored



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

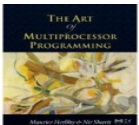
Otherwise, not present



Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

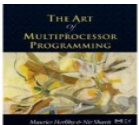
- pred reachable from head
- curr is pred.next
- So curr.item is in the set



Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point if
item is present**



Why remove() is linearizable


```
while (curr.key <= key) {
```

```
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }
```

```
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();
```

```
}  
return false;
```

Node locked, so no other thread
can remove it




Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

return false;

Item not present




Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

return false;

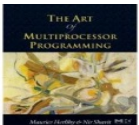
- **pred** reachable from head
- **curr** is pred.next
- **pred.key < key**
- **key < curr.key**



Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Linearization point



Adding Nodes

- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted
 - (Is successor lock actually required?)



Same Abstraction Map

- $S(\text{head}) =$
 - $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$
 - $\}$



Rep Invariant

- Easy to check that
 - tail always reachable from head
 - Nodes sorted, no duplicates



Drawbacks

- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Long chain of acquire/release
 - Inefficient

