



Introduction




Hyungsoo Jung




Multicore Programming: Course Overview


- Fundamentals
 - Models, algorithms, impossibility
- Real-world programming
 - Architectures
 - Techniques



Sequential Computation



Concurrent Computation



Asynchrony




Sudden unpredictable delays

- Cache misses (*short*)
- Page faults (*long*)
- Scheduling quantum used up (*really long*)


Model Summary

- Multiple *threads*
 - Sometimes called *processes*
- Single shared *memory*
- *Objects* live in memory
- Unpredictable asynchronous delays




Road Map

- We are going to focus on principles first, then practice
 - Start with idealized models
 - Look at simplistic problems
 - Emphasize correctness over pragmatism
 - “Correctness may be theoretical, but incorrectness has practical impact”




Concurrency Jargon

- Hardware
 - Processors
- Software
 - Threads, processes
- Sometimes OK to confuse them, sometimes not.




Parallel Primality Testing


- Challenge
 - Print primes from 1 to 10^{10}
- Given
 - Ten-processor multiprocessor
 - One thread per processor
- Goal
 - Get ten-fold speedup (or close)



Load Balancing




- Split the work evenly
- Each thread tests range of 10^9




Procedure for Thread i

```
void primePrint {  
    int i = ThreadID.get(); // IDs in {0..9}  
    for (j = i*109+1, j<(i+1)*109; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```



Issues


- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict




Issues

- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict
- Need *dynamic* load balancing

rejected




Shared Counter



Procedure for Thread *i*

```
int counter = new Counter(1);

void primePrint {
    long j = 0;
    while (j < 1010) {
        j = counter.getAndIncrement();
        if (isPrime(j))
            print(j);
    }
}
```




Procedure for Thread *i*

```
Counter counter = new Counter(1);
```

```
void primePrint {  
    long j = 0;  
    while (j < 1010) {  
        j = counter.getAndIncrement();  
        if (isPrime(j))  
            print(j);  
    }  
}
```



**Shared counter
object**



Where Things Reside

```
void primePrint {  
    int i = ThreadID.get();  
    // IDs in {0..9}  
    for (j = i*10+1;  
         j<(i+1)*10; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```

code



Procedure for Thread *i*

```
Counter counter = new Counter(1);
```

```
void primePrint {
```

```
    long j = 0;
```

```
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```


```
        if (isPrime(j))
```

```
            print(j);
```

```
    }
```

```
}
```

**Stop when every
value taken**



Procedure for Thread *i*

```
Counter counter = new Counter(1);
```

```
void primePrint {
```

```
    long j = 0;
```

```
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```


```
        if (isPrime(j))
```

```
            print(j);
```

```
    }
```


```
}
```

**Increment & return
each new value**



Counter Implementation


```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```



Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

**OK for single thread,
not for concurrent threads**



What It Means


```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```



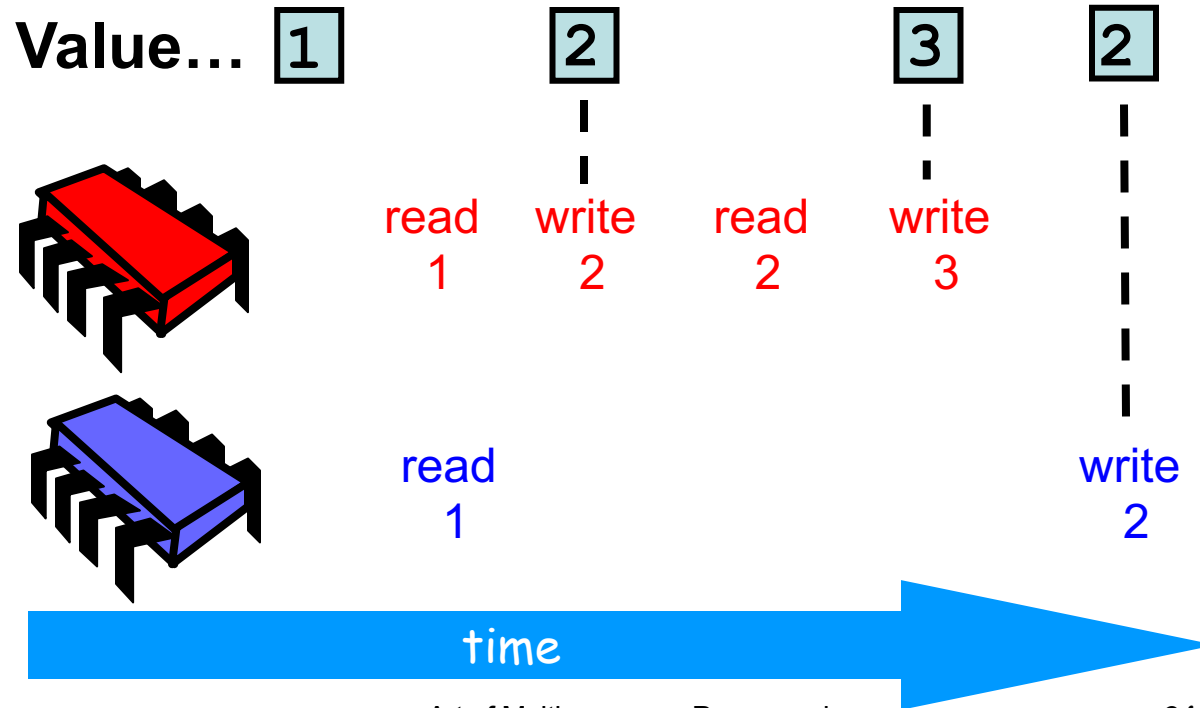
What It Means

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

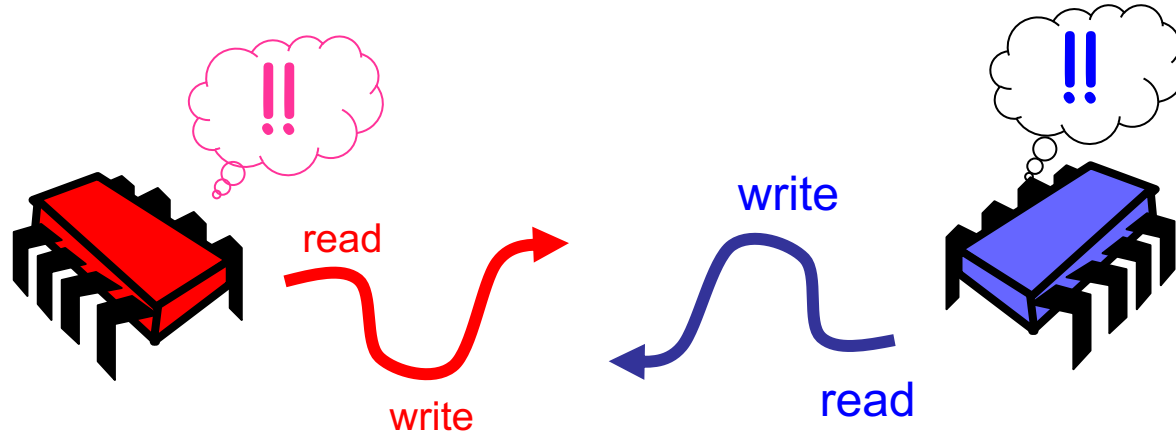
temp = value;
value = temp + 1;
return temp;



Not so good...




Is this problem inherent?



If we could only glue reads and writes together...

Challenge


```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```



Challenge


```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

**Make these steps
atomic (indivisible)**




Hardware Solution

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```




**ReadModifyWrite()
instruction**



An Aside: Java™


```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```



An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```


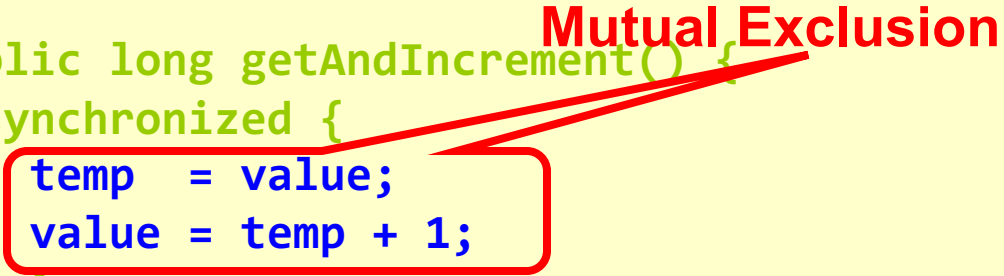
Synchronized block




An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```

Mutual Exclusion



Mutual Exclusion or “Alice & Bob share a pond”



Alice has a pet



Bob has a pet




The Problem




Formalizing the Problem

- Two types of formal properties in asynchronous computation:
- Safety Properties
 - Nothing bad happens ever
- Liveness Properties
 - Something good happens eventually




Formalizing our Problem

- Mutual Exclusion
 - Both pets never in pond simultaneously
 - This is a **safety** property
- No Deadlock
 - if only one wants in, it gets in
 - if both want in, one gets in.
 - This is a **liveness** property




Simple Protocol

- Idea
 - Just look at the pond
- Gotcha
 - Not atomic
 - Trees obscure the view




Interpretation

- Threads can't “see” what other threads are doing
- Explicit communication required for coordination




Cell Phone Protocol

- Idea
 - Bob calls Alice (or vice-versa)
- Gotcha
 - Bob takes shower
 - Alice recharges battery
 - Bob out shopping for pet food ...




Interpretation


- Message-passing doesn't work
- Recipient might not be
 - Listening
 - There at all
- Communication must be
 - Persistent (like writing)
 - Not transient (like speaking)




Flag Protocol



Alice's Protocol (sort of)




Bob's Protocol (sort of)




Alice's Protocol

- Raise flag
- Wait until Bob's flag is down
- Unleash pet
- Lower flag when pet returns




Bob's Protocol

- Raise flag
- Wait until Alice's flag is down
- Unleash pet
- Lower flag when pet returns



Bob's Protocol (2nd try)


- Raise flag
- While Alice's flag is up
 - Lower flag
 - Wait for Alice's flag to go down
 - Raise flag
- Unleash pet
- Lower flag when pet returns



Bob's Protocol


- Raise flag
- While Alice's flag is up
 - Lower flag
 - Wait for Alice's flag to go down
 - Raise flag
- Unleash pet
- Lower flag when pet returns

**Bob defers
to Alice**




The Flag Principle

- Raise the flag
- Look at other's flag
- Flag Principle:
 - If each raises and looks, then
 - Last to look must see both flags up




Proof of Mutual Exclusion


- Assume both pets in pond
 - Derive a contradiction
 - By reasoning **backwards**
- Consider the last time Alice and Bob each looked before letting the pets in
- Without loss of generality assume Alice was the last to look...



Proof




Alice must have seen Bob's Flag. A Contradiction




Proof of No Deadlock

- If only one pet wants in, it gets in.



Proof of No Deadlock

- If only one pet wants in, it gets in.
- Deadlock requires both continually trying to get in.



Proof of No Deadlock


- If only one pet wants in, it gets in.
- Deadlock requires both continually trying to get in.
- If Bob sees Alice's flag, he gives her priority (a gentleman...)

QED




Remarks

- Protocol is *unfair*
 - Bob's pet might never get in
- Protocol uses *waiting*
 - If Bob is eaten by his pet, Alice's pet might never get in




Moral of Story

- Mutual Exclusion cannot be solved by
 - transient communication (cell phones)
 - interrupts (cans)
- It can be solved by
 - one-bit shared variables
 - that can be read or written




The Fable Continues

- Alice and Bob fall in love & marry




The Fable Continues

- Alice and Bob fall in love & marry
- Then they fall out of love & divorce
 - She gets the pets
 - He has to feed them




The Fable Continues

- Alice and Bob fall in love & marry
- Then they fall out of love & divorce
 - She gets the pets
 - He has to feed them
- Leading to a new coordination problem: Producer-Consumer



Bob Puts Food in the Pond




Alice releases her pets to Feed




Producer/Consumer

- Alice and Bob can't meet
 - Each has restraining order on other
 - So he puts food in the pond
 - And later, she releases the pets
- Avoid
 - Releasing pets when there's no food
 - Putting out food if uneaten food remains




Producer/Consumer


- Need a mechanism so that
 - Bob lets Alice know when food has been put out
 - Alice lets Bob know when to put out more food




Surprise Solution





Bob puts food in Pond




Bob knocks over Can



Alice Releases Pets




Alice Resets Can when Pets are Fed



Pseudocode

```
while (true) {  
    while (can.isUp()){};  
    pet.release();  
    pet.recapture();  
    can.reset();  
}
```

Alice's code




Pseudocode

```
while (true) {  
    while (can.isUp()){};  
    pet.release();  
    pet.recapture();  
    can.reset();  
}
```

Alice's code


Bob's code

```
while (true) {  
    while (can.isDown()){};  
    pond.stockWithFood();  
    can.knockOver();  
}
```




Correctness

- Mutual Exclusion
 - Pets and Bob never together in pond




Correctness

- Mutual Exclusion
 - Pets and Bob never together in pond
- No Starvation
 - if Bob always willing to feed, and pets always famished, then pets eat infinitely often.




Correctness

- **Mutual Exclusion** — safety
 - Pets and Bob never together in pond
- **No Starvation** — liveness
 - if Bob always willing to feed, and pets always famished, then pets eat infinitely often.
- **Producer/Consumer** — safety
 - The pets never enter pond unless there is food, and Bob never provides food if there is unconsumed food.




Could Also Solve Using Flags




Waiting

- Both solutions use waiting
 - `while(mumble){}`
- In some cases waiting is ***problematic***
 - If one participant is delayed
 - So is everyone else
 - But delays are common & unpredictable




The Fable drags on ...

- Bob and Alice still have issues




The Fable drags on ...

- Bob and Alice still have issues
- So they need to communicate





The Fable drags on ...


- Bob and Alice still have issues
- So they need to communicate
- They agree to use billboards ...




Billboards are Large





Write One Letter at a Time ...





To post a message



Let's send another message




Uh-Oh




Amdahl's Law

$$\text{Speedup} = \frac{\text{1-thread execution time}}{\text{\textit{n}-thread execution time}}$$



Amdahl's Law

$$\text{Speedup} = \frac{1}{1 + \frac{p}{n}}$$




Amdahl's Law

Speedup=

$$\frac{1}{1 + \frac{p}{n}}$$

Parallel fraction




Amdahl's Law

Sequential fraction

Speedup=

$$\frac{1}{1 + \frac{p}{n}}$$

Parallel fraction



Amdahl's Law


Sequential fraction

Speedup=

Parallel fraction


$$\frac{1}{1 + \frac{p}{n}}$$

Number of threads



Example


- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?



Example


- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 2.17 = \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$



Example


- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?



Example


- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 3.57 = \frac{1}{1 - 0.8 + \frac{0.8}{10}}$$



Example


- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?



Example


- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 5.26 = \frac{1}{1 - 0.9 + \frac{0.9}{10}}$$



Example


- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?



Example


- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 9.17 = \frac{1}{1 - 0.99 + \frac{0.99}{10}}$$






Back to Real-World Multicore Scaling



Speedup




User code



Multicore



**Not reducing
sequential % of code**



Why?

Amdahl's Law:


$$\text{Speedup} = 1/(\text{ParallelPart}/N + \text{SequentialPart})$$

Pay for N = 8 cores

SequentialPart = 25%

As num cores grows the effect of 25%
becomes more accute

2.3/4, 2.9/8, 3.4/16, 3.7/32....



Shared Data Structures

