
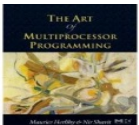


# Mutual Exclusion



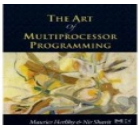
Hyungsoo Jung



# Mutual Exclusion



- We will clarify our understanding of mutual exclusion
- We will also show you how to reason about various properties in an asynchronous concurrent setting




# Mutual Exclusion



In his 1965 paper E. W. Dijkstra wrote:

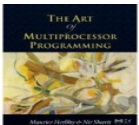
"Given in this paper is a solution to a problem which, to the knowledge of the author, has been an open question since at least 1962, irrespective of the solvability. [...] Although the setting of the problem might seem somewhat academic at first, the author trusts that anyone familiar with the logical problems that arise in computer coupling will appreciate the significance of the fact that this problem indeed can be solved."



# Mutual Exclusion



- Formal problem definitions
- Solutions for 2 threads
- Solutions for  $n$  threads
- Fair solutions
- Inherent costs




# Warning

- You will never use these protocols
  - Get over it
- You are advised to understand them
  - The same issues show up everywhere
  - Except hidden and more complex




# Time

- “Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external.” (I. Newton, 1689)
- “Time is, like, Nature’s way of making sure that everything doesn’t happen all at once.” (Anonymous, circa 1968)




# Events

- An *event*  $a_0$  of thread A is
  - Instantaneous
  - No simultaneous events (break ties)



# Threads

- A *thread* A is (formally) a sequence  $a_0, a_1, \dots$  of events
  - “Trace” model
  - Notation:  $a_0 \rightarrow a_1$  indicates order






# Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...



# Threads are State Machines

Events are  
transitions



# States

- Thread State
  - Program counter
  - Local variables
- System state
  - Object fields (shared variables)
  - Union of thread states



# Concurrency

- Thread A



# Concurrency

- Thread A




- Thread B




# Interleavings

- Events of two or more threads
  - Interleaved
  - Not necessarily independent (why?)




# Intervals

- An *interval*  $A_0 = (a_0, a_1)$  is
  - Time between events  $a_0$  and  $a_1$




# Intervals may Overlap






# Intervals may be Disjoint




# Precedence

Interval  $A_0$  precedes interval  $B_0$




# Precedence




- Notation:  $A_0 \rightarrow B_0$
- Formally,
  - End event of  $A_0$  before start event of  $B_0$
  - Also called “happens before” or “precedes”

# Precedence Ordering



- Remark:  $A_0 \rightarrow B_0$  is just like saying
  - 1066 AD  $\rightarrow$  1492 AD,
  - Middle Ages  $\rightarrow$  Renaissance,
- Oh wait,
  - what about this week vs this month?

# Precedence Ordering




- Never true that  $A \rightarrow A$
- If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- If  $A \rightarrow B$  &  $B \rightarrow C$  then  $A \rightarrow C$
- Funny thing:  $A \rightarrow B$  &  $B \rightarrow A$  might both be false!

# Partial Orders

(review)

- Irreflexive:
  - Never true that  $A \rightarrow A$
- Antisymmetric:
  - If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- Transitive:
  - If  $A \rightarrow B$  &  $B \rightarrow C$  then  $A \rightarrow C$



# Total Orders

(review)

- Also
  - Irreflexive
  - Antisymmetric
  - Transitive
- Except that for every distinct  $A, B$ ,
  - Either  $A \rightarrow B$  or  $B \rightarrow A$



# Repeated Events


```
while (mumble) {  
    a0; a1;  
}
```

*k*-th occurrence of  
event  $a_0$

$a_0^k$

*k*-th occurrence of  
interval  $A_0 = (a_0, a_1)$

$A_0^k$

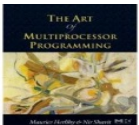




# Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Make these steps  
*indivisible* using locks



# Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```



# Locks (Mutual Exclusion)


```
public interface Lock {
```

```
    public void lock();
```

**acquire lock**

```
    public void unlock();
```

```
}
```



# Locks (Mutual Exclusion)

```
public interface Lock {
```


```
    public void lock();
```

**acquire lock**

```
    public void unlock();
```

**release lock**

```
}
```



# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```



# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

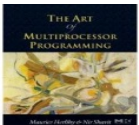
**acquire Lock**



# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

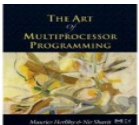
Release lock  
(no matter what)



# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

critical section







# Mutual Exclusion







- Let  $CS_i^k \iff$  be thread  $i$ 's  $k$ -th critical section execution





# Mutual Exclusion

- Let  $CS_i^k$   be thread i's k-th critical section execution
- And  $CS_j^m$   be thread j's m-th critical section execution

# Mutual Exclusion

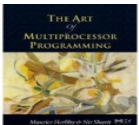
- Let  $CS_i^k$   be thread i's k-th critical section execution
- And  $CS_j^m$   be j's m-th execution
- Then either
  -   or  

# Mutual Exclusion

- Let  $CS_i^k$   be thread i's k-th critical section execution
- And  $CS_j^m$   be j's m-th execution
- Then either

—   or  

$CS_i^k \rightarrow CS_j^m$




# Mutual Exclusion

- Let  $CS_i^k$   $\longleftrightarrow$  be thread i's k-th critical section execution
- And  $CS_j^m$   $\longleftrightarrow$  be j's m-th execution
- Then either

—  $\longleftrightarrow$   $\longleftrightarrow$  or  $\longleftrightarrow$   $\longleftrightarrow$

$CS_i^k \rightarrow CS_j^m$


$CS_j^m \rightarrow CS_i^k$



# Deadlock-Free



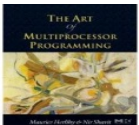
- If some thread calls **lock()**
  - And never returns
  - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
  - Even if individuals starve



# Starvation-Free



- If some thread calls `lock()`
  - It will eventually return
- Individual threads make progress



# Two-Thread vs $n$ -Thread Solutions

- 2-thread solutions first
  - Illustrate most basic ideas
  - Fits on one slide
- Then  $n$ -thread solutions





# Two-Thread Conventions

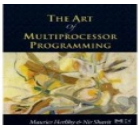
```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```



# Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

Henceforth: *i* is current  
thread, *j* is other thread



# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

**Each thread has flag**

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

**Set my flag**

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

**Wait for other flag to  
become false**

# LockOne Satisfies Mutual Exclusion


- Assume  $CS_A^j$  overlaps  $CS_B^k$
- Consider each thread's last (j-th and k-th) read and write in the `lock()` method before entering
- Derive a contradiction



# From the Code

- **write<sub>A</sub>(flag[A]=true) →  
read<sub>A</sub>(flag[B]==false) → CS<sub>A</sub>**
- **write<sub>B</sub>(flag[B]=true) →  
read<sub>B</sub>(flag[A]==false) → CS<sub>B</sub>**

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```





# From the Assumption

- **$\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$**
- **$\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$**




# Combining

- Assumptions:
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

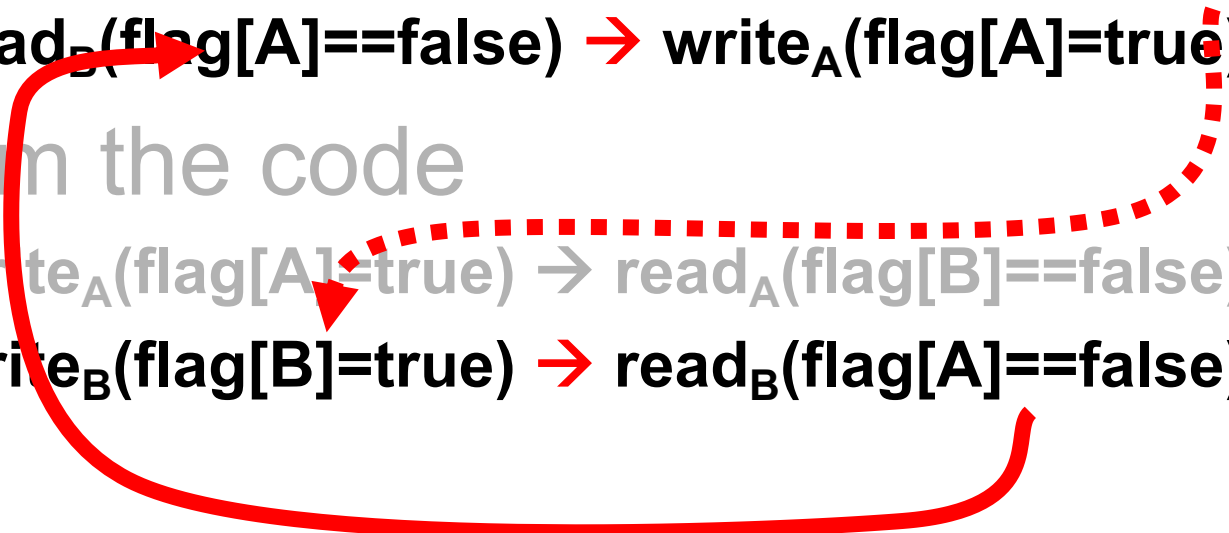



# Combining

- Assumptions:
  - **read<sub>A</sub>(flag[B]==false) → write<sub>B</sub>(flag[B]=true)**
  - read<sub>B</sub>(flag[A]==false) → write<sub>A</sub>(flag[A]=true)
- From the code
  - write<sub>A</sub>(flag[A]=true) → read<sub>A</sub>(flag[B]==false)
  - **write<sub>B</sub>(flag[B]=true) → read<sub>B</sub>(flag[A]==false)**



# Combining

- Assumptions:
    - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
    - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
  - From the code
    - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
    - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 




# Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$




# Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$




# Combining

- Assumptions:


- `readA(flag[B]==false) → writeB(flag[B]=true)`
- `readB(flag[A]==false) → writeA(flag[A]=true)`

- From the code

- `writeA(flag[A]=true) → readA(flag[B]==false)`
- `writeB(flag[B]=true) → readB(flag[A]==false)`



# Cycle!






# Deadlock Freedom

- LockOne Fails deadlock-freedom
  - Concurrent execution can deadlock

```
flag[i] = true;    flag[j] = true;
while (flag[j]){   while (flag[i]){
```

- Sequential executions OK



# LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

# LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Let other go  
first**

# LockTwo

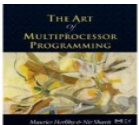
```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Wait for permission**

# LockTwo

```
public class Lock2 implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
    public void unlock() {}  
}
```

**Nothing to do**



# LockTwo Claims


- Satisfies mutual exclusion

- If thread **i** in CS
- Then **victim == j**
- Cannot be both 0 and 1

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {}  
}
```

- Not deadlock free

- Sequential execution deadlocks
- Concurrent execution does not



# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```



# Peterson's Algorithm

**Announce I'm  
interested**

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```




# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

**Announce I'm  
interested**

**Defer to other**



# Peterson's Algorithm

```
public void lock() {
```

```
    flag[i] = true;
```

```
    victim = i;
```

```
    while (flag[j] && victim == i) {};
```

```
}
```

```
public void unlock() {
```


```
    flag[i] = false;
```

```
}
```

**Announce I'm  
interested**

**Defer to other**

**Wait while other  
interested & I'm  
the victim**



# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

**Announce I'm interested**

**Defer to other**

**Wait while other interested & I'm the victim**

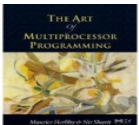
**No longer interested**

# Mutual Exclusion

(1)  $\text{write}_B(\text{Flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

From the Code



# Also from the Code


(2)  $\text{write}_A(\text{victim}=\text{A}) \rightarrow \text{read}_A(\text{flag}[\text{B}])$   
 $\rightarrow \text{read}_A(\text{victim})$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

# Assumption

(3)  $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

W.L.O.G. assume **A** is the last  
thread to write **victim**



# Combining Observations

- (1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$
- (3)  $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$
- (2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$




# Combining Observations

(1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3)  $\text{write}_B(\text{victim}=B) \rightarrow$

(2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$





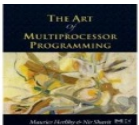
# Combining Observations

(1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3)  $\text{write}_B(\text{victim}=B) \rightarrow$

(2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$

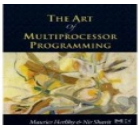
A **read**  $\text{flag}[B] == \text{true}$  **and**  $\text{victim} == A$ , so it  
could not have entered the CS (**QED**)



# Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};
```

- Thread blocked
  - only at **while** loop
  - only if other's flag is true
  - only if it is the **victim**
- Solo: other's flag is false
- Both: one or the other not the victim



# Starvation Free


- Thread *i* blocked only if *j* repeatedly re-enters so that `flag[j] == true` and `victim == i`
- When *j* re-enters
  - it sets `victim` to *j*.
  - So *i* gets in

```
public void lock() {  
    flag[i] = true;  
    victim  = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

# The Filter Algorithm for $n$ Threads

There are  $n-1$  “waiting rooms” called levels


- At each level
  - At least one enters level
  - At least one blocked if many try
- Only one thread makes it through




# Filter

```
class Filter implements Lock {  
    int[] level; // level[i] for thread i  
    int[] victim; // victim[L] for level L
```

```
    public Filter(int n) {  
        level = new int[n];  
        victim = new int[n];  
        for (int i = 1; i < n; i++) {  
            level[i] = 0;  
        }  
    }  
    ...  
}
```



**Thread 2 at level 4**



# Filter

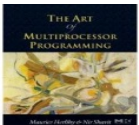
```
class Filter implements Lock {  
    ...  
  
    public void lock(){  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
  
            while (( $\exists$  k != i level[k] >= L) &&  
                    victim[L] == i ) {};  
        }  
    }  
    public void unlock() {  
        level[i] = 0;  
    }  
}
```



# Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i) {};  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

One level at a time



# Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i)  
            }  
        }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```


**Announce  
intention to enter  
level L**



# Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i) {};  
        }  
    }  
    public void release(int i)  
        level[i] = 0;  
}
```

**Give priority to  
anyone but me**



# Filter

**Wait as long as someone else is at same or higher level, and I'm designated victim**

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L) &&  
            victim[L] == i) {};  
    }  
}  
public void release(int i) {  
    level[i] = 0;  
}
```


# Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i) {};  
        }  
    }  
}
```

**Thread enters level L when it completes the loop**

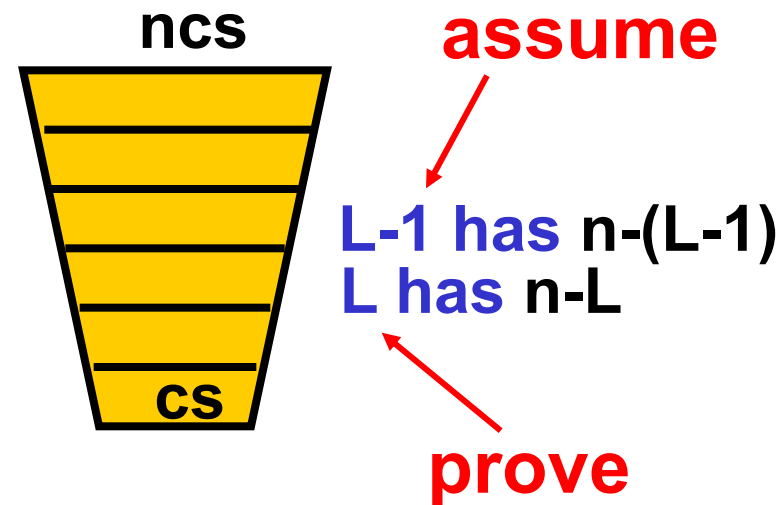
# Claim

- Start at level  $L=0$
- At most  $n-L$  threads enter level  $L$
- Mutual exclusion at level  $L=n-1$




# Induction Hypothesis


- No more than  $n-(L-1)$  at level  $L-1$
- Induction step: by contradiction
- Assume all at level  $L-1$  enter level  $L$
- A last to write `victim[L]`
- B is any other thread at level  $L$



# Proof Structure



Show that **A** must have seen **B** in level **L** and since **victim[L] == A** could not have entered




# Just Like Peterson

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L)  
            && victim[L] == i) {};  
    }  
}
```

From the Code



# From the Code

(2)  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$   
 $\rightarrow \text{read}_A(\text{victim}[L])$

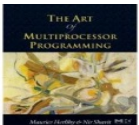
```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
                && victim[L] == i) {};  
    }  
}
```



# By Assumption

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

By assumption, A is the last thread  
to write **victim[L]**



# Combining Observations

- (1)  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$
- (3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2)  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$   
 $\rightarrow \text{read}_A(\text{victim}[L])$




# Combining Observations

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow$

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2)  $\rightarrow \text{read}_A(\text{level}[B])$

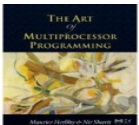
$\rightarrow \text{read}_A(\text{victim}[L])$



# Combining Observations

- (1)  $\text{write}_B(\text{level}[B]=L) \rightarrow$
- (3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2)  $\rightarrow \text{read}_A(\text{level}[B])$
- $\rightarrow \text{read}_A(\text{victim}[L])$

**A read  $\text{level}[B] \geq L$ , and  $\text{victim}[L] = A$ , so it could not have entered level L!**



# No Starvation

- Filter Lock satisfies properties:
  - Just like Peterson Alg at any level
  - So no one starves
- But what about fairness?
  - Threads can be overtaken by others



# Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- If A starts before B, then A enters before B?
- But what does “start” mean?
- Need to adjust definitions ....



# Bounded Waiting

- Divide `lock()` method into 2 parts:
  - Doorway interval:
    - Written  $D_A$
    - always finishes in finite steps
  - Waiting interval:
    - Written  $W_A$
    - may take unbounded steps



# First-Come-First-Served

- For threads A and B:
  - If  $D_A^k \rightarrow D_B^j$ 
    - A's k-th doorway precedes B's j-th doorway
  - Then  $CS_A^k \rightarrow CS_B^j$ 
    - A's k-th critical section precedes B's j-th critical section
    - B cannot overtake A





# Fairness Again

- Filter Lock satisfies properties:
  - No one starves
  - But very weak fairness
    - Can be overtaken **arbitrary** # of times
  - That's pretty lame...



# Bakery Algorithm

- Provides First-Come-First-Served
- How?
  - Take a “number”
  - Wait until lower numbers have been served
- Lexicographic order
  - $(a,i) > (b,j)$ 
    - If  $a > b$ , or  $a = b$  and  $i > j$



# Bakery Algorithm

```
class Bakery implements Lock {  
    boolean[] flag;  
    Label[] label;  
    public Bakery (int n) {  
        flag = new boolean[n];  
        label = new Label[n];  
        for (int i = 0; i < n; i++) {  
            flag[i] = false; label[i] = 0;  
        }  
    }  
    ...  
}
```



# Bakery Algorithm

```
class Bakery implements Lock {
```

```
    boolean[] flag;
```

```
    Label[] label;
```

```
    public Bakery (int n) {
```

```
        flag = new boolean[n];
```

```
        label = new Label[n];
```


```
        for (int i = 0; i < n; i++) {
```

```
            flag[i] = false; label[i] = 0;
```

```
        }
```

```
    }
```

```
    ...
```



# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```



# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**Doorway**

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**I'm interested**

# Bakery Algorithm

**Take increasing  
label (read labels  
in some arbitrary  
order)**

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```



# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```


**Someone is interested**

# Bakery Algorithm

```
class Bakery implements Lock {  
    boolean flag[n];  
    int label[n];  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }
```

**Someone is  
interested ...**

**... whose (label,i) in  
lexicographic order is lower**



# Bakery Algorithm

```
class Bakery implements Lock {  
  
    ...  
  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

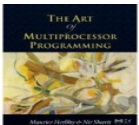


# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

**No longer interested**

**labels are always increasing**



# No Deadlock


- There is always one thread with earliest label
- Ties are impossible (why?)



# First-Come-First-Served

- If  $D_A \rightarrow D_B$  then
  - A's label is smaller
- And:
  - $\text{write}_A(\text{label}[A]) \rightarrow$
  - $\text{read}_B(\text{label}[A]) \rightarrow$
  - $\text{write}_B(\text{label}[B]) \rightarrow \text{read}_B(\text{flag}[A])$
- So B sees
  - smaller label for A
  - locked out while  $\text{flag}[A]$  is true

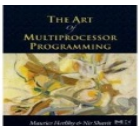
```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) >  
                (label[k], k));  
    }  
}
```



# Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
  - flag[A] is false, or
  - label[A] > label[B]

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1])+1;  
        while (∃k flag[k]  
                && (label[i], i) >  
                (label[k], k));  
    }  
}
```



# Mutual Exclusion

- Labels are strictly increasing so
- B must have seen `flag[A] == false`





# Mutual Exclusion

- Labels are strictly increasing so
- B must have seen  $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$



# Mutual Exclusion

- Labels are strictly increasing so
- B must have seen  $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has an earlier label



# Deep Philosophical Question

- The Bakery Algorithm is
  - Succinct,
  - Elegant, and
  - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read **N** distinct variables



# Shared Memory

- Shared read/write memory locations called **Registers** (historical reasons)
- Come in different flavors
  - Multi-Reader-Single-Writer (**Flag**[])
  - Multi-Reader-Multi-Writer (**Victim**[])
  - Not that interesting: SRMW and SRSW



# Summary of Lecture

- In the 1960's several **incorrect** solutions to starvation-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO **N** thread mutual exclusion using **2N** RW-Registers

