



Foundations of Shared Memory




Hyungsoo Jung



Art of Multiprocessor Programming


Last Lecture

- Defined concurrent objects using linearizability and sequential consistency
- Fact: implemented linearizable objects (Two thread FIFO Queue) in read-write memory without mutual exclusion
- Fact: hardware does not provide linearizable read-write memory



Fundamentals


- What is the weakest form of communication that supports mutual exclusion?
- What is the weakest shared object that allows shared-memory computation?




Alan Turing



- Showed what is and is not computable on a sequential machine.
- Still best model there is.




Turing Computability



- Mathematical model of computation
- What is (and is not) computable
- Efficiency (mostly) irrelevant


Shared-Memory Computability?



- Mathematical model of **concurrent** computation
- What is (and is not) concurrently computable
- Efficiency (mostly) irrelevant


Foundations of Shared Memory

To understand modern
multiprocessors we need to ask some
basic questions ...




Foundations of Shared Memory

To understand modern
What is the weakest useful form of
shared memory?




Foundations of Shared Memory

To understand modern
What is the weakest useful form of
What can it do?




Register*


Holds a
(binary) value




*** A memory location: name is historical**



Register



Register

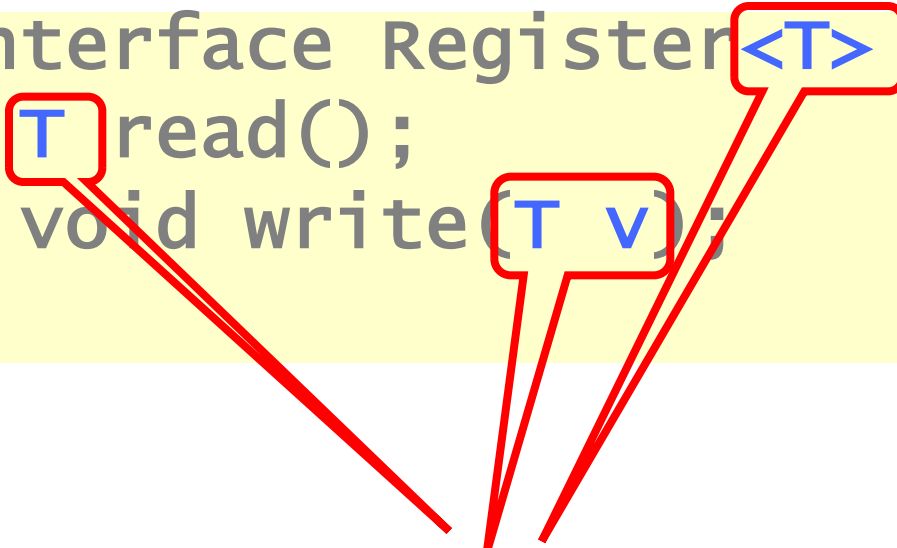


Registers

```
public interface Register<T> {  
    public T read();  
    public void write(T v);  
}
```


Registers

```
public interface Register<T> {  
    public T read();  
    public void write(T v);  
}
```




Type of register
(usually Boolean or *m*-bit Integer)

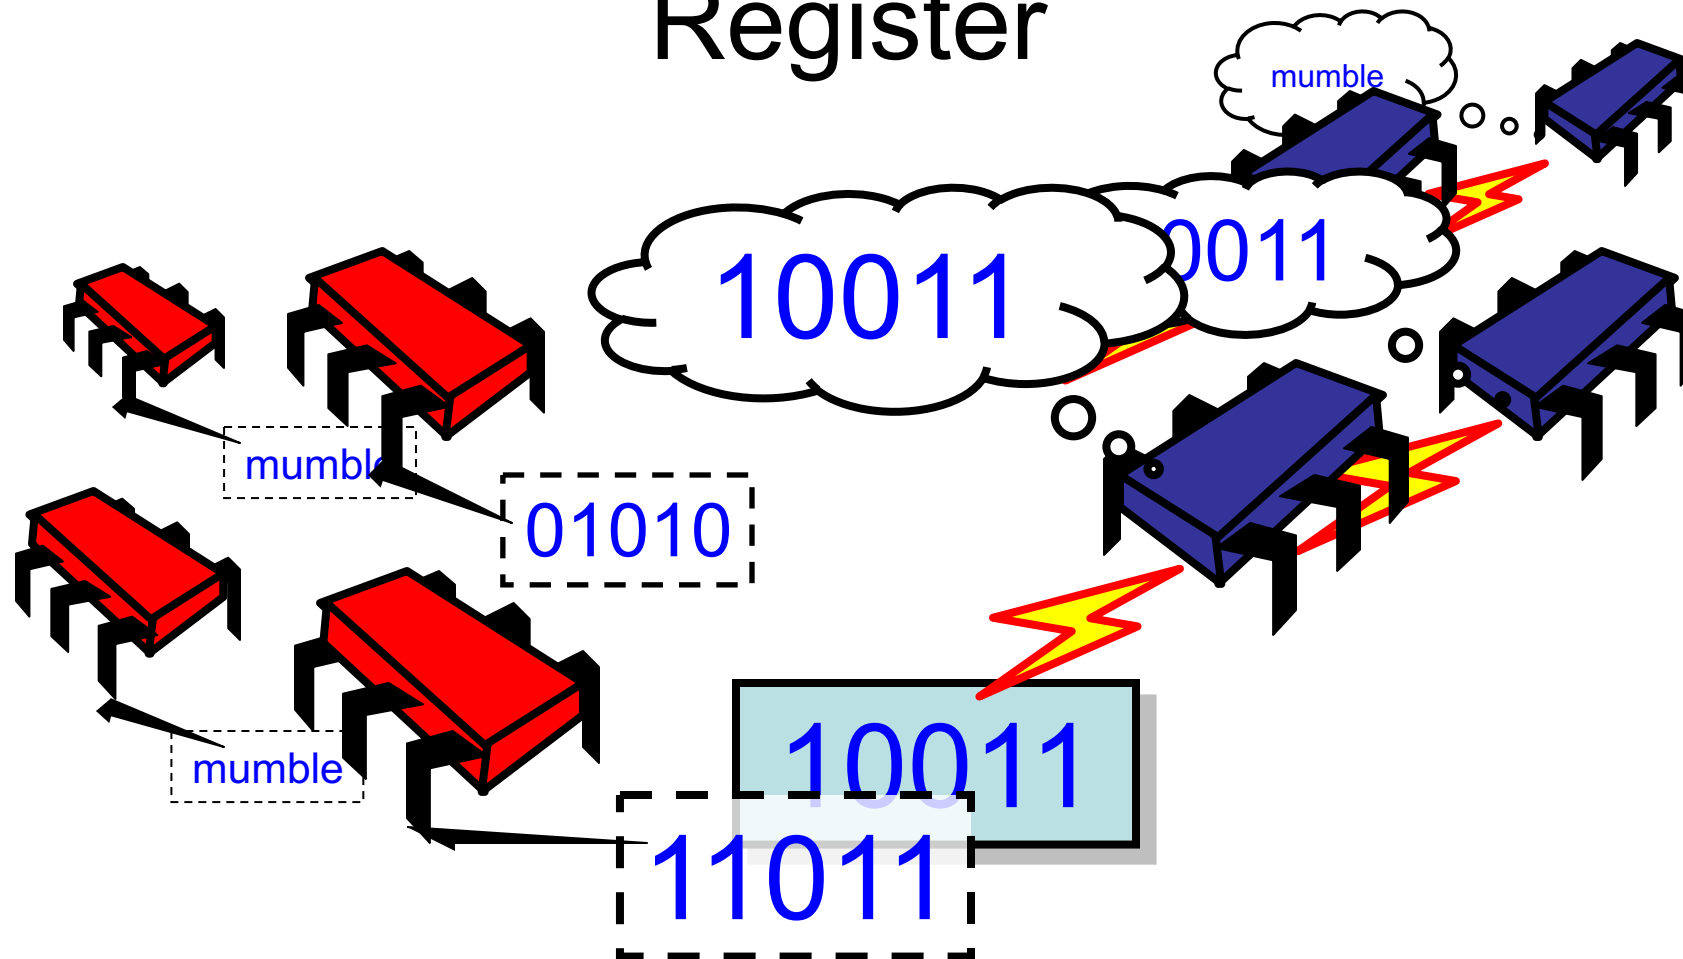
Single-Reader/Single-Writer Register



Multi-Reader/Single-Writer Register




Multi-Reader/Multi-Writer Register






Jargon Watch

- SRSW
 - Single-reader single-writer
- MRSW
 - Multi-reader single-writer
- MRMW
 - Multi-reader multi-writer



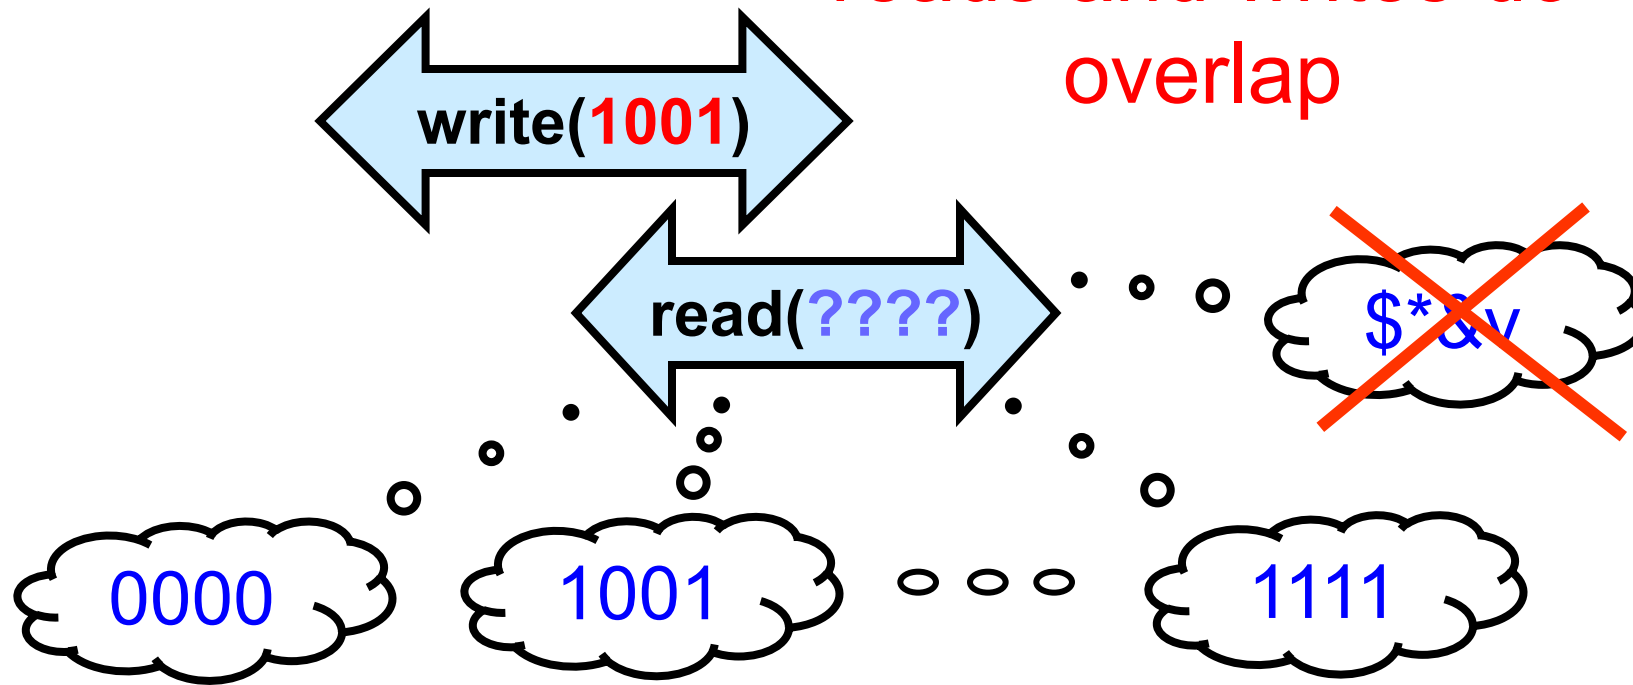
Safe Register

OK if reads
and writes
don't overlap




Safe Register

Some valid value if
reads and writes do
overlap




Regular Register




- Single Writer
- Readers return:
 - Old value if no overlap (safe)
 - Old or one of new values if overlap


Regular or Not?



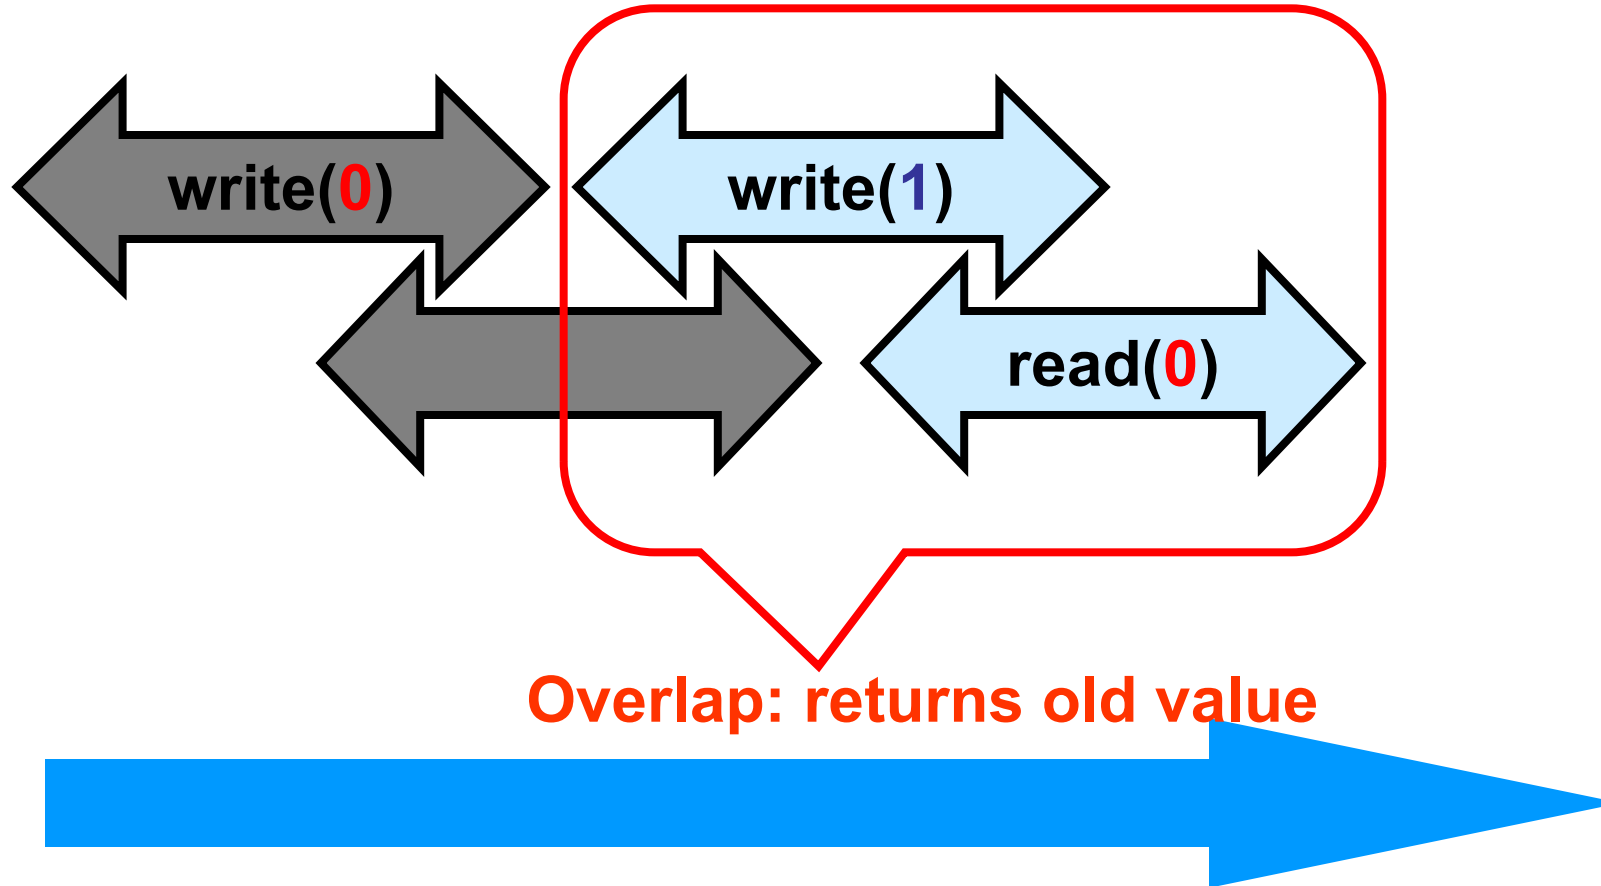
Regular or Not?



Overlap: returns new value




Regular or Not?

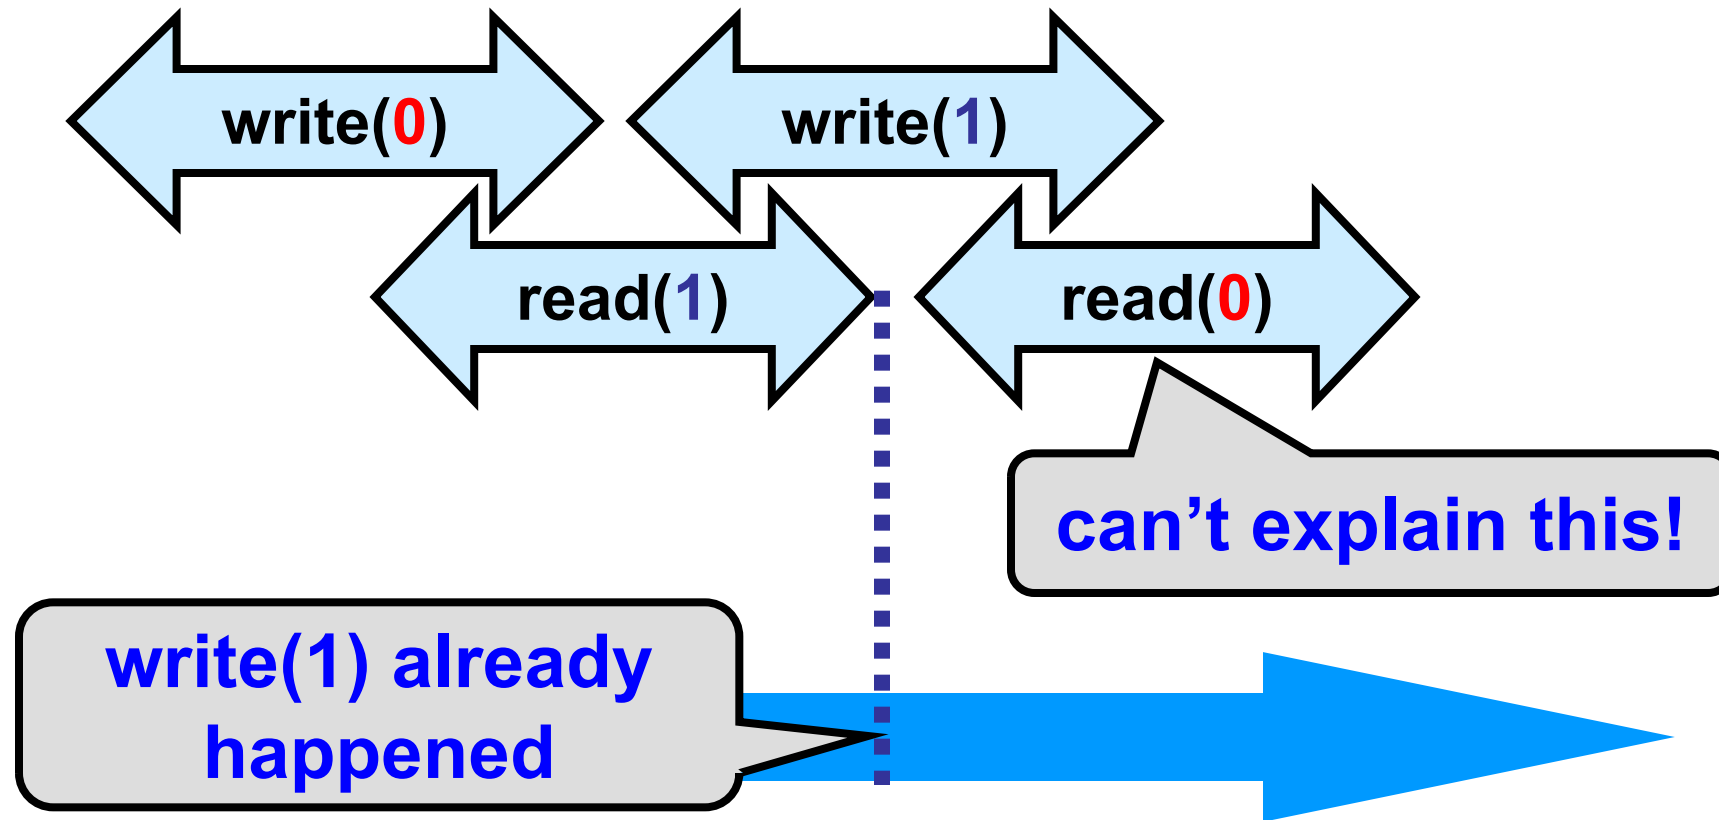


Overlap: returns old value


Regular or Not?



Regular \neq Linearizable

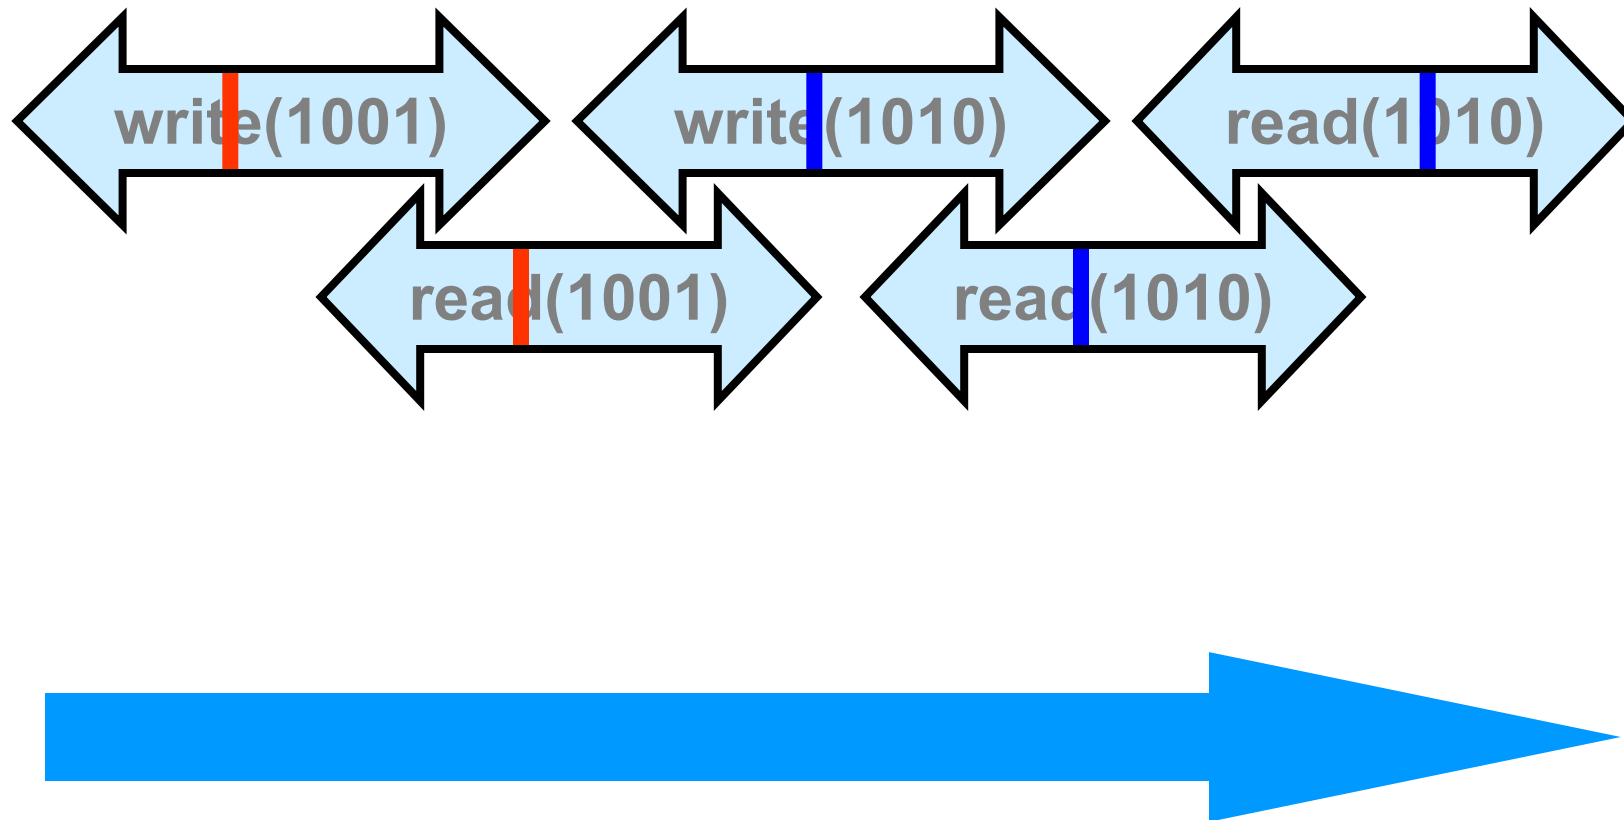


Atomic Register





Linearizable to sequential safe
register


Atomic Register



Register Space




Weakest Register



Weakest Register

Single writer

Single reader



Get correct reading if not during state transition

Results

- From SRSW safe Boolean register


- All the other registers
- Mutual exclusion

Foundations
of the field

- But not everything!


- Consensus hierarchy

The really cool stuff ...



Locking within Registers

- Not interesting to rely on mutual exclusion in register constructions
- We want registers to implement mutual exclusion!
- It's cheating to use mutual exclusion to implement itself!




Definition



An object implementation is ***wait-free*** if every method call completes in a finite number of steps

No mutual exclusion

- Thread could halt in critical section
- Build mutual exclusion from registers




From Safe SRSW Boolean to Atomic Snapshots




Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot



Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot



Register Names

```
public class SafeBoolMRSWRegister  
    implements Register<Boolean> {  
    public boolean read() { ... }  
    public void write(boolean x) { ... }  
}
```

Register Names

```
public class SafeBoolMRSWRegister  
    implements Register<Boolean> {  
    public boolean read() { ... }  
    public void write(boolean x) { ... }  
}
```


property

Register Names

```
public class SafeBooleanMRSWRegister  
    implements Register<Boolean> {  
    public boolean read() { ... }  
    public void write(boolean x) { ... }  
}
```

property

type




Register Names

```
public class SafeBoolMRSWRegister  
    implements Register<Boolean> {  
    public boolean read() { ... }  
    public void write(boolean x) { ... }  
}
```


property

type


how many readers &
writers?




Safe Boolean **MR**SW from Safe Boolean **SR**SW




Safe Boolean MRSW from Safe Boolean SRSW




Safe Boolean MRSW from Safe Boolean SRSW




Safe Boolean MRSW from Safe Boolean SRSW



Safe Boolean MRSW from Safe Boolean SRSW



Safe Boolean MRSW from Safe Boolean SRSW



Safe Boolean MRSW from Safe Boolean SRSW

```
public class SafeBoolMRSWRegister
implements Register<Boolean> {
private SafeBoolSRSWRegister[] r =
    new SafeBoolSRSWRegister[N];
public void write(boolean x) {
    for (int j = 0; j < N; j++)
        r[j].write(x);
}
public boolean read() {
    int i = ThreadID.get();
    return r[i].read();
}}
```


Safe Boolean MRSW from Safe Boolean SRSW


```
public class SafeBoolMRSWRegister  
    implements BooleanRegister {  
    private SafeBoolSRSWRegister[] r =  
        new SafeBoolSRSWRegister[N];  
    public void write(boolean x) {  
        for (int j = 0; j < N; j++)  
            r[j].write(x);  
    }  
    public boolean read() {  
        int i = ThreadID.get();  
        return r[i].read();  
    }  
}
```

**Each thread has own
safe SRSW register**

Safe Boolean MRSW from Safe Boolean SRSW

```
public class SafeBoolMRSWRegister
    implements BooleanRegister {
    private SafeBoolSRSWRegister[] r =
        new SafeBoolSRSWRegister[N];
    public void write(boolean x) {
        for (int j = 0; j < N; j++)
            r[j].write(x);
        }
    public boolean read() {
        int i = ThreadID.get();
        return r[i].read();
    }}
}
```

write method



Safe Boolean MRSW from Safe Boolean SRSW


```
public class SafeBoolMRSWRegister  
    implements BooleanRegister {  
    private SafeBoolSRSWRegister[] r =  
        new SafeBoolSRSWRegister[N];  
    public void write(boolean x) {  
        for (int j = 0; j < N; j++)  
            r[j].write(x);  
    }  
    public boolean read() {  
        int i = ThreadID.get();  
        return r[i].read();  
    }  
}
```

**Write each
thread's register
one at a time**

Safe Boolean MRSW from Safe Boolean SRSW

```
public class SafeBoolMRSWRegister
implements BooleanRegister {
private SafeBoolSRSWRegister[] r =
    new SafeBoolSRSWRegister[N];
public void write(boolean x) {
    for (int j = 0; j < N; j++)
        r[j].write(x);
}
public boolean read() {
    int i = ThreadID.get();
    return r[i].read();
}}
```

read method

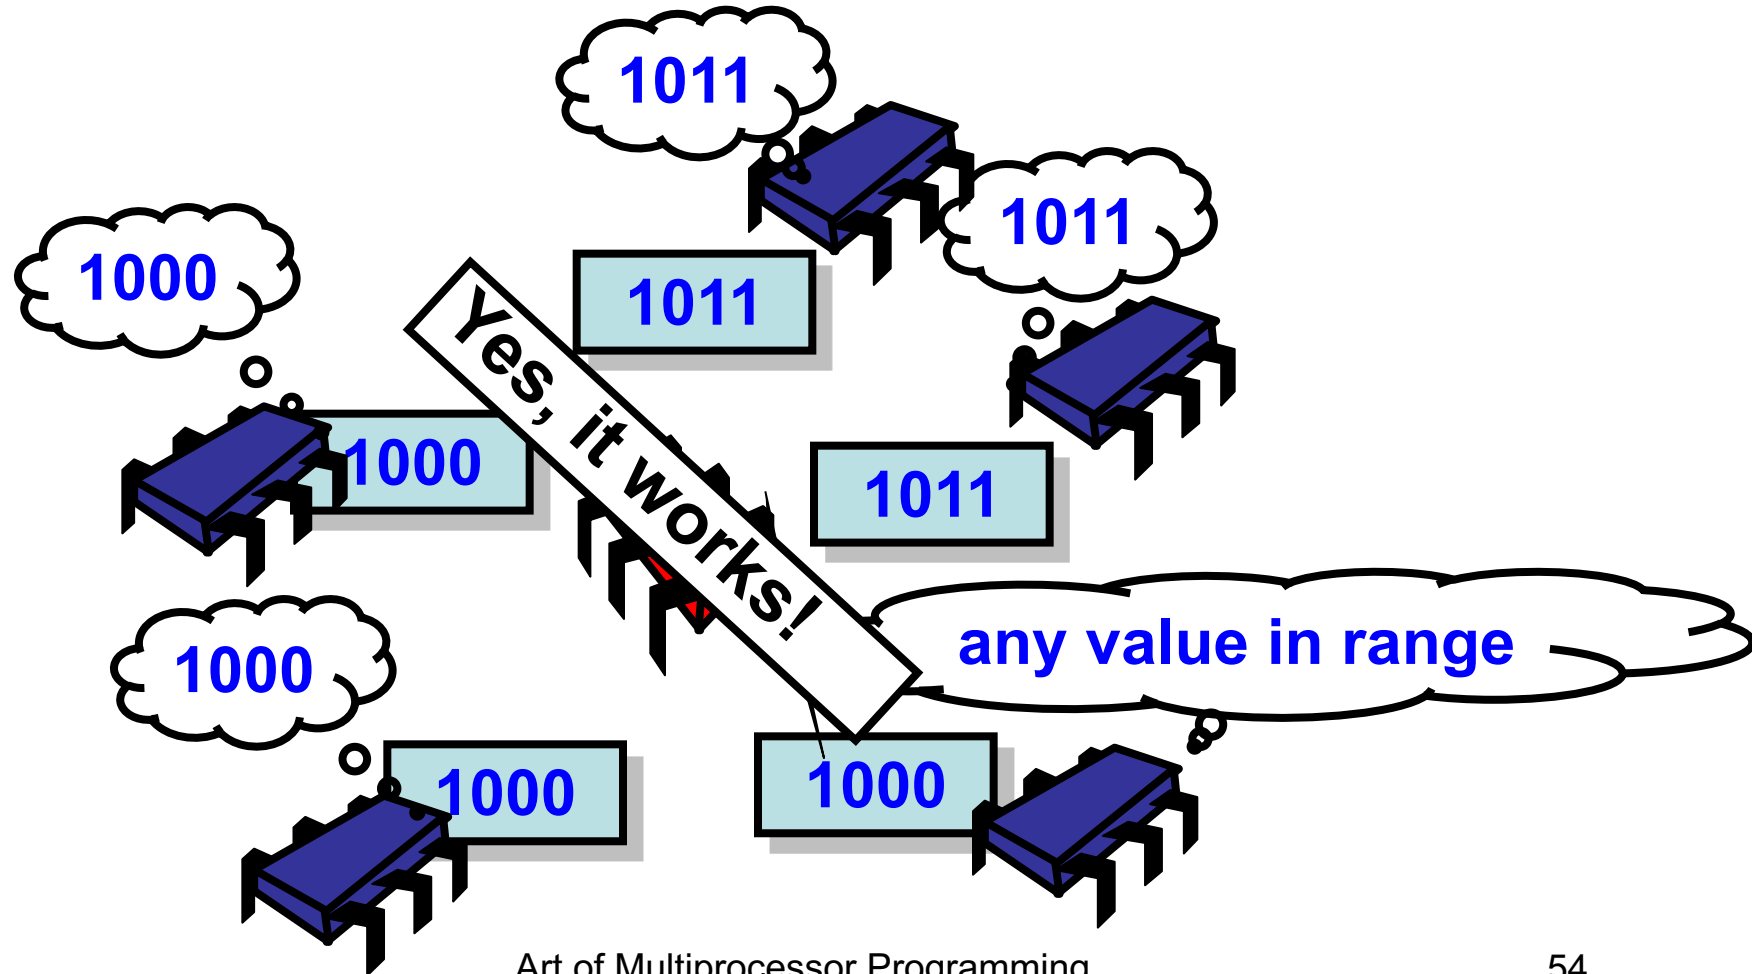


Safe Boolean MRSW from Safe Boolean SRSW

```
public class SafeBoolMRSWRegister
    implements BooleanRegister {
    private SafeBoolSRSWRegister[] r =
        new SafeBoolSRSWRegister[N];
    public void write(boolean x) {
        for (int j = 0; j < N; j++)
            r[j].write(x);
    }
    public boolean read() {
        int i = ThreadID.get();
        return r[i].read();
    }
}
```

**Read my own
register**

Safe Multi-Valued MRSW from Safe Multi-Valued SRSW?




Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot




Questions?




Road Map


- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot




Regular Boolean MRSW from Safe Boolean MRSW



Regular Boolean MRSW from Safe Boolean MRSW




Regular Boolean MRSW from Safe Boolean MRSW



Regular Boolean MRSW from Safe Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    private boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```




Regular Boolean MRSW from Safe Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Last bit this thread wrote

(made-up syntax)



Regular Boolean MRSW from Safe Boolean MRSW


```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Actual value

Regular Boolean MRSW from Safe Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```


**Is new value different
from last value I wrote?**



Regular Boolean MRSW from Safe Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

**If so, change it
(otherwise don't!)**




Regular Boolean MRSW from Safe Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean>{
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

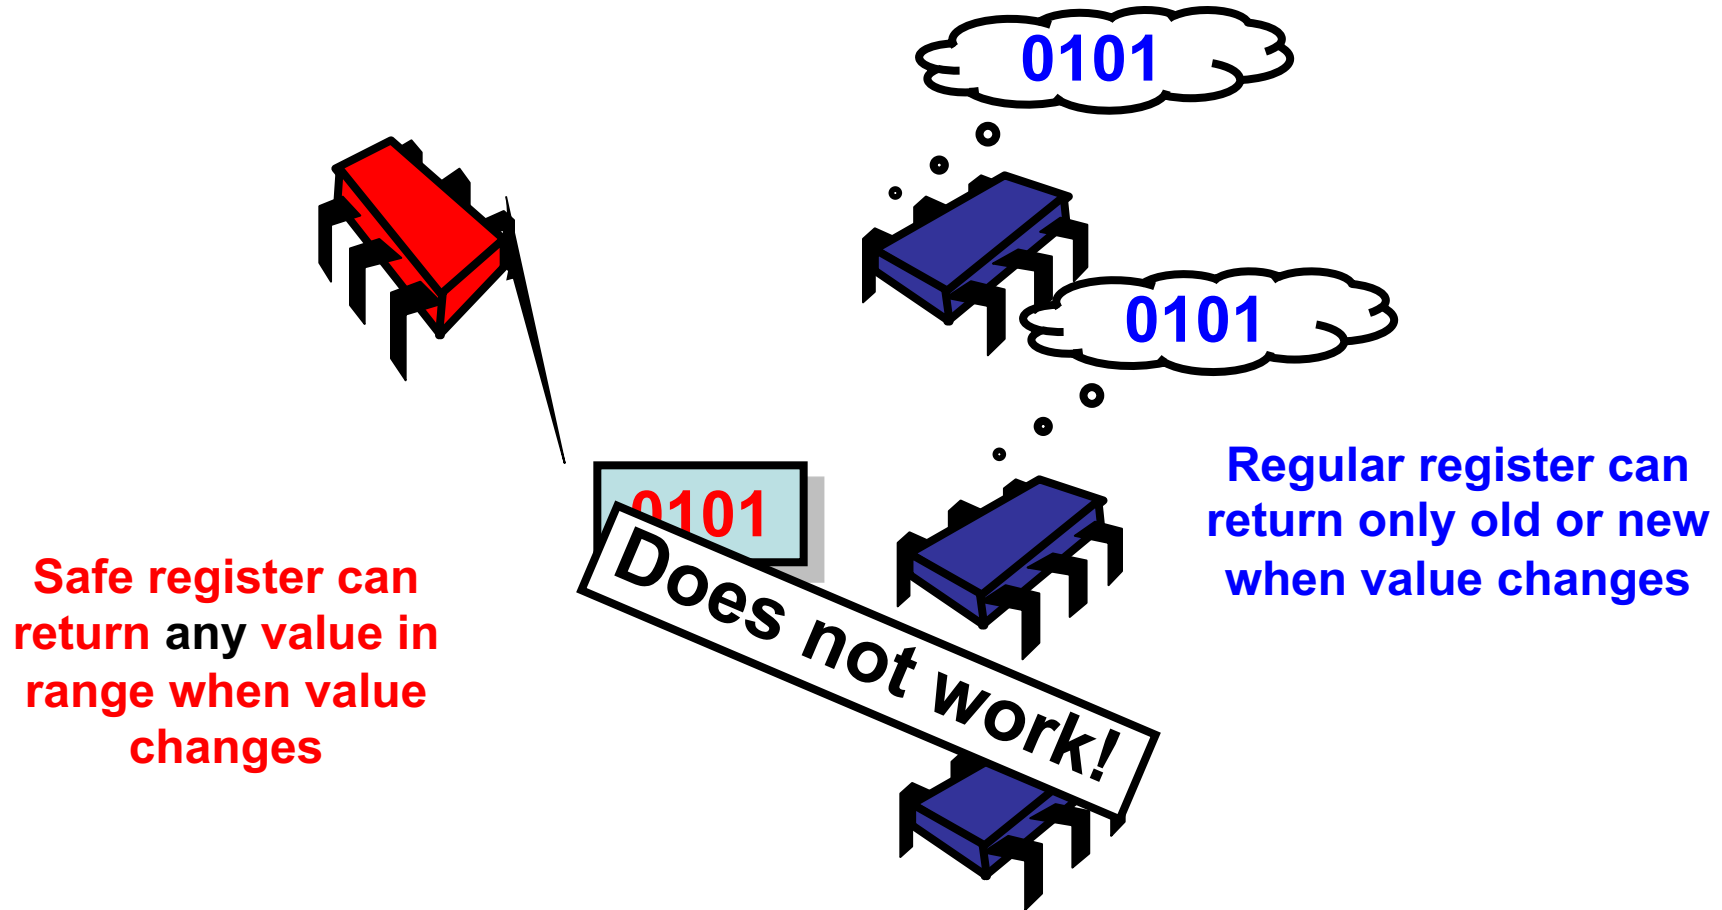
Overlap? What overlap?

No problem

either Boolean value works



Regular Multi-Valued MRSW from Safe Multi-Valued MRSW?




Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot




Questions?



Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- **MRSW regular**
- MRSW atomic
- MRMW atomic
- Atomic snapshot


 **Next**




Representing m Values

Unary representation:


bit[i] means value i




Writing m -Valued Register



Writing m -Valued Register



Writing m -Valued Register



MRSW Regular m -valued from MRSW Regular Boolean

```
public class RegMRSWRegister implements Register{
    RegBoolMRSWRegister[M] bit;

    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }

    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
}
```

MRSW Regular m -valued from MRSW Regular Boolean

```
public class RegMRSWRegister implements Register{  
    RegBoolMRSWRegister[M] bit;  
  
    public void write(int x) {  
        bit[x].write(true);  
        for (int i=x-1; i>=0; i--)  
            bit[i].write(false);  
    }  
  
    public int read() {  
        for (int i=0; i < M; i++)  
            if (bit[i].read())  
                return i;  
    }  
}
```

Unary representation:
bit[i] means value i

MRSW Regular m -valued from MRSW Regular Boolean

```
public class RegMRSWRegisterimplements Register {
    RegBoolMRSWRegister[m] bit;

    public void write(int x) {
        bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            bit[i].write(false);
    }

    public int read() {
        for (int i=0; i < M; i++)
            if (bit[i].read())
                return i;
    }
}
```

set bit x


MRSW Regular m -valued from MRSW Regular Boolean

```
public class RegMRSWRegisterimplements Register {
    RegBoolMRSWRegister[m] bit;

    public void write(int x) {
        bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            bit[i].write(false);
    }

    public int read() {
        for (int i=0; i < M; i++)
            if (bit[i].read())
                return i;
    }
}
```

**Clear bits
from higher
to lower**




MRSW Regular m -valued from MRSW Regular Boolean

```
public class RegMRSWRegisterimplements Register {
    RegBoolMRSWRegister[m] bit;

    public void write(int x) {
        bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            bit[i].write(false);
    }

    public int read() {
        for (int i=0; i < M; i++)
            if (bit[i].read())
                return i;
    }
}
```

**Scan from lower
to higher & return
first bit set**




Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot




Questions?



Road Map


- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot



Road Map (Slight Detour)


- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot

 SRSW Atomic




SRSW Atomic From SRSW Regular

Regular writer



5678



Regular reader

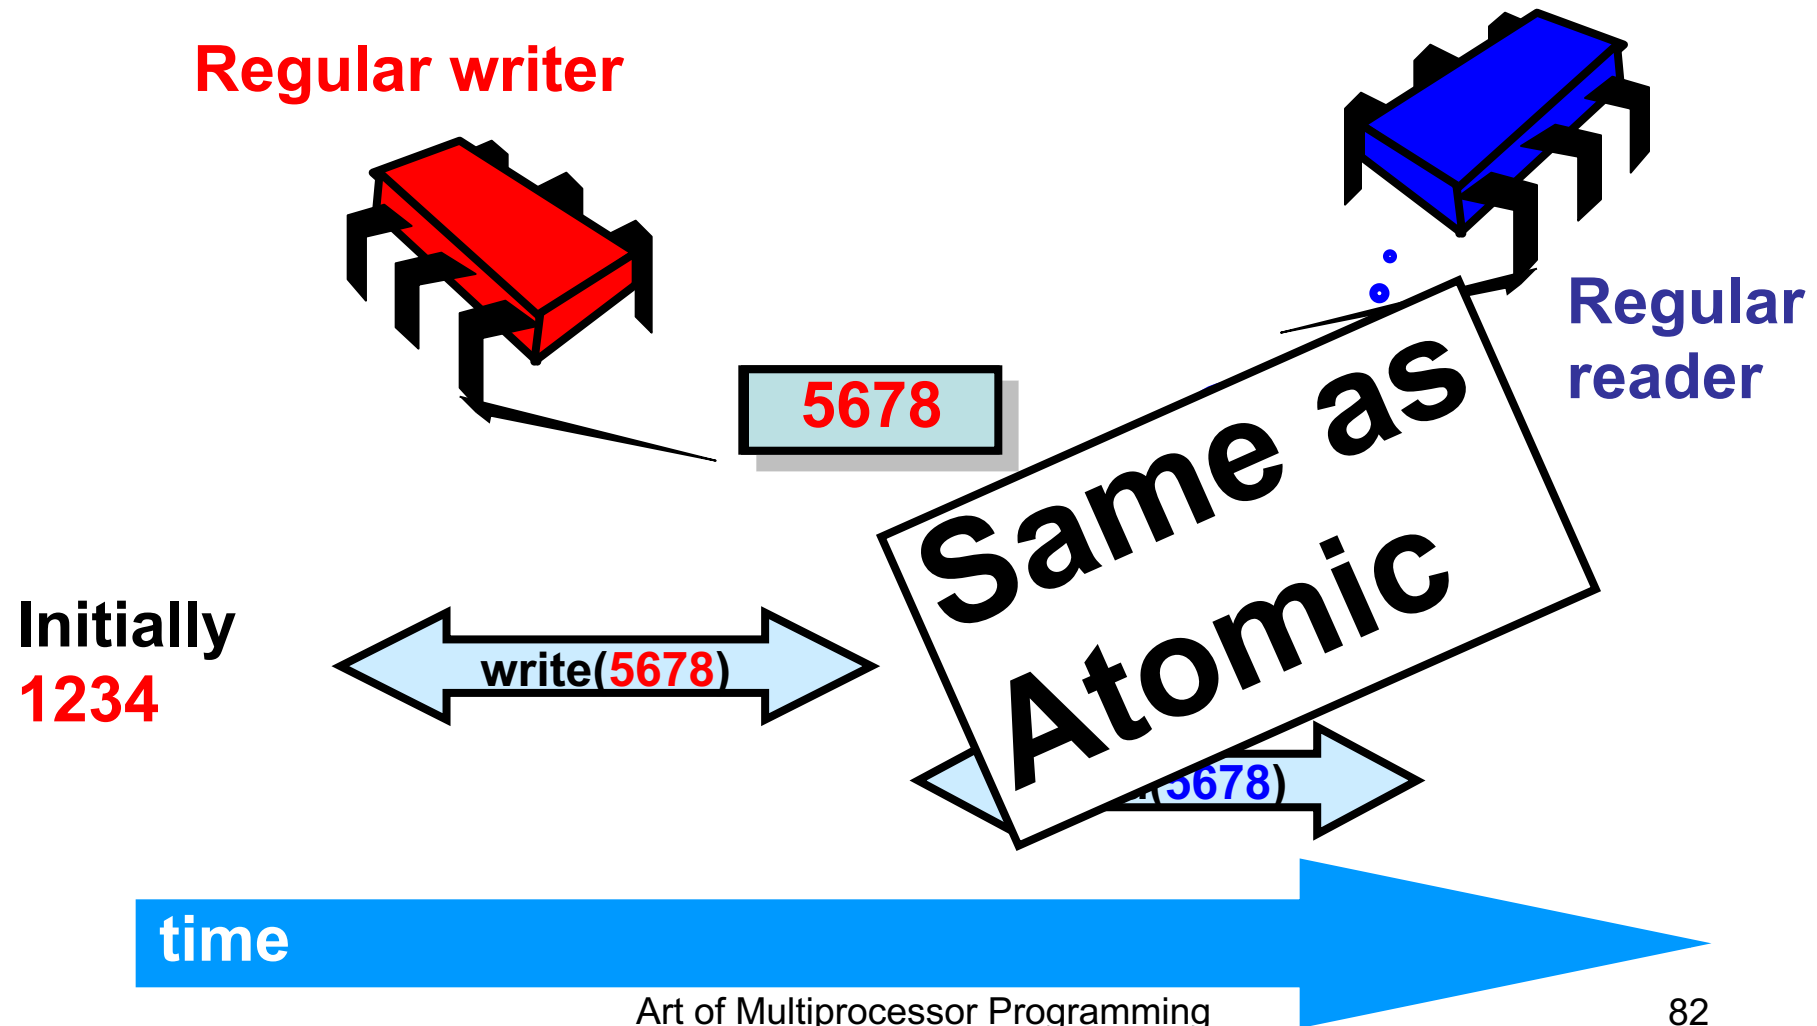
1234

Instead of 5678...

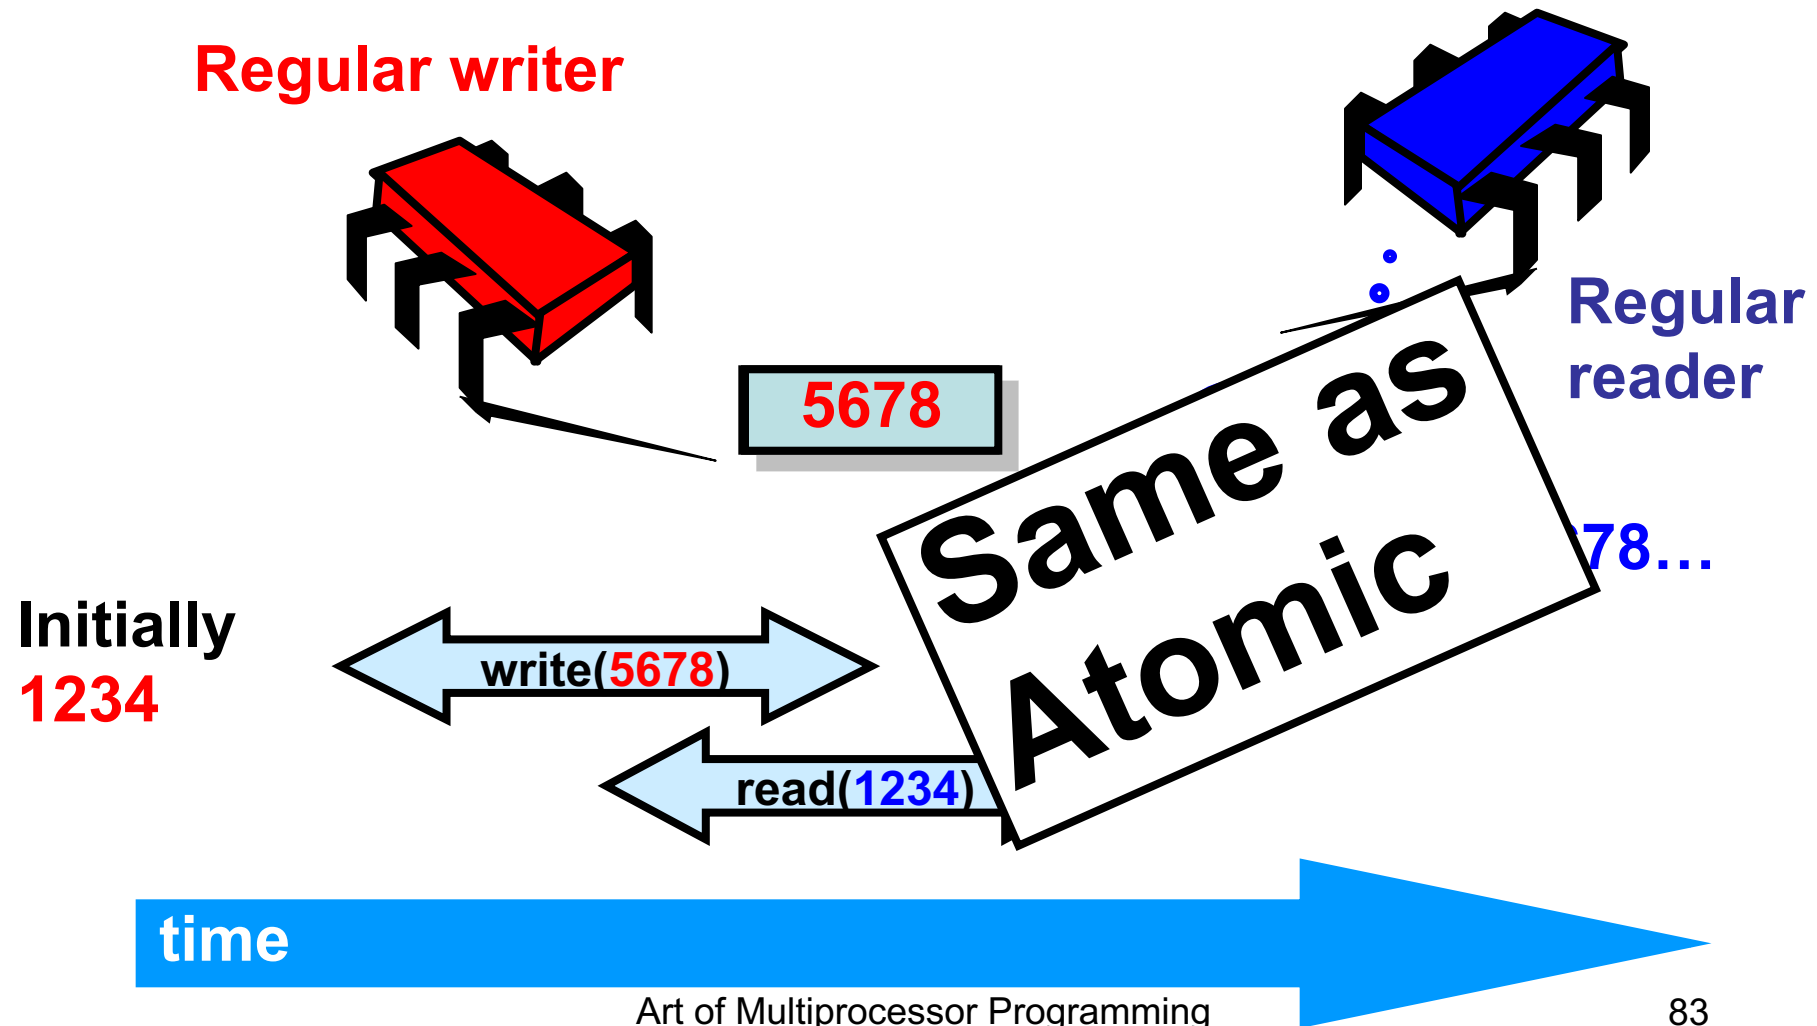
Concurrent Reading

When is this a problem?

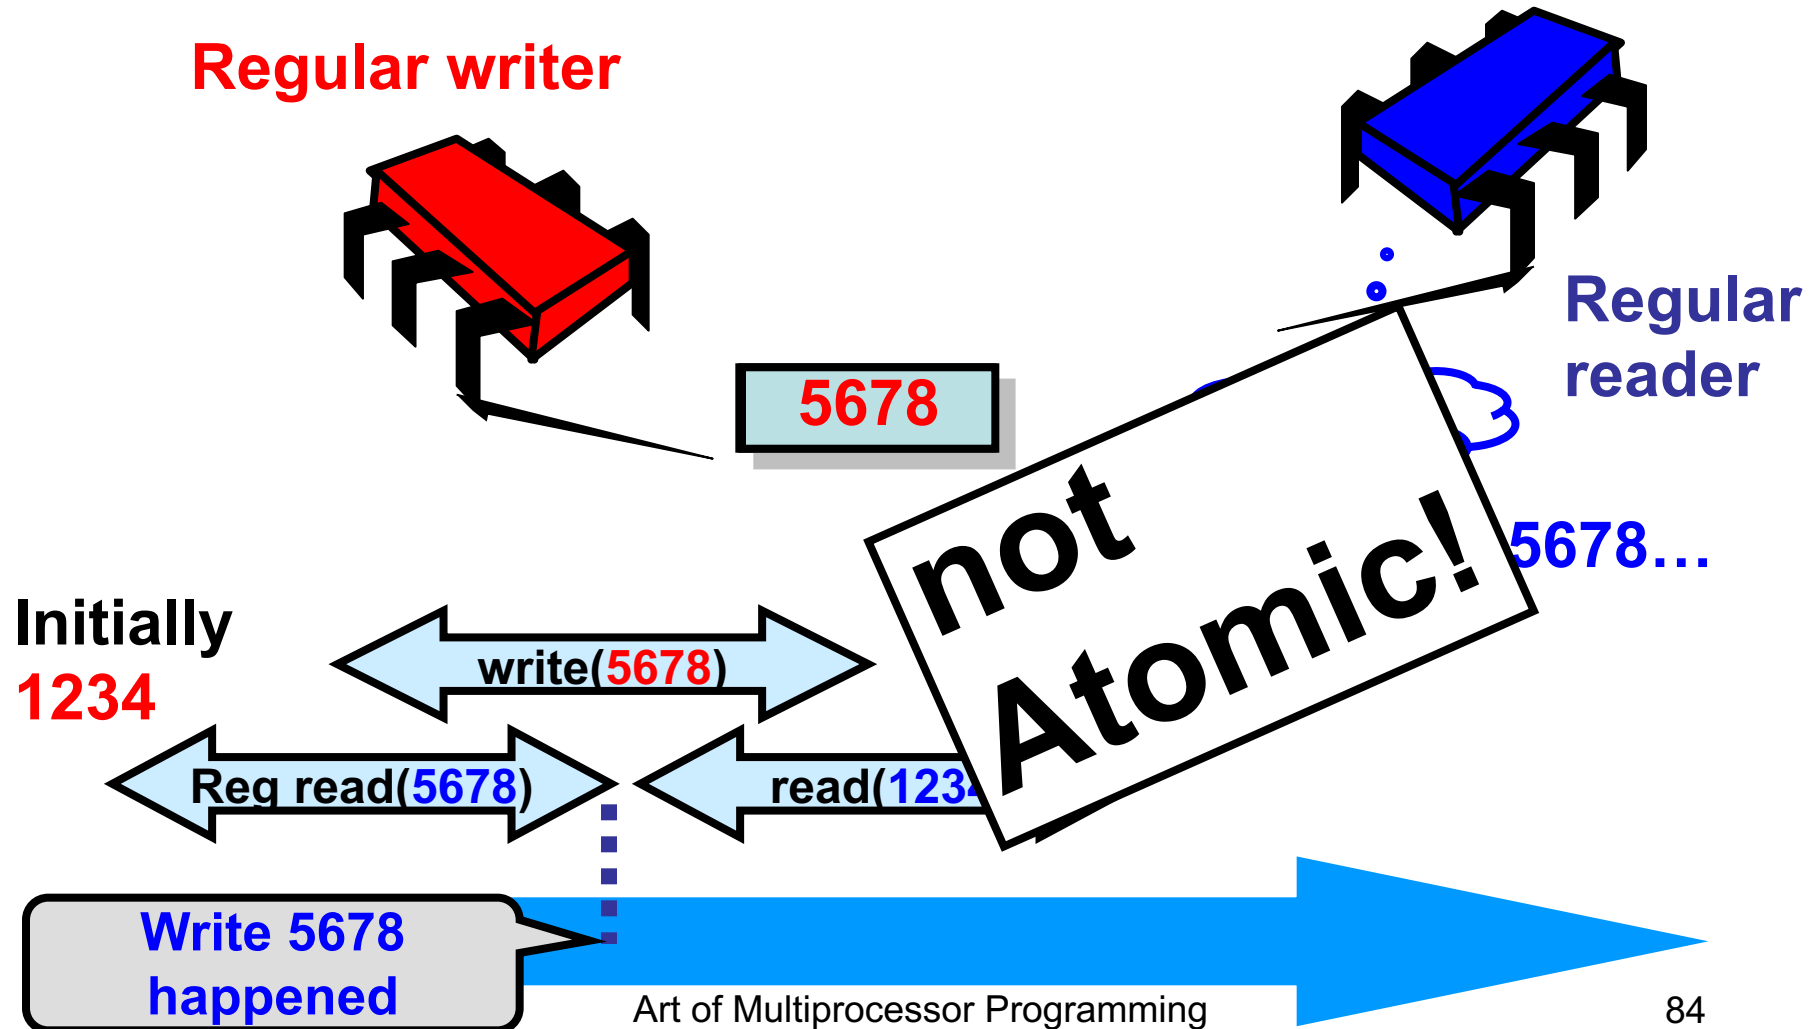
SRSW Atomic From SRSW Regular



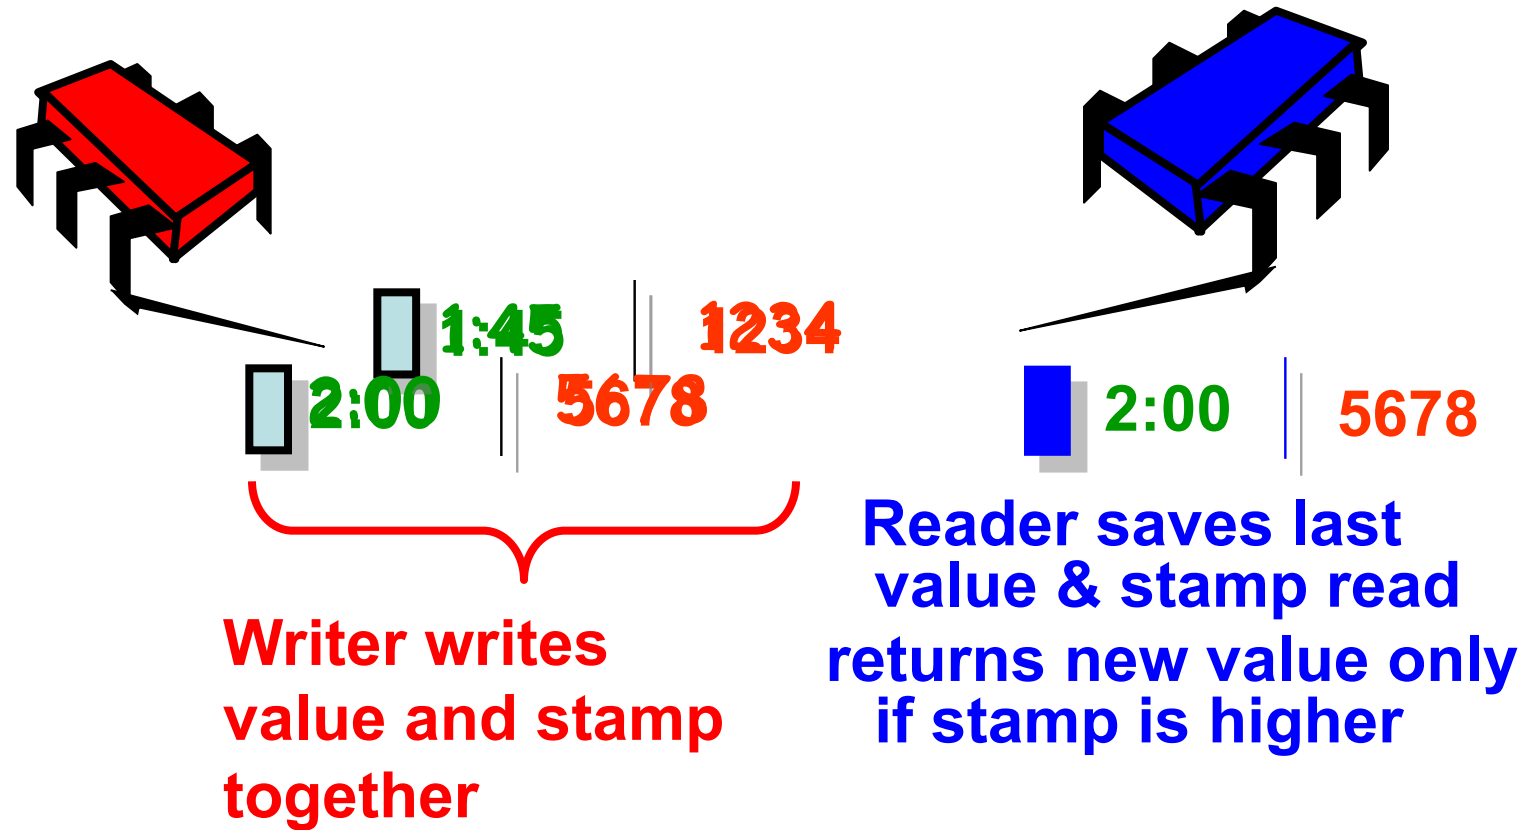
SRSW Atomic From SRSW Regular



SRSW Atomic From SRSW Regular

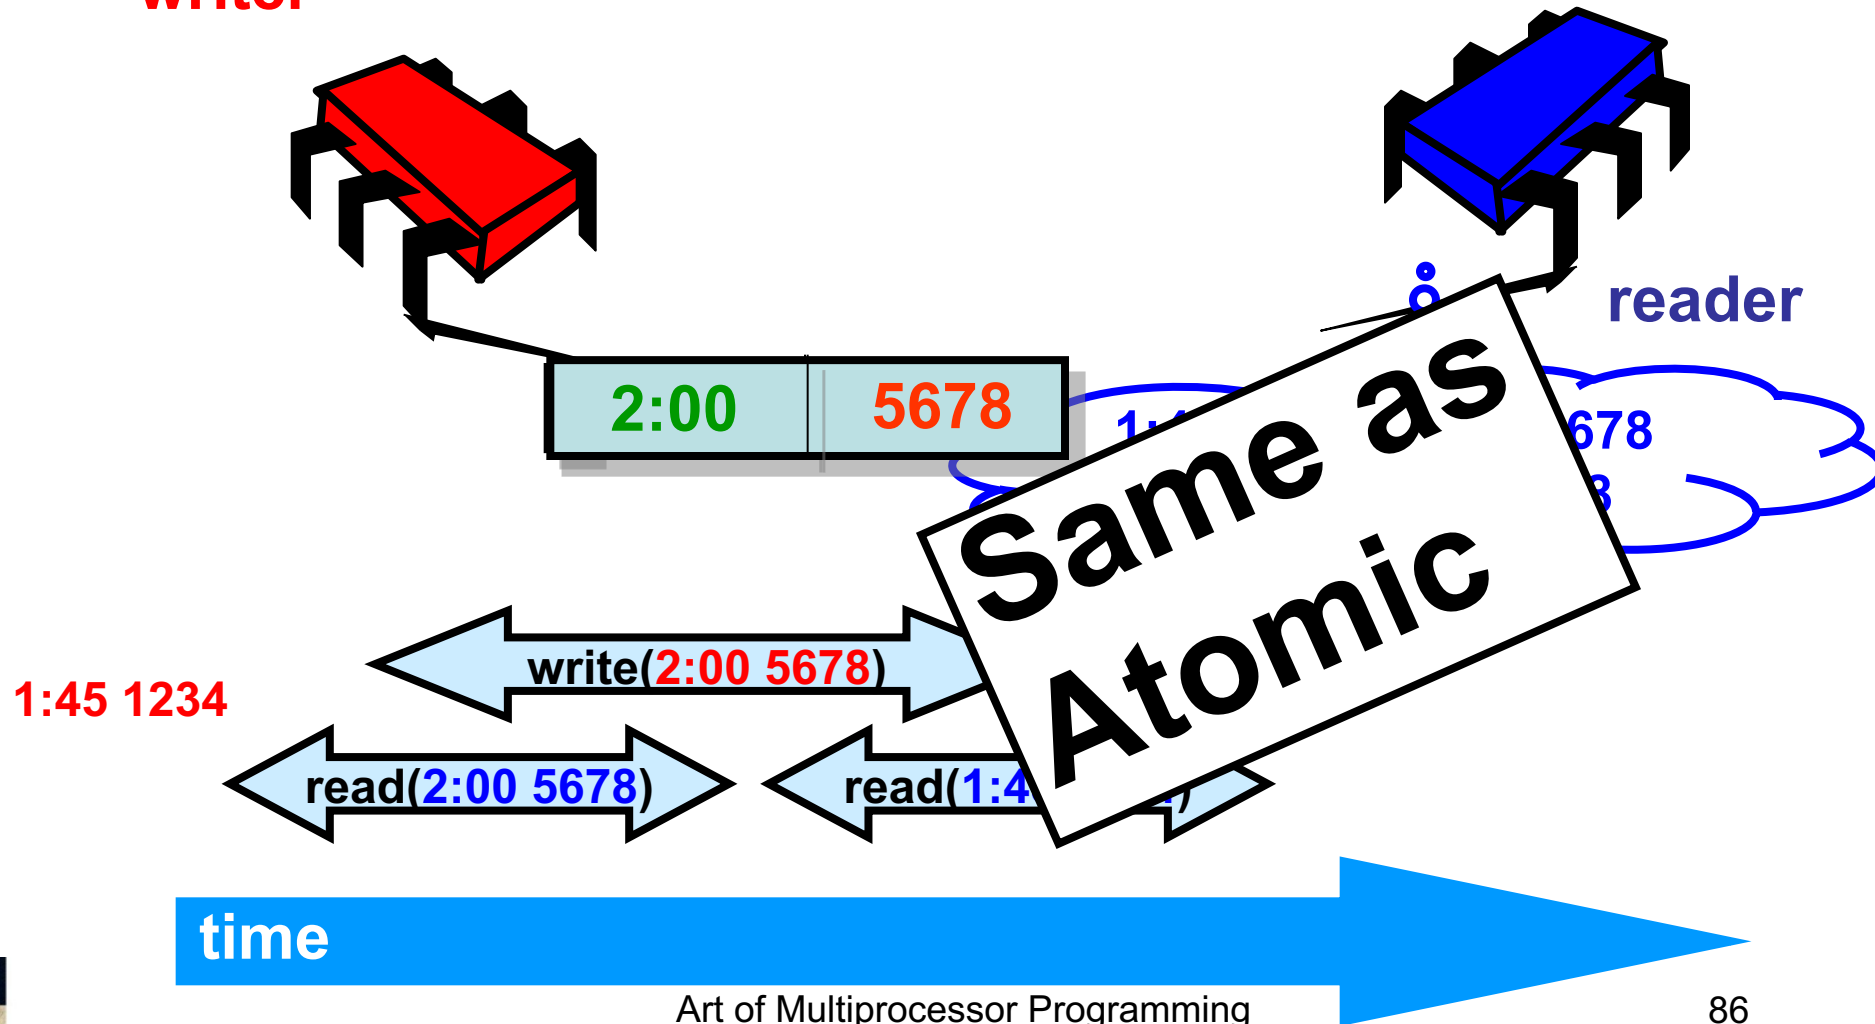


Timestamped Values



SRSW Atomic From SRSW Regular

writer




Atomic Single-Reader to Atomic Multi-Reader


stamp	value
1:45	1234
1:45	1234
1:45	1234
1:45	1234

} One per reader

Another Scenario




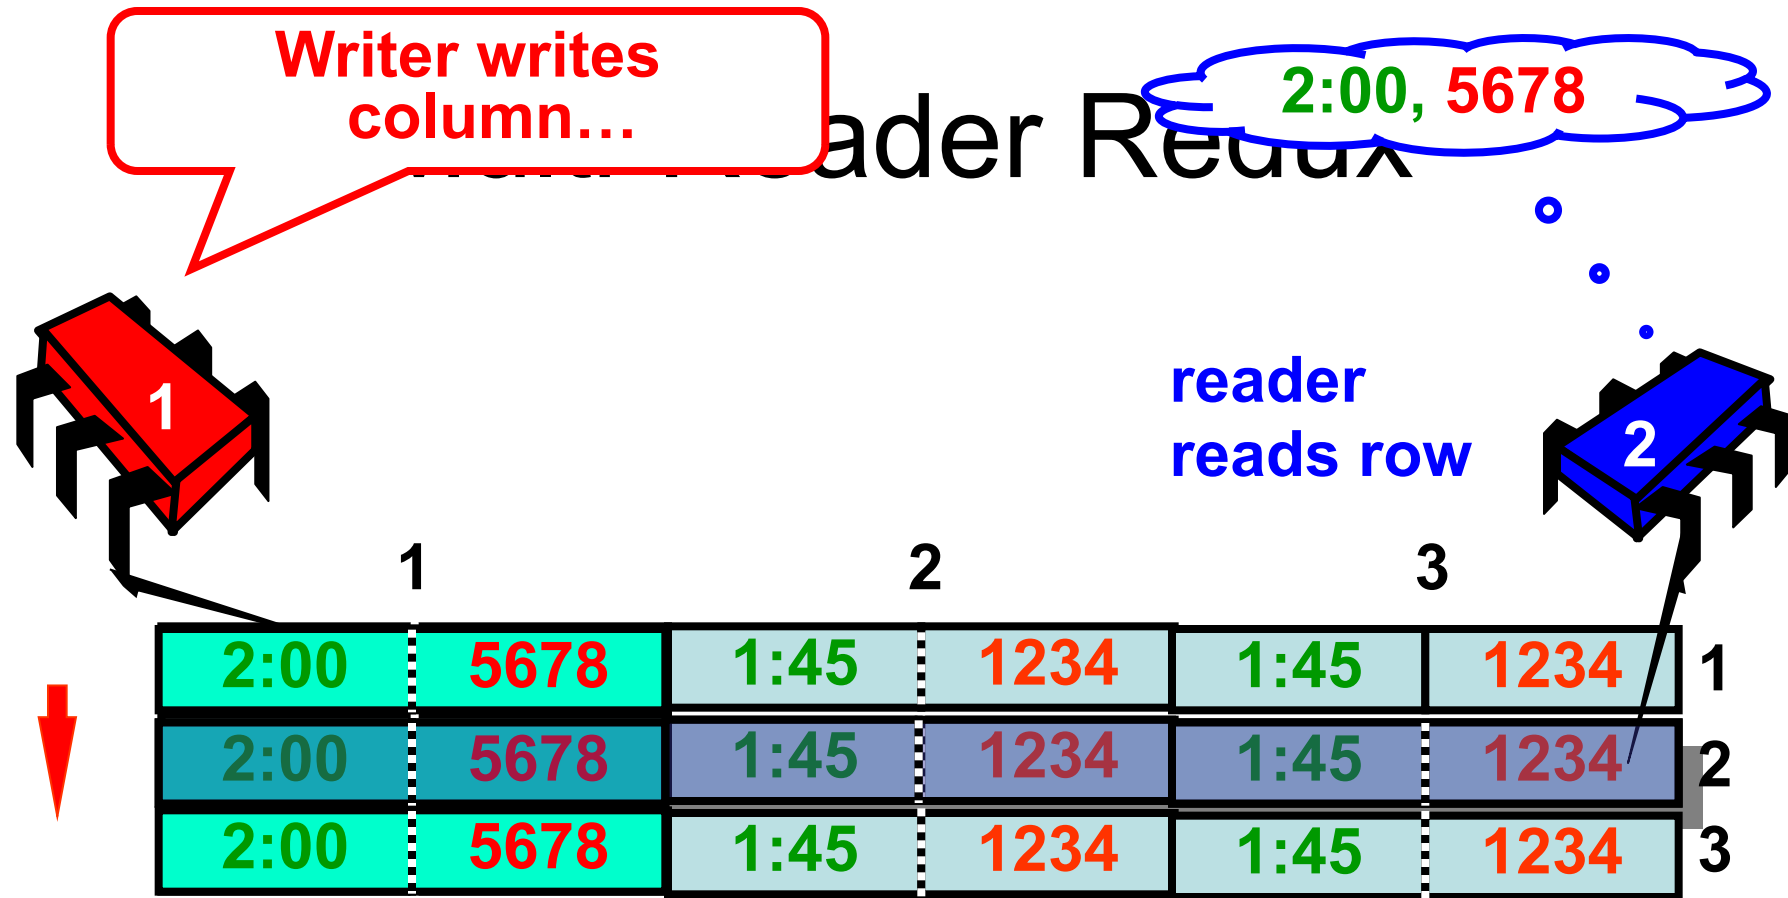
Another Scenario




**Yellow was completely after Blue but
read earlier value...not linearizable!**

Multi-Reader Redux

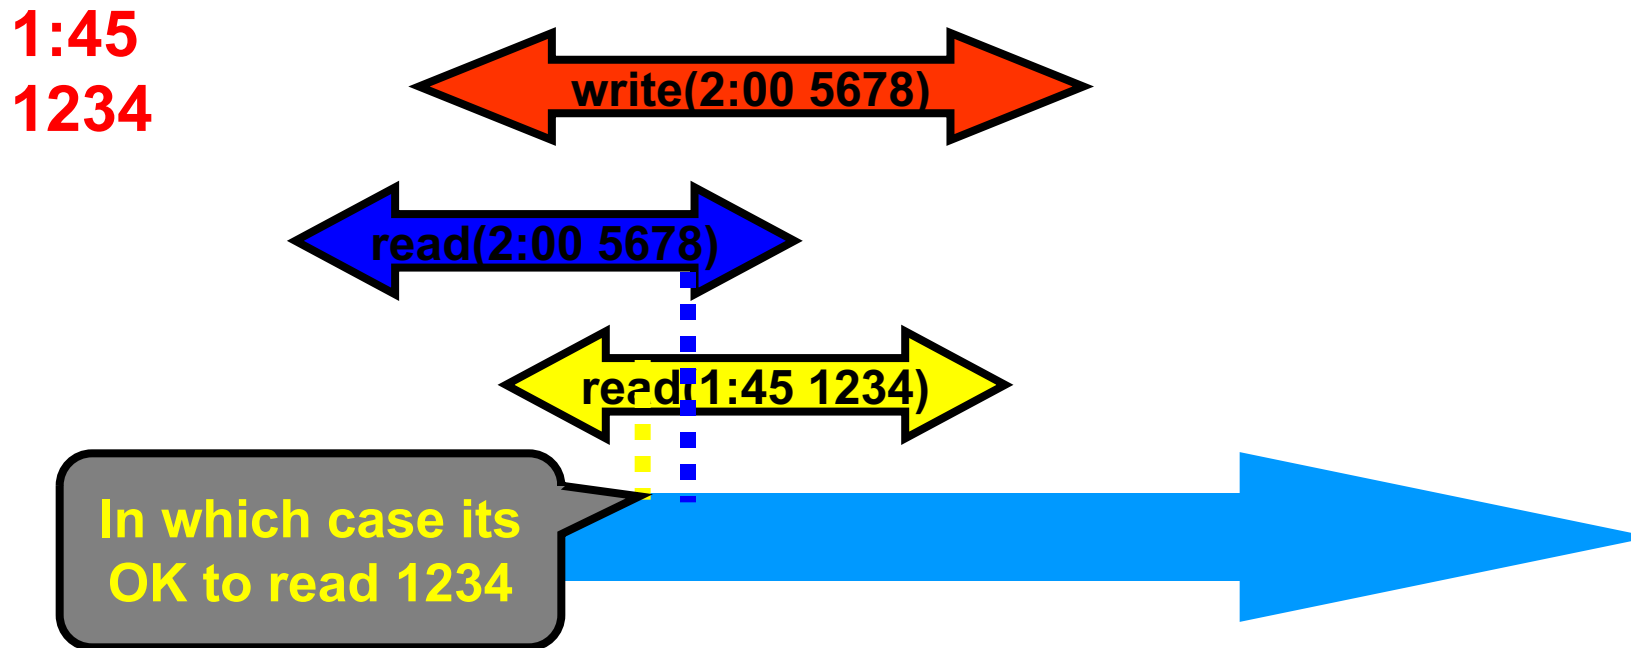




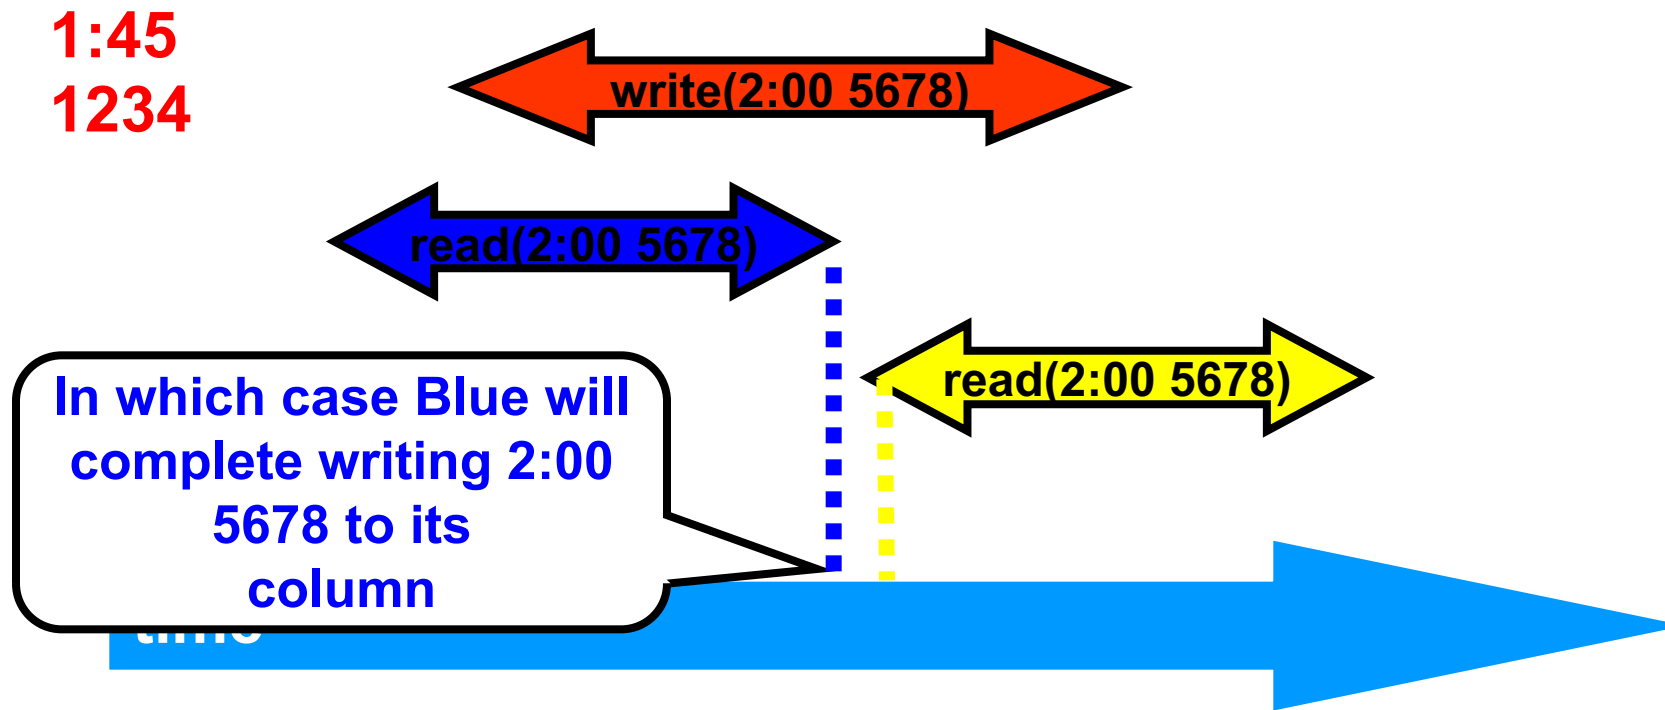
Multi Reader Redux



Can't Yellow Miss Blue's Update? ... Only if Readers Overlap...




Bad Case Only When Readers Don't Overlap

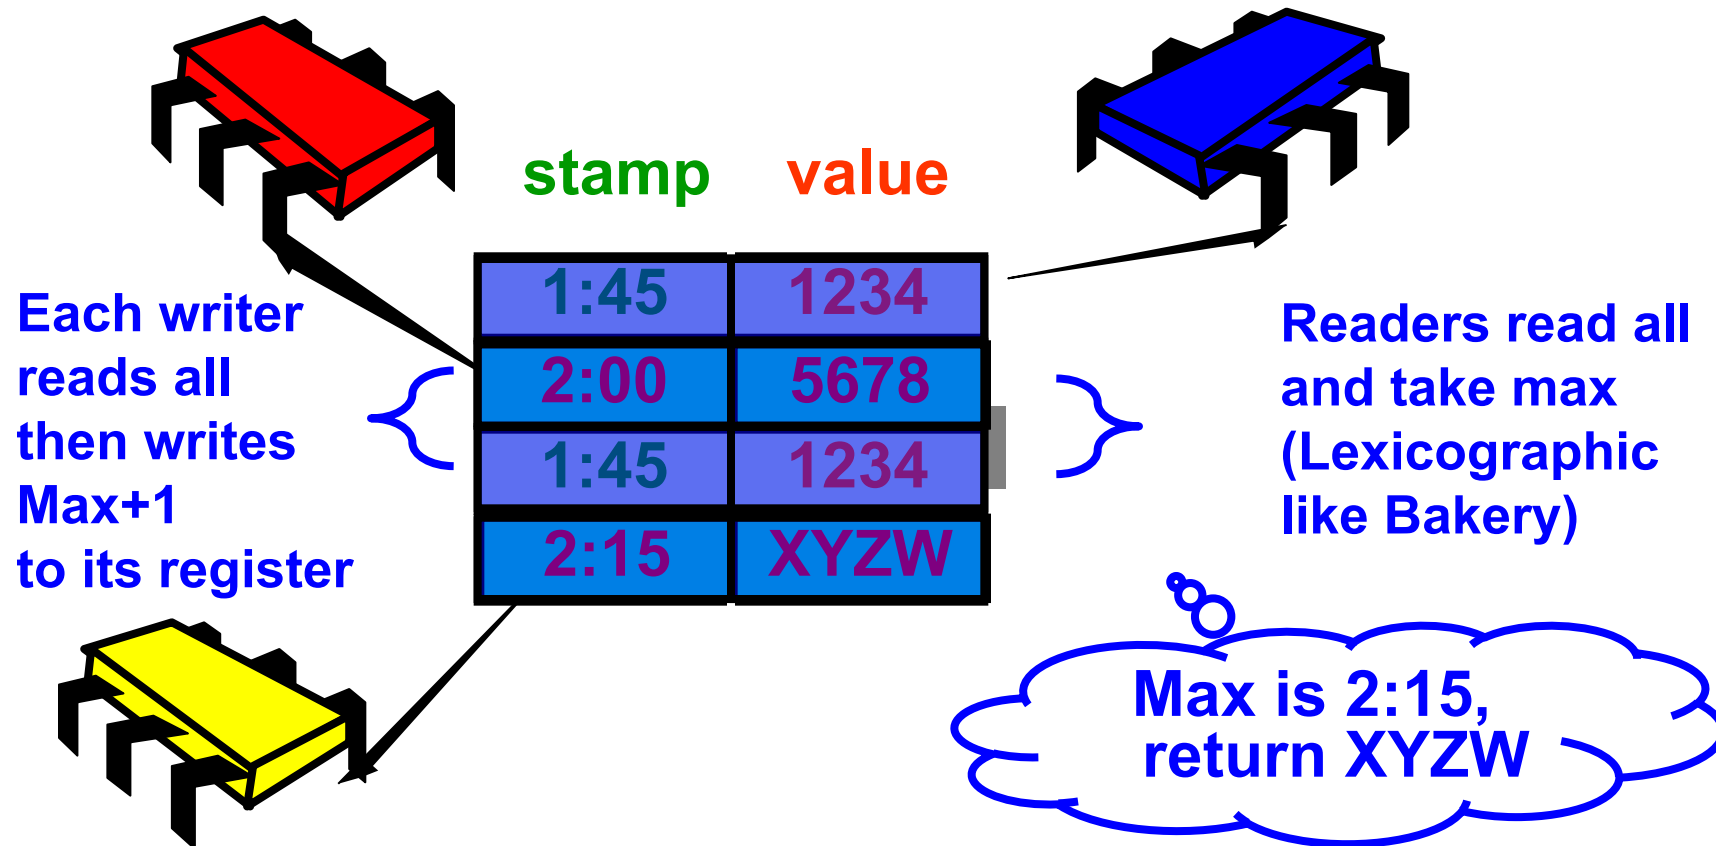


Road Map


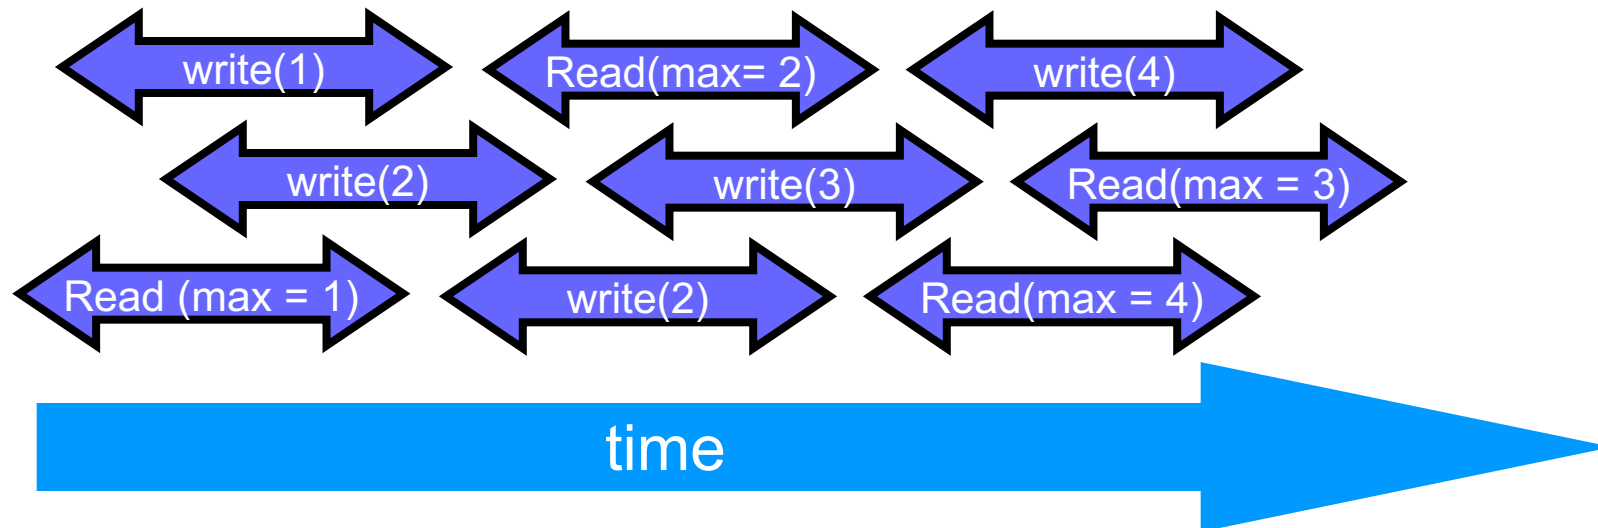
- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot




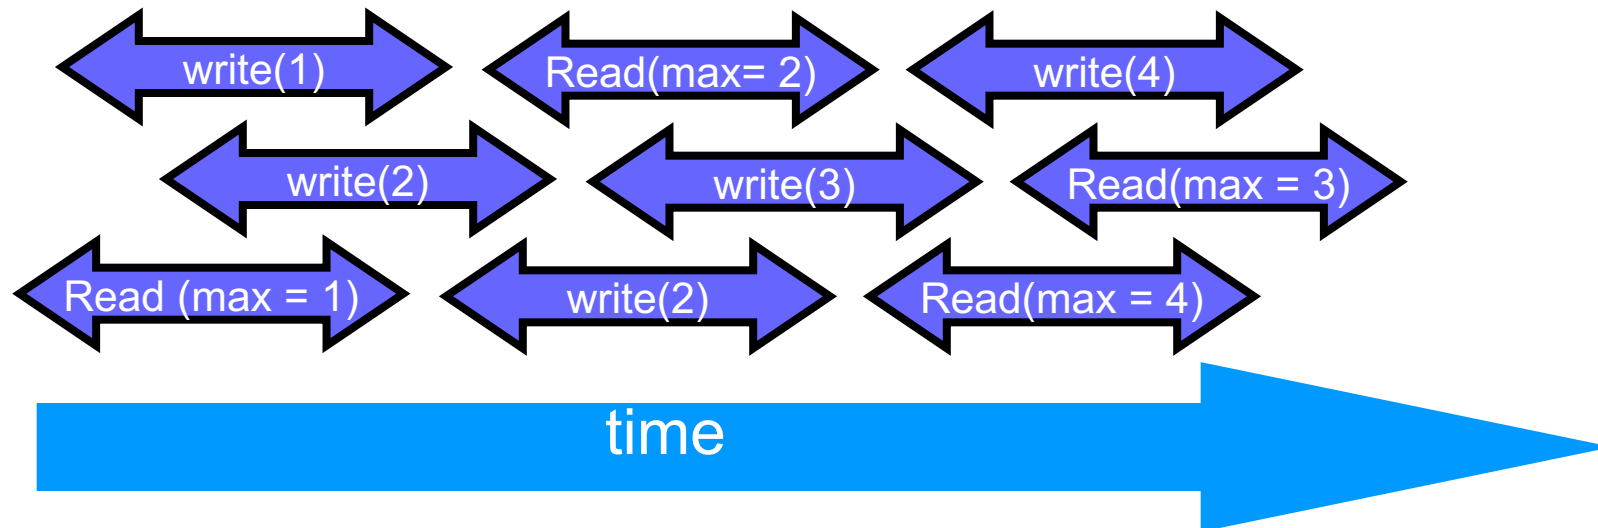
Multi-Writer Atomic From Multi-Reader Atomic



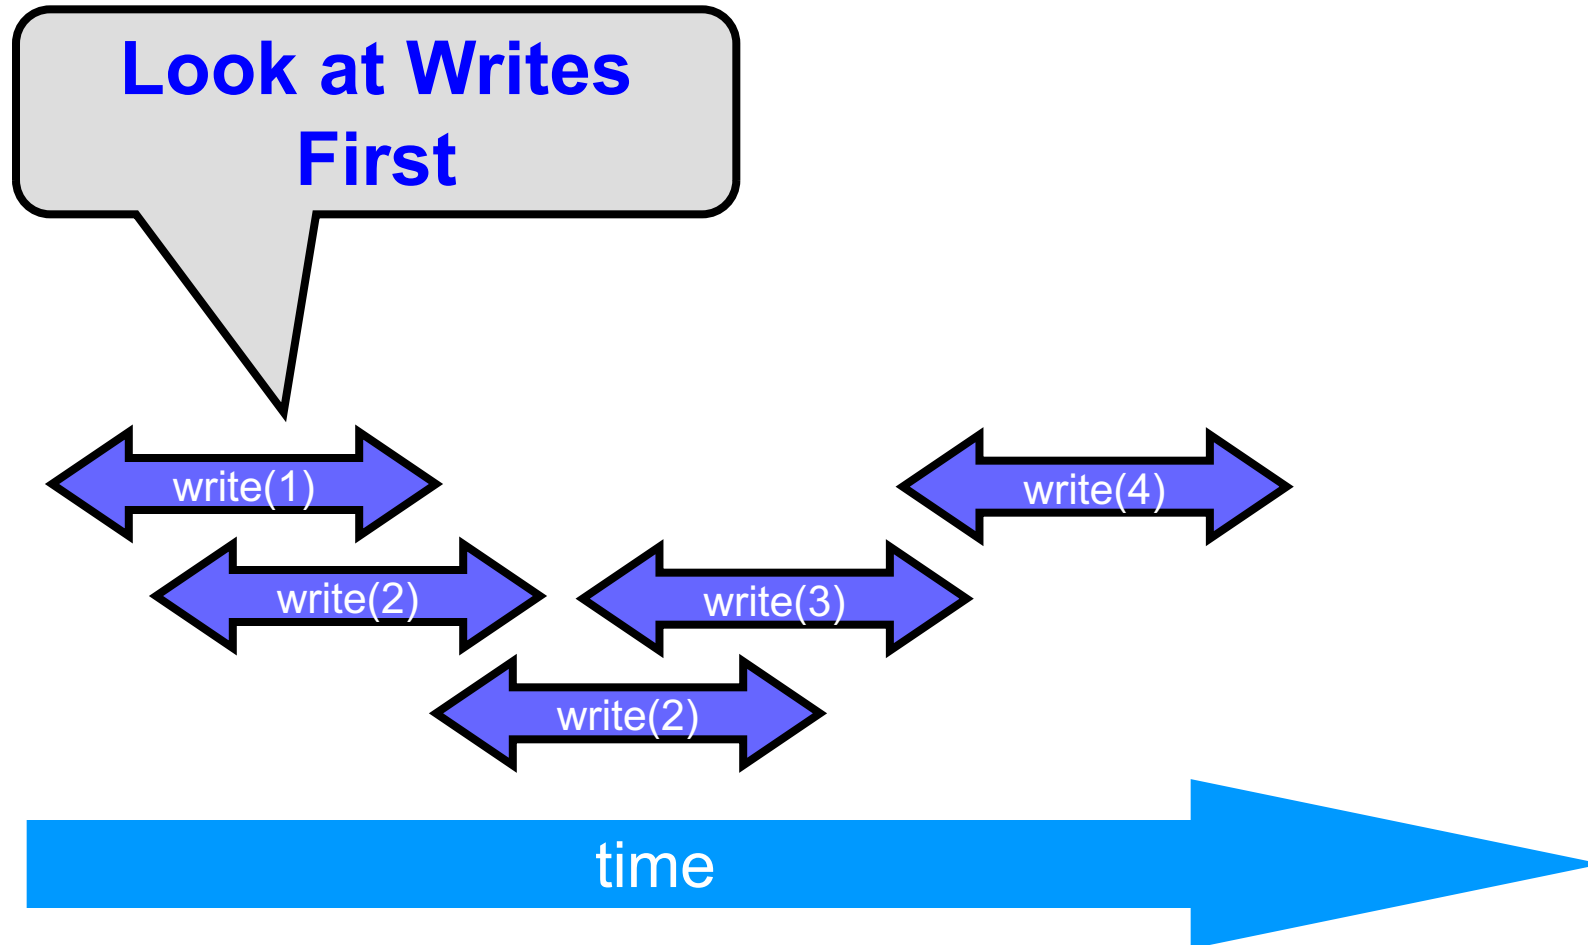
Atomic Execution Means it is Linearizable



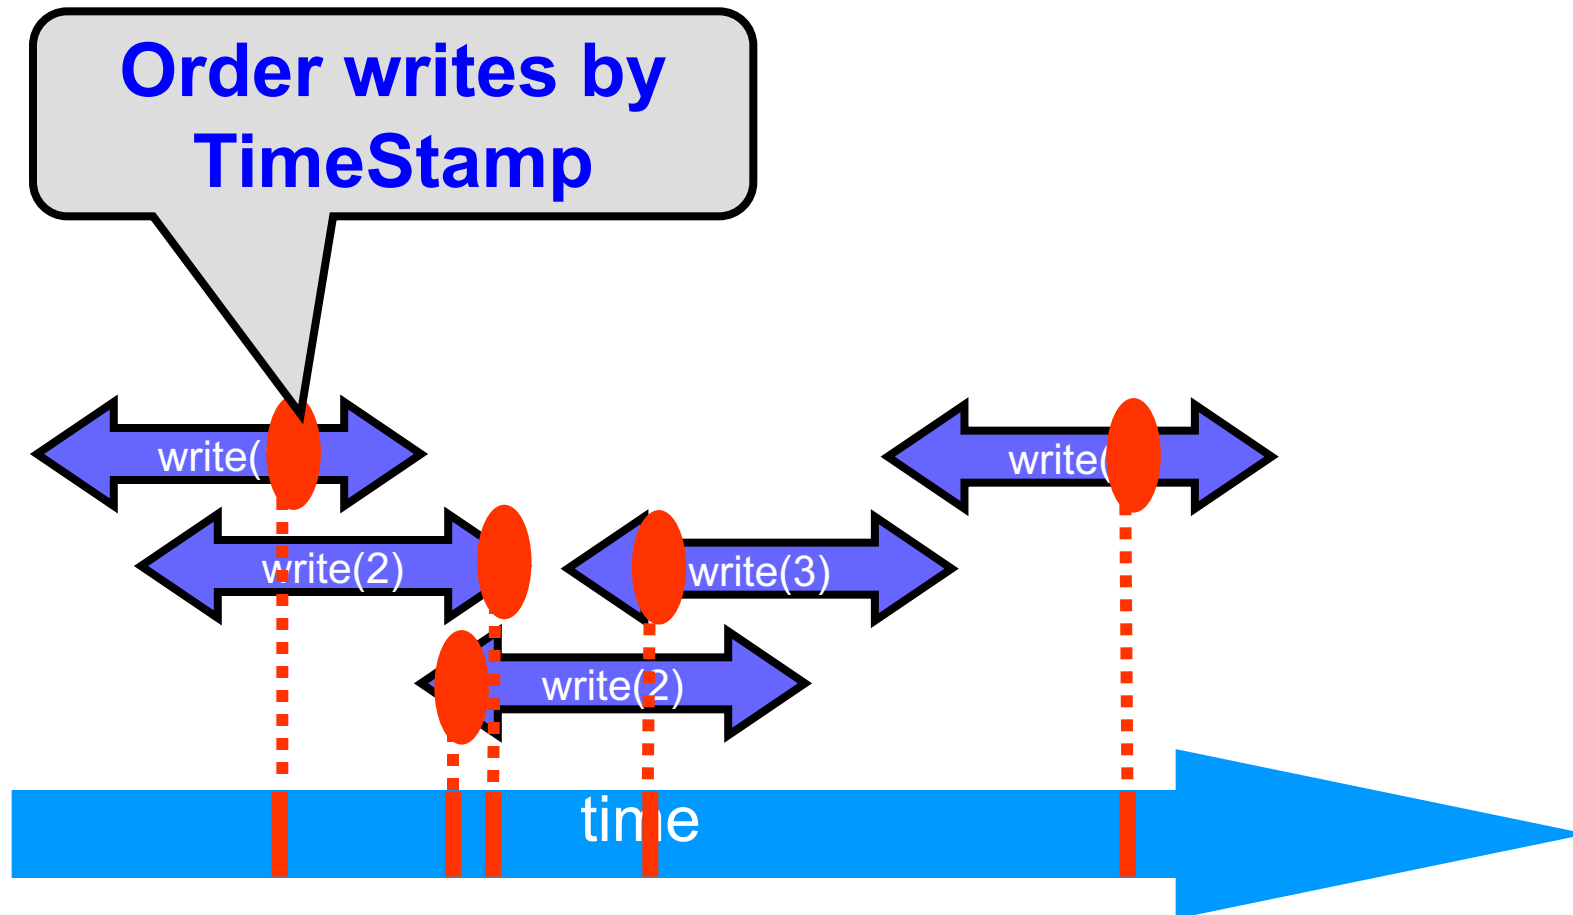
Linearization Points



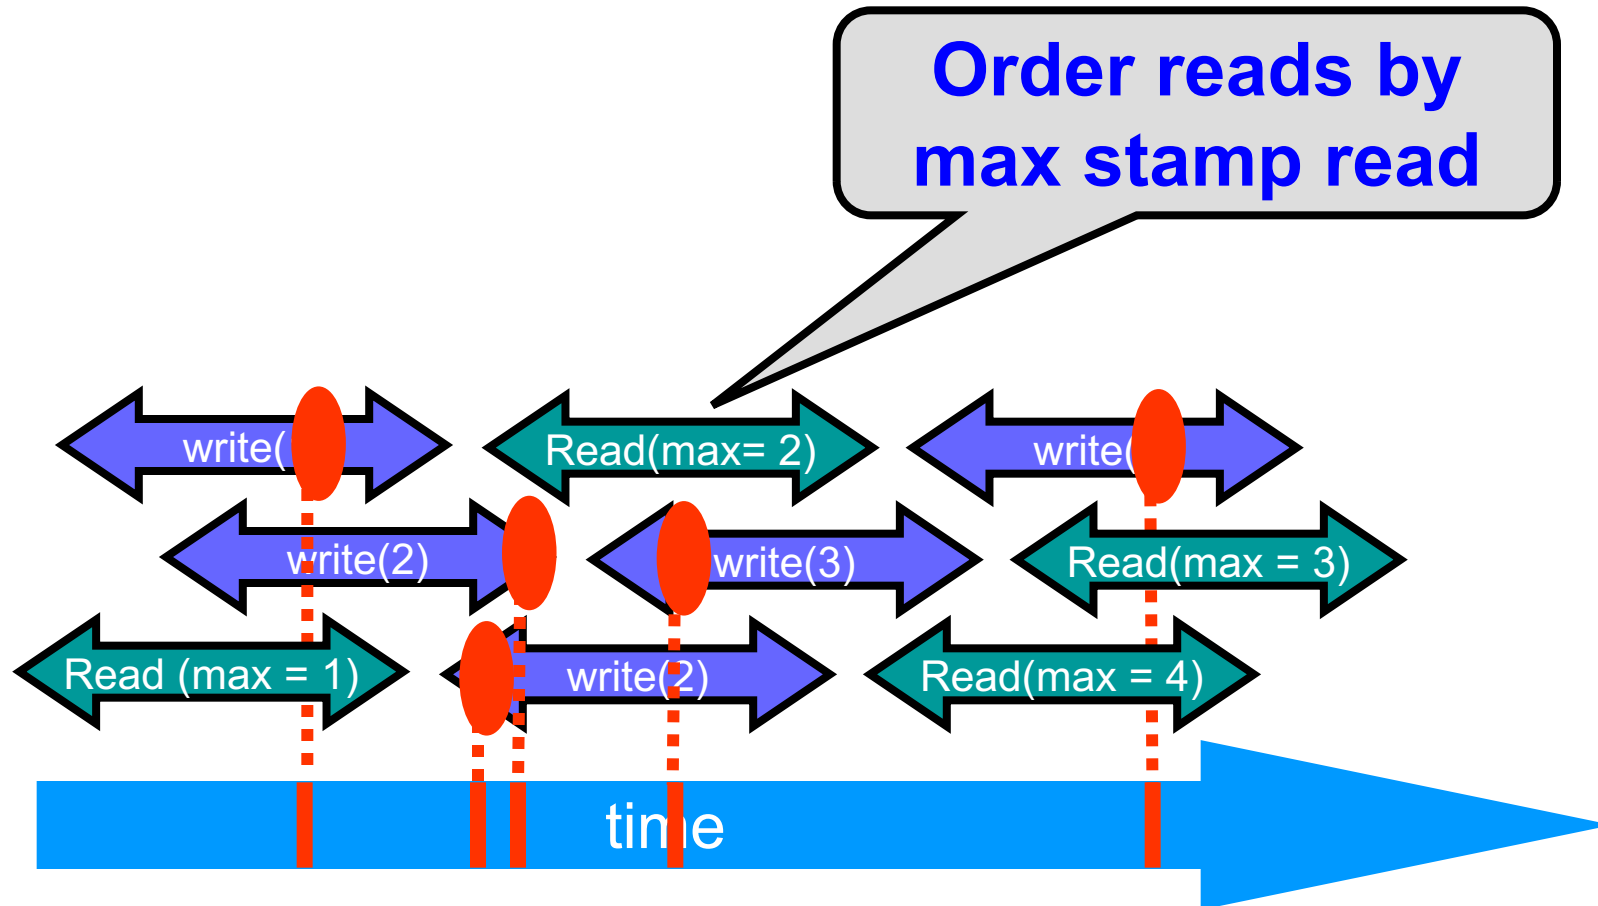
Linearization Points



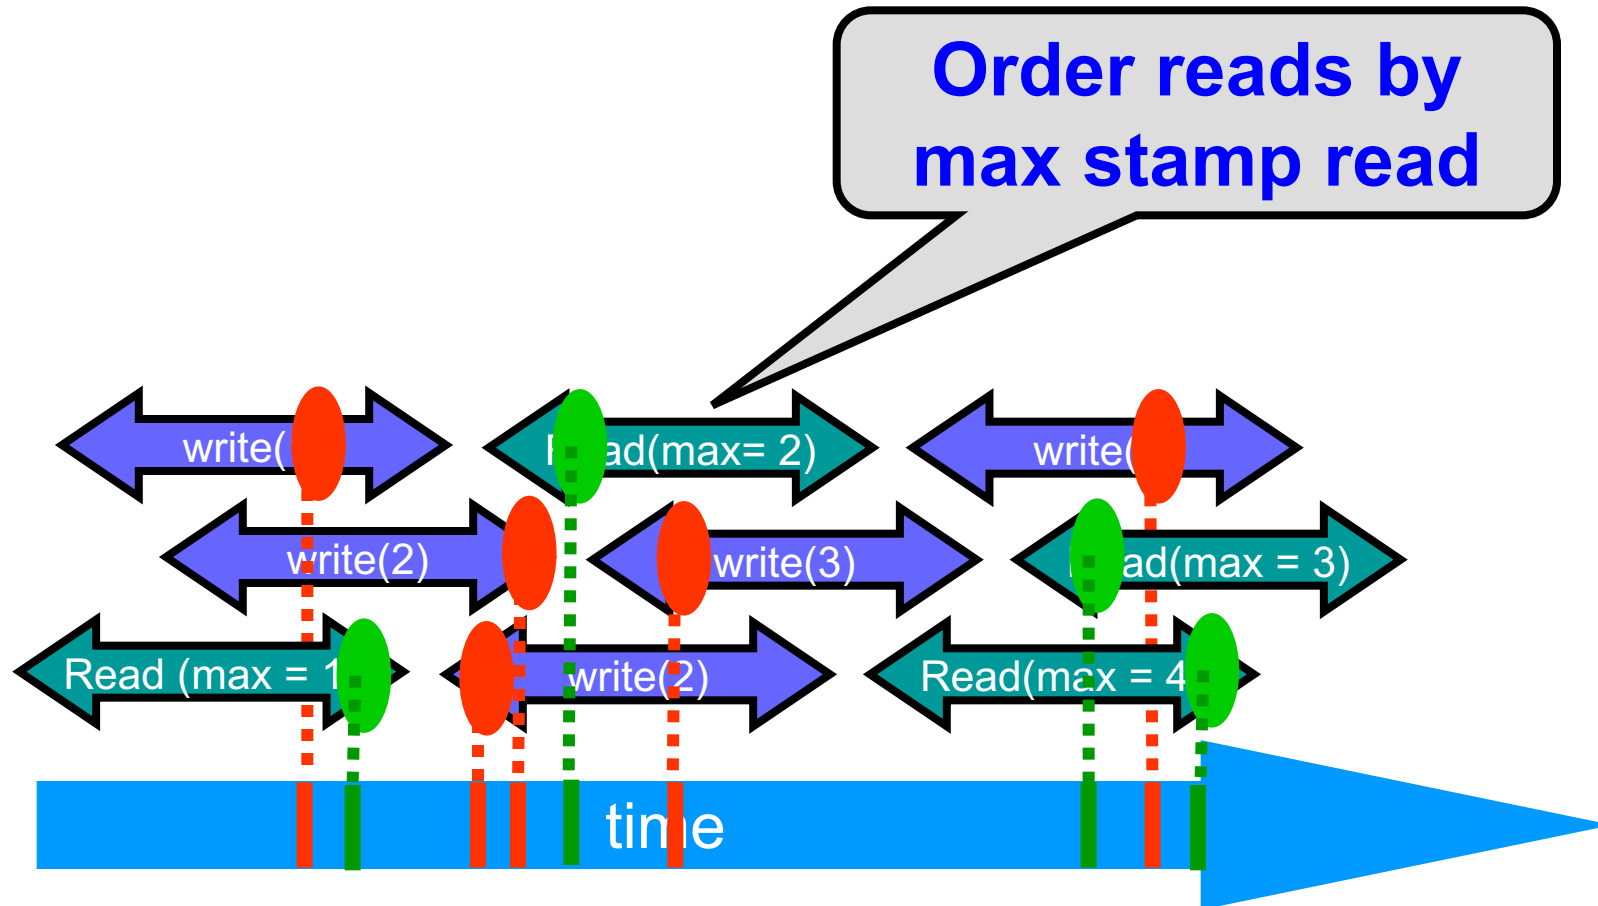
Linearization Points



Linearization Points

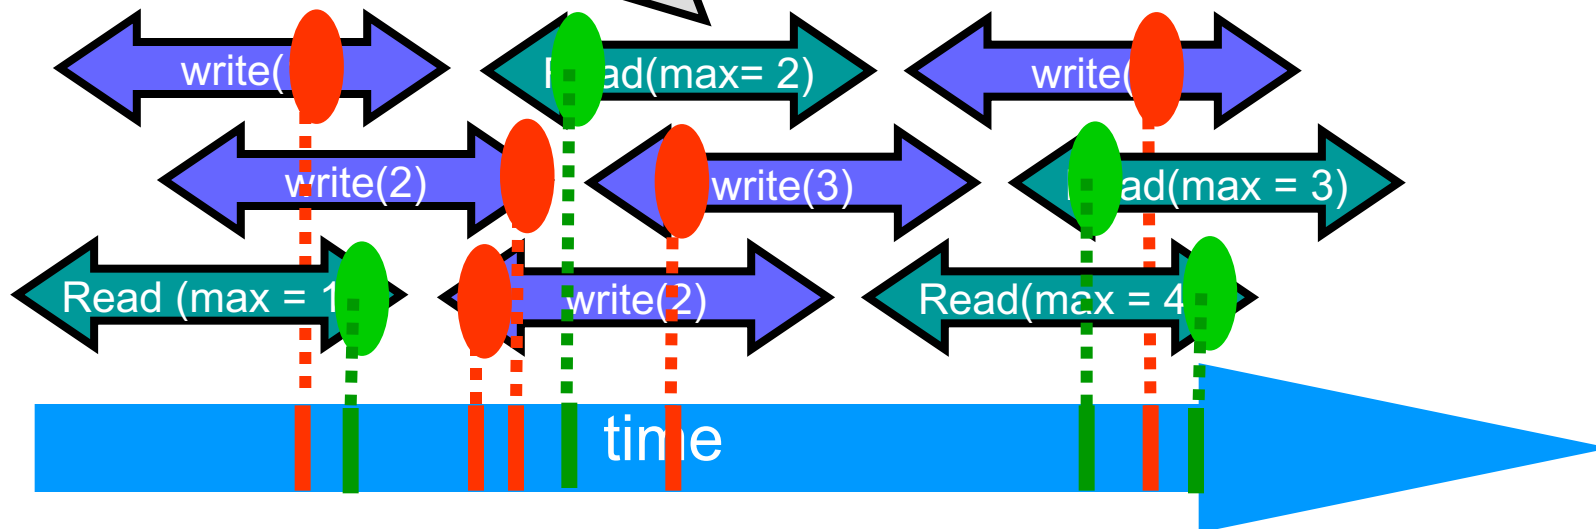


Linearization Points



Linearization Points

The linearization point depends on the execution (not a line in the code)!




Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot



Questions?




Road Map


- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic
- Atomic snapshot



Next




Atomic Snapshot



Atomic Snapshot

- Array of SWMR atomic registers
- Take instantaneous snapshot of all
- Generalizes to MRMW registers ...



Snapshot Interface

```
public interface Snapshot {  
    public int update(int v);  
    public int[] scan();  
}
```

Snapshot Interface

Thread *i* writes *v* to its register

```
public interface Snapshot {  
    public int update(int v);  
    public int[] scan();  
}
```


Snapshot Interface

Instantaneous snapshot of all threads' registers

```
public interface Snapshot {  
    public int update(int v);  
    public int[] scan();  
}
```


Atomic Snapshot

- Collect
 - Read values one at a time
- Problem
 - Incompatible concurrent collects
 - Result not linearizable



Clean Collects



- Clean Collect
 - Collect during which nothing changed
 - Can we make it happen?
 - Can we detect it?



Simple Snapshot


- Put increasing labels on each entry
- Collect
- If both
 - We're done
- Otherwise,
 - Try again

Problem: Scanner might not be collecting a snapshot!




Claim:

But scanner sees x and z together!




time

x and z are never
in memory
together





Scanner reads x and z with different labels and recognizes collect not clean




Simple Snapshot

- Collect twice
- If both agree,
 - We're done
- Otherwise,
 - Try again



Simple Snapshot: Update


```
public class SimpleSnapshot implements Snapshot {  
    private AtomicMRSWRegister[] register;  
  
    public void update(int value) {  
        int i = Thread.myIndex();  
        LabeledValue oldValue = register[i].read();  
        LabeledValue newValue =  
            new LabeledValue(oldValue.label+1, value);  
        register[i].write(newValue);  
    }  
}
```



Simple Snapshot: Update

```
public class SimpleSnapshot implements Snapshot {  
    private AtomicMRSWRegister[] register;  
  
    public void update(int value) {  
        int i = Thread.myIndex();  
        LabeledValue oldValue = register[i].read();  
        LabeledValue newValue =  
            new LabeledValue(oldValue.label+1, value);  
        register[i].write(newValue);  
    }  
}
```


One single-writer register per thread



Simple Snapshot: Update


```
public class SimpleSnapshot implements Snapshot {  
    private AtomicMRSWRegister[] register;  
  
    public void update(int value) {  
        int i = Thread.myIndex();  
        LabeledValue oldValue = register[i].read();  
        LabeledValue newValue =  
            new LabeledValue(oldValue.label+1, value);  
        register[i].write(newValue);  
    }  
}
```

Write each time with higher label



Simple Snapshot: Collect

```
private LabeledValue[] collect() {  
    LabeledValue[] copy =  
        new LabeledValue[n];  
    for (int j = 0; j < n; j++)  
        copy[j] = this.register[j].read();  
    return copy;  
}
```




Simple Snapshot

```
private LabeledValue[] collect() {  
    LabeledValue[] copy =  
        new LabeledValue[n];  
    for (int j = 0; j < n; j++)  
        copy[j] = this.register[j].read();  
    return copy;  
}
```

Just read each register into array

Simple Snapshot: Scan

```
public int[] scan() {  
    LabeledValue[] oldCopy, newCopy;  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        if (!equals(oldCopy, newCopy)) {  
            oldCopy = newCopy;  
            continue collect;  
        }  
        return getValues(newCopy);  
    }  
}
```



Simple Snapshot: Scan

```
public int[] scan() {  
    LabeledValue[] oldCopy, newCopy;  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        if (!equals(oldCopy, newCopy)) {  
            oldCopy = newCopy;  
            continue collect;  
        }  
        return getValues(newCopy);  
    }  
}
```

Collect once

Simple Snapshot: Scan

```
public int[] scan() {  
    LabeledValue[] oldCopy, newCopy;  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        if (!equals(oldCopy, newCopy)) {  
            oldCopy = newCopy;  
            continue collect;  
        }  
        return getValues(newCopy);  
    }  
}
```

Collect once

Collect twice

Simple Snapshot: Scan

```
public int[] scan() {  
    LabeledValue[] oldCopy, newCopy;  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        if (!equals(oldCopy, newCopy)) {  
            oldCopy = newCopy;  
            continue collect;  
        }  
        return getValues(newCopy);  
    }  
}
```

Collect once

Collect twice

On mismatch, try again

Simple Snapshot: Scan

```
public int[] scan() {  
    LabeledValue[] oldCopy, newCopy;  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        if (!equals(oldCopy, newCopy)) {  
            oldCopy = newCopy;  
            continue collect;  
        }  
        return getValues(newCopy);  
    }  
}
```


Collect once

Collect twice

On match, return values


Simple Snapshot

- Linearizable
- Update is wait-free
 - No unbounded loops
- But Scan can starve
 - If interrupted by concurrent update




Wait-Free Snapshot

- Add a scan before every update
- Write resulting snapshot together with update value
- If scan is continuously interrupted by updates, scan can take the update's snapshot

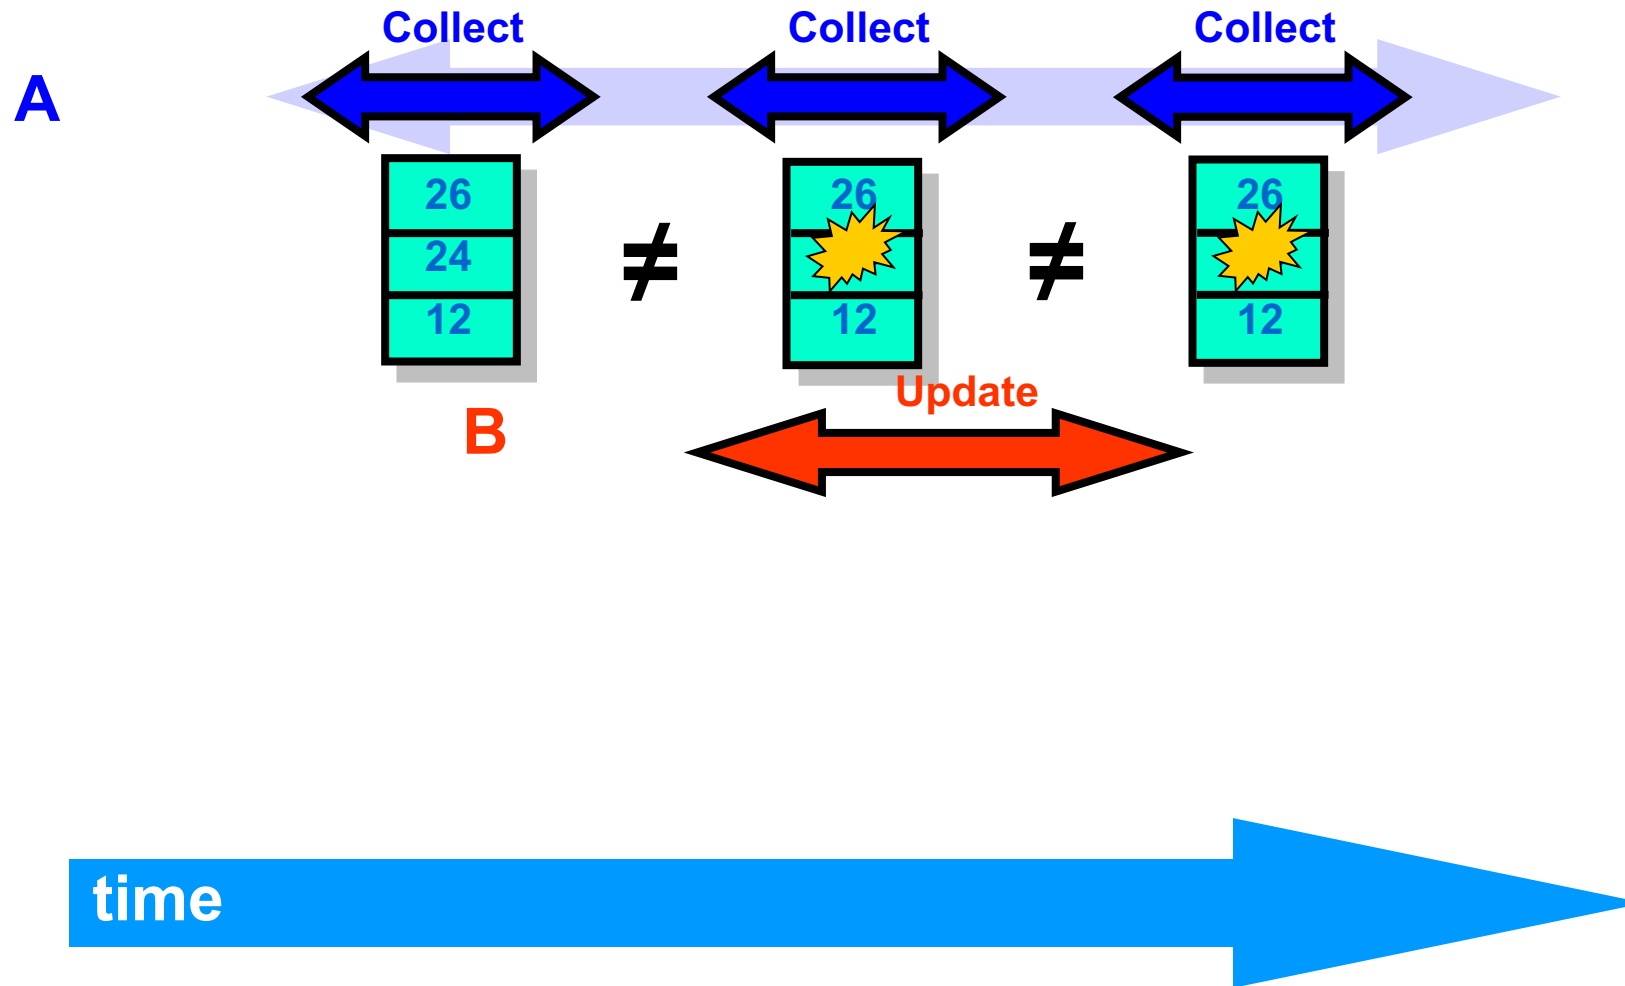


Wait-free Snapshot

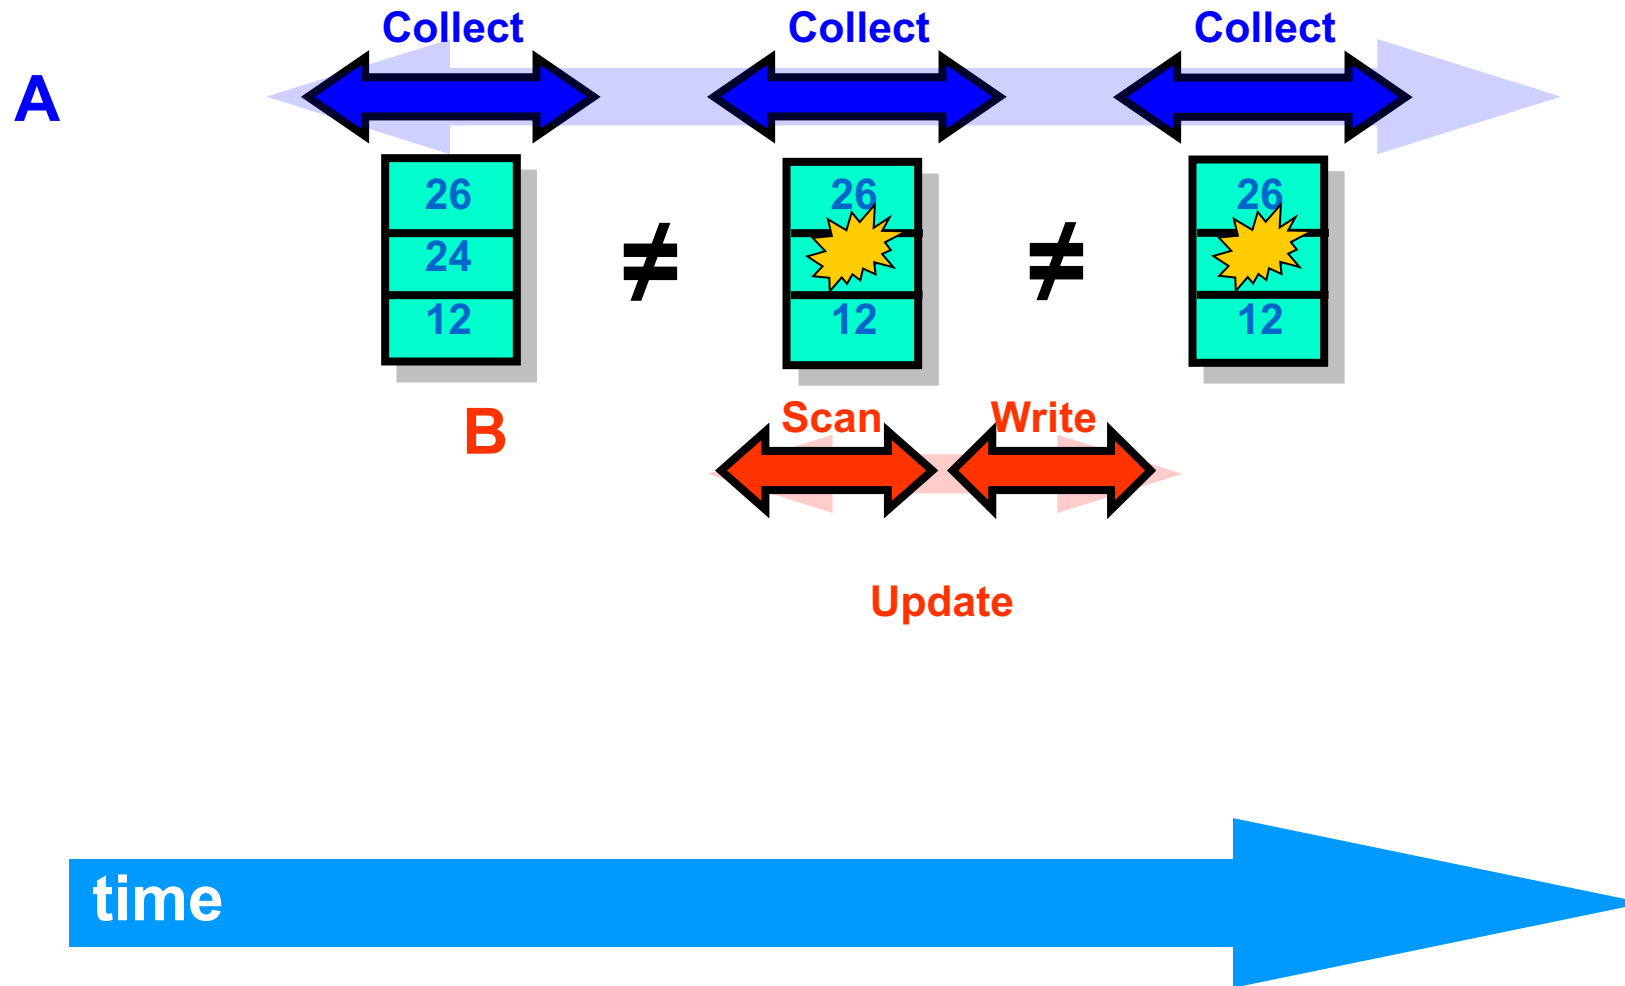
If A's scan observes that B moved **twice**, then B completed an update while A's scan was in progress



Wait-free Snapshot

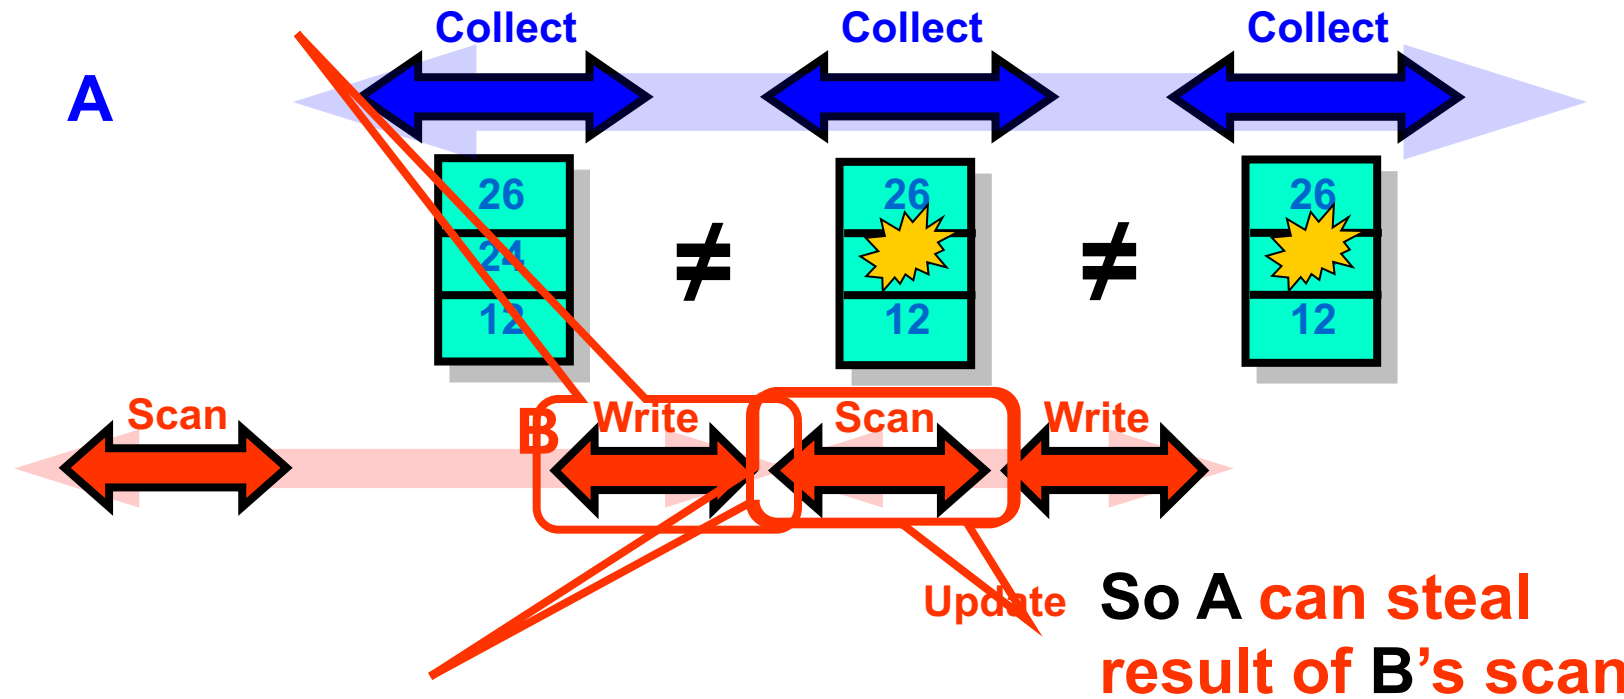


Wait-free Snapshot




Wait-free Snapshot

B's 1st update must have written during 1st collect




So scan of B's second update must be within interval of A's scan

time




Wait-free Snapshot

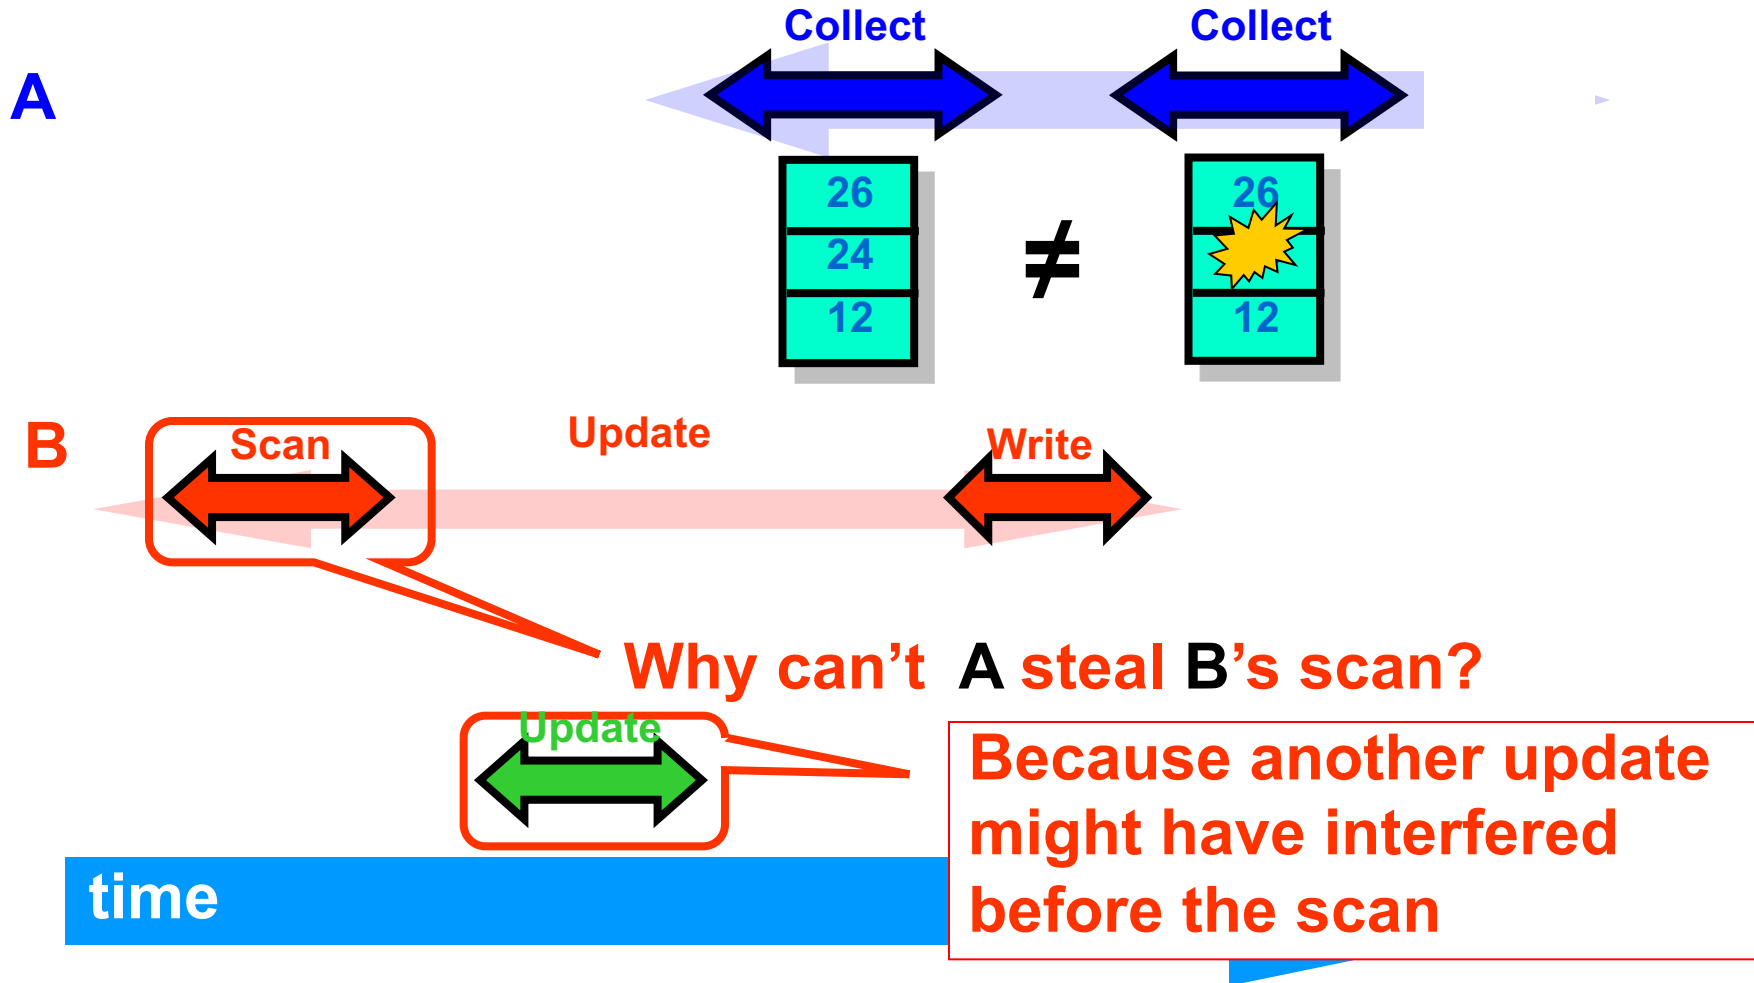


But no guarantee that scan
of B's 1st update can be used...
Why?


time



Once is not Enough

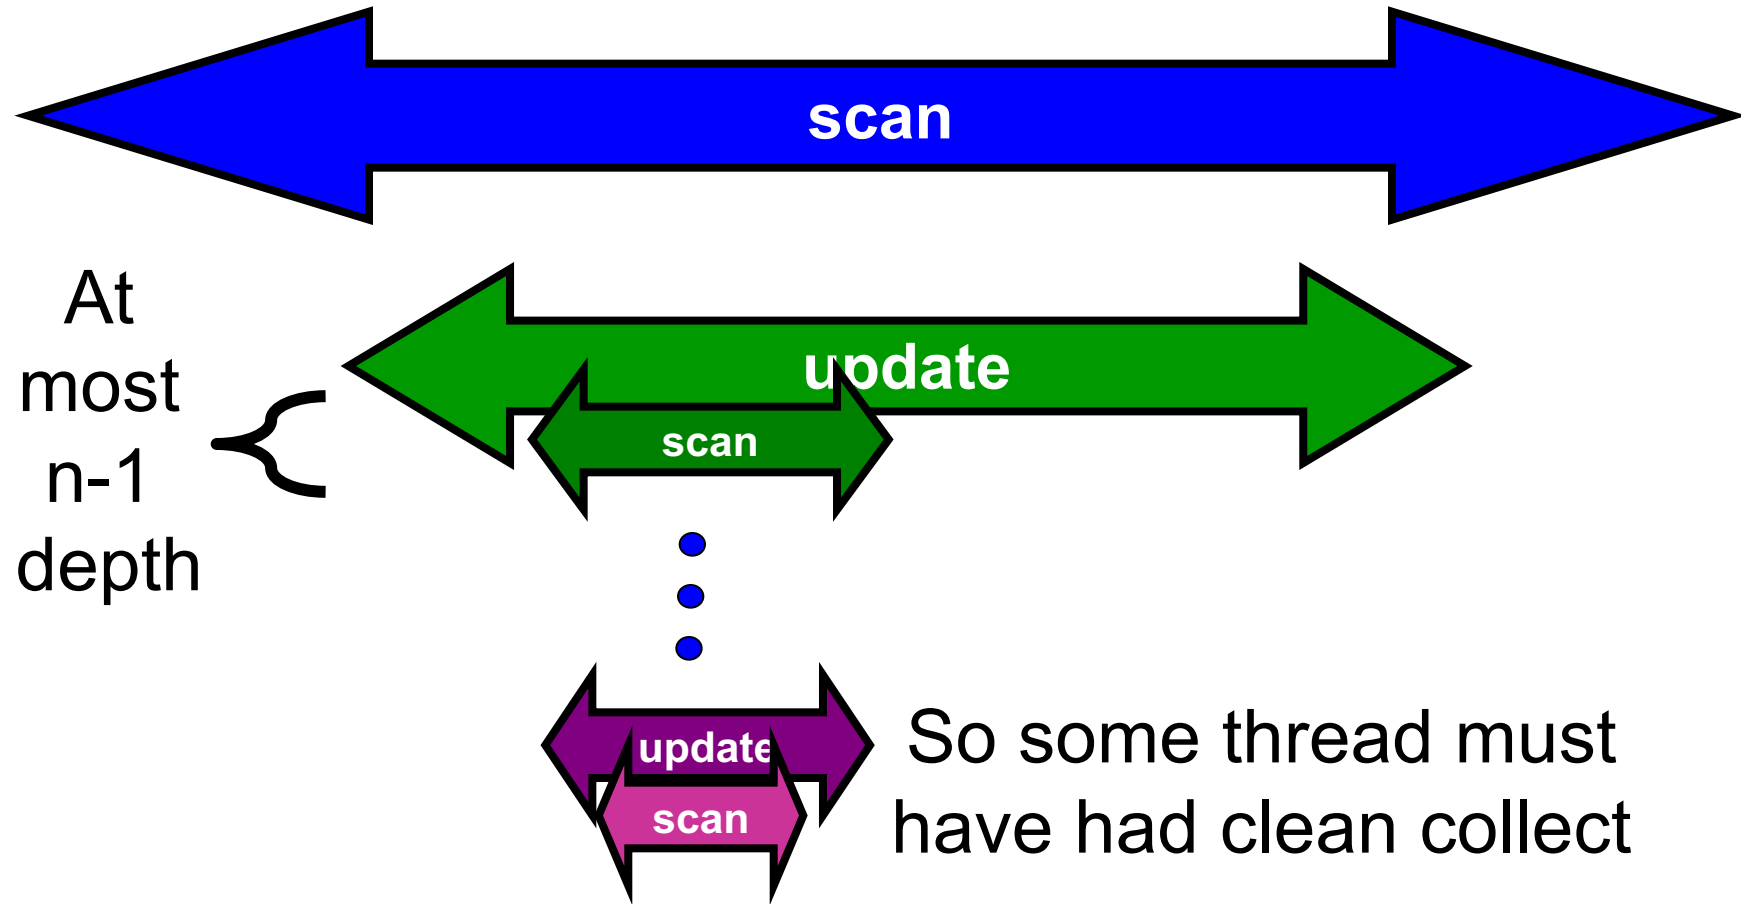


Someone Must Move Twice




If we collect n times...some thread must move twice (pigeonhole principle)

Scan is Wait-free



Wait-Free Snapshot Label

```
public class SnapValue {  
    public int    label;  
    public int    value;  
    public int[]  snap;  
}
```



Wait-Free Snapshot Label

```
public class SnapValue {  
    public int    label;  
    public int    value;  
    public int[]  snap;  
}
```

**Counter incremented
with each snapshot**

Wait-Free Snapshot Label

```
public class SnapValue {  
    public int    label;  
    public int    value;  
    public int[]  snap;  
}
```


Actual value

Wait-Free Snapshot Label

```
public class SnapValue {  
    public int    label;  
    public int    value;  
    public int[]  snap;  
}
```


most recent snapshot

Wait-Free Snapshot Label



Wait-free Update

```
public void update(int value) {  
    int i = Thread.myIndex();  
    int[] snap = this.scan();  
    SnapValue oldValue = r[i].read();  
    SnapValue newValue =  
        new SnapValue(oldValue.label+1,  
                       value, snap);  
    r[i].write(newValue);  
}
```



Wait-free Scan

```
public void update(int value) {  
    int i = Thread.myIndex();  
    int[] snap = this.scan(); Take scan  
    SnapValue oldValue = r[i].read();  
    SnapValue newValue =  
        new SnapValue(oldValue.label+1,  
                        value, snap);  
    r[i].write(newValue);  
}
```

Wait-free Scan


```
public void update(int value) {  
    int i = Thread.myIndex();  
    int[] snap = this.scan();  
    SnapValue oldValue = r[i].read();  
    SnapValue newValue =  
        new SnapValue(oldValue.label+1,  
            value, snap);  
    r[i].write(newValue);  
}
```

Take scan

Label value with scan

Wait-free Scan

```
public int[] scan() {  
    SnapValue[] oldCopy, newCopy;  
    boolean[] moved = new boolean[n];  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        for (int j = 0; j < n; j++) {  
            if (oldCopy[j].label != newCopy[j].label) {  
                ...  
            }  
        }  
        return getValues(newCopy);  
    }  
}
```



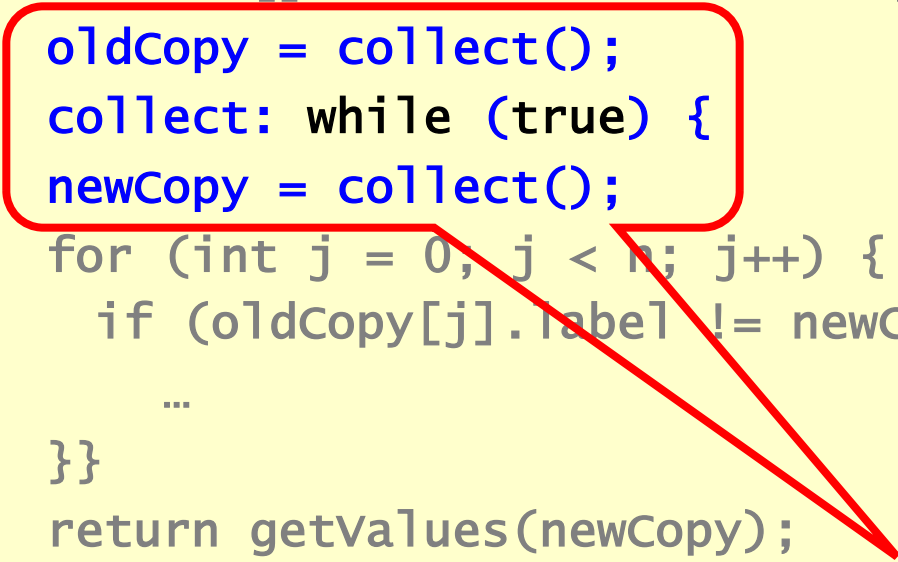
Wait-free Scan

```
public int[] scan() {  
    SnapValue[] oldCopy, newCopy;  
    boolean[] moved = new boolean[n];  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        for (int j = 0; j < n; j++) {  
            if (oldCopy[j].label != newCopy[j].label) {  
                ...  
            }  
        }  
        return getValues(newCopy);  
    }  
}
```

Keep track of who moved

Wait-free Scan

```
public int[] scan() {  
    SnapValue[] oldCopy, newCopy;  
    boolean[] moved = new boolean[n];  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        for (int j = 0; j < n; j++) {  
            if (oldCopy[j].label != newCopy[j].label) {  
                ...  
            }  
        }  
        return getValues(newCopy);  
    }  
}
```




Repeated double collect

Wait-free Scan

```
public int[] scan() {  
    SnapValue[] oldCopy, newCopy;  
    boolean[] moved = new boolean[n];  
    oldCopy = collect();  
    collect: while (true) {  
        newCopy = collect();  
        for (int i = 0; i < n; i++) {  
            if (oldCopy[i].label != newCopy[i].label) {  
                ...  
            }  
        }  
        return getValues(newCopy);  
    }  
}
```

If mismatch detected...



Mismatch Detected

```
if (oldCopy[j].label != newCopy[j].label) {  
    if (moved[j]) {          // second move  
        return newCopy[j].snap;  
    } else {  
        moved[j] = true;  
        oldCopy = newCopy;  
        continue collect;  
    }  
}  
return getValues(newCopy);  
}
```

Mismatch Detected

```
if (oldCopy[j].label != newCopy[j].label) {  
    if (moved[j]) {  
        return newCopy[j].snap;  
    } else {  
        moved[j] = true;  
        oldCopy = newCopy;  
        continue collect;  
    }  
}  
return getValues(newCopy);  
}
```

**If thread moved twice,
just steal its second
snapshot**


Mismatch Detected

```
if (oldCopy[j].label != newCopy[j].label) {  
    if (moved[j]) {        // second move  
        return newCopy[j].snap;  
    } else {  
        moved[j] = true;  
        oldCopy = newCopy;  
        continue collect;  
    }  
}  
return getValues(newCopy);  
}
```

**Remember that
thread moved**

Observations

- Uses unbounded counters
 - can be replaced with 2 bits
- Assumes SWMR registers
 - for labels
 - can be extended to MRMW



Summary

- We saw we could implement MRMW multi valued snapshot objects
- From SRSW binary safe registers (simple flipflops)
- But what is the next step to attempt with read-write registers?

