




Universality of Consensus



Hyungsoo Jung




Turing Computability



- A mathematical model of computation
- Computable = Computable on a T-Machine

Shared-Memory Computability



- Model of asynchronous concurrent computation
- Computable = Wait-free/Lock-free computable on a multiprocessor


Consensus Hierarchy

1 Read/Write Registers, Snapshots...

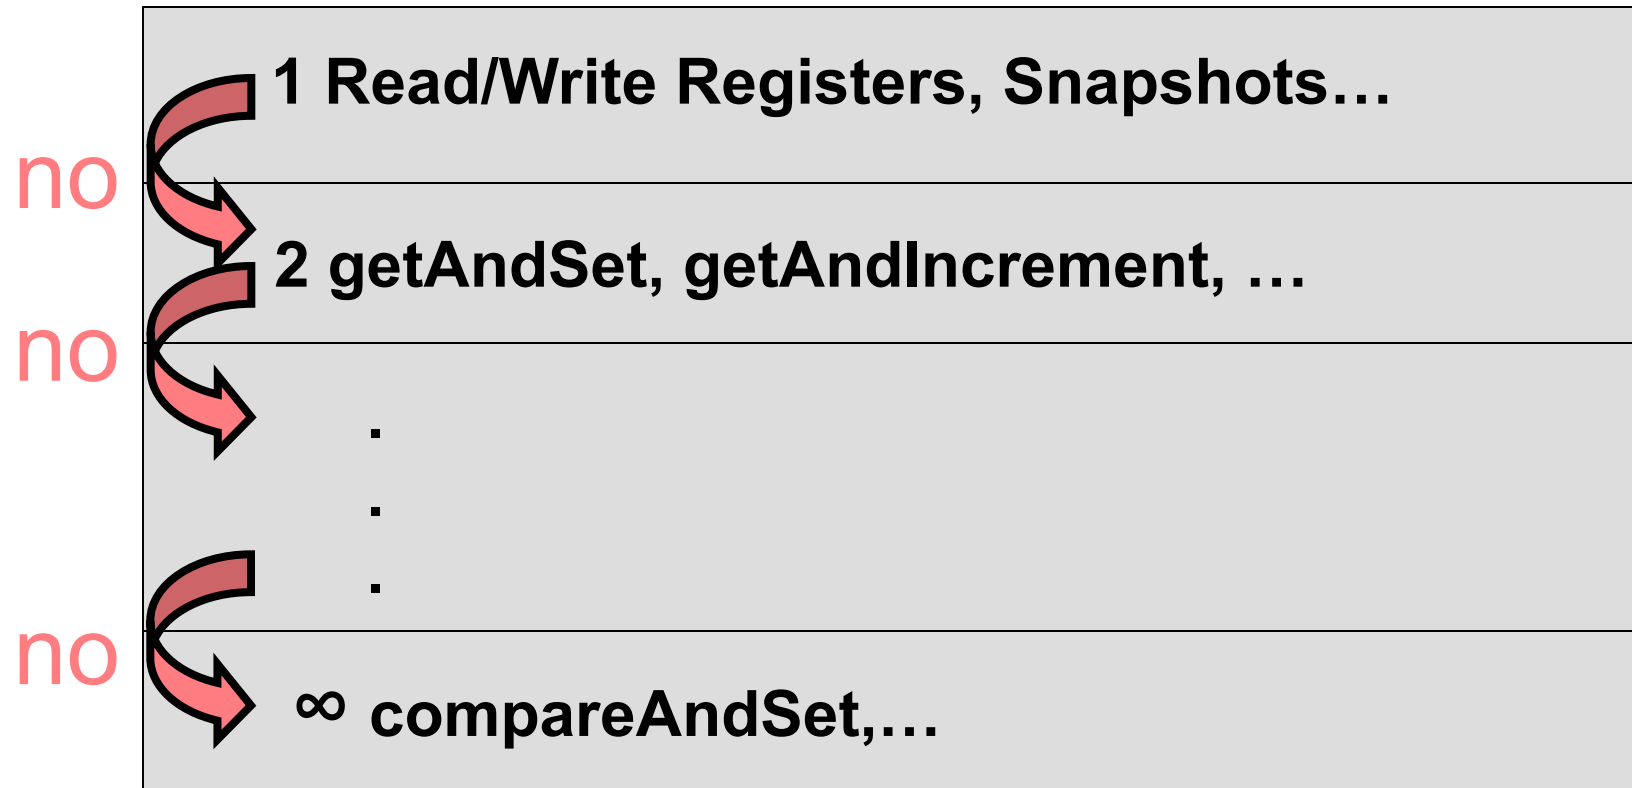
2 getAndSet, getAndIncrement, ...

▪
▪
▪

∞ compareAndSet,...



Who Implements Whom?



Hypothesis



Theorem: Universality

- Consensus is **universal**
- From n -thread consensus build a
 - Wait-free
 - Linearizable
 - n -threaded implementation
 - Of **any** sequentially specified object



Proof Outline

- A universal construction
 - From ***n***-consensus objects
 - And atomic registers
- Any wait-free linearizable object
 - Not a practical construction
 - But we know where to start looking ...



Like a Turing Machine

- This construction
 - Illustrates what needs to be done
 - Optimization fodder
- Correctness, not efficiency
 - **Why** does it work? (Asks the scientist)
 - **How** does it work? (Asks the engineer)
 - Would you like fries with that? (Asks the liberal arts major)



A Generic Sequential Object

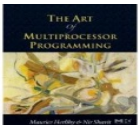
```
public interface SeqObject {  
    public abstract  
        Response apply(Invocation invoc);  
}
```



A Generic Sequential Object

```
public interface SeqObject {  
    public abstract  
    Response apply(Invocation invoc);  
}
```

Push:5, Pop:null



Invocation

```
public class Invoc {  
    public String method;  
    public Object[] args;  
}
```



Invocation


```
public class Invoc {  
    public String method;  
    public Object[] args;  
}
```

Method name

Invocation

```
public class Invoc {  
    public String method;  
    public Object[] args;  
}
```

Arguments



A Generic Sequential Object

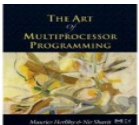
```
public interface SeqObject {  
    public abstract  
        Response apply(Invocation invoc);  
}
```



A Generic Sequential Object

```
public interface SeqObject {  
    public abstract  
    Response apply(Invocation invoc);  
}
```

OK, 4



Response

```
public class Response {  
    public Object value;  
}
```



Response

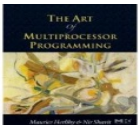
```
public class Response {  
    public object value;  
}
```

Return value

Universal Concurrent Object

```
public interface SeqObject {  
    public abstract  
        Response apply(Invocation invoc);  
}
```

A universal concurrent object is
linearizable to the generic
sequential object



Start with Lock-Free Universal Construction




- First Lock-free: infinitely often some method call finishes.
- Then Wait-Free: each method call takes a finite number of steps to finish

Naïve Idea

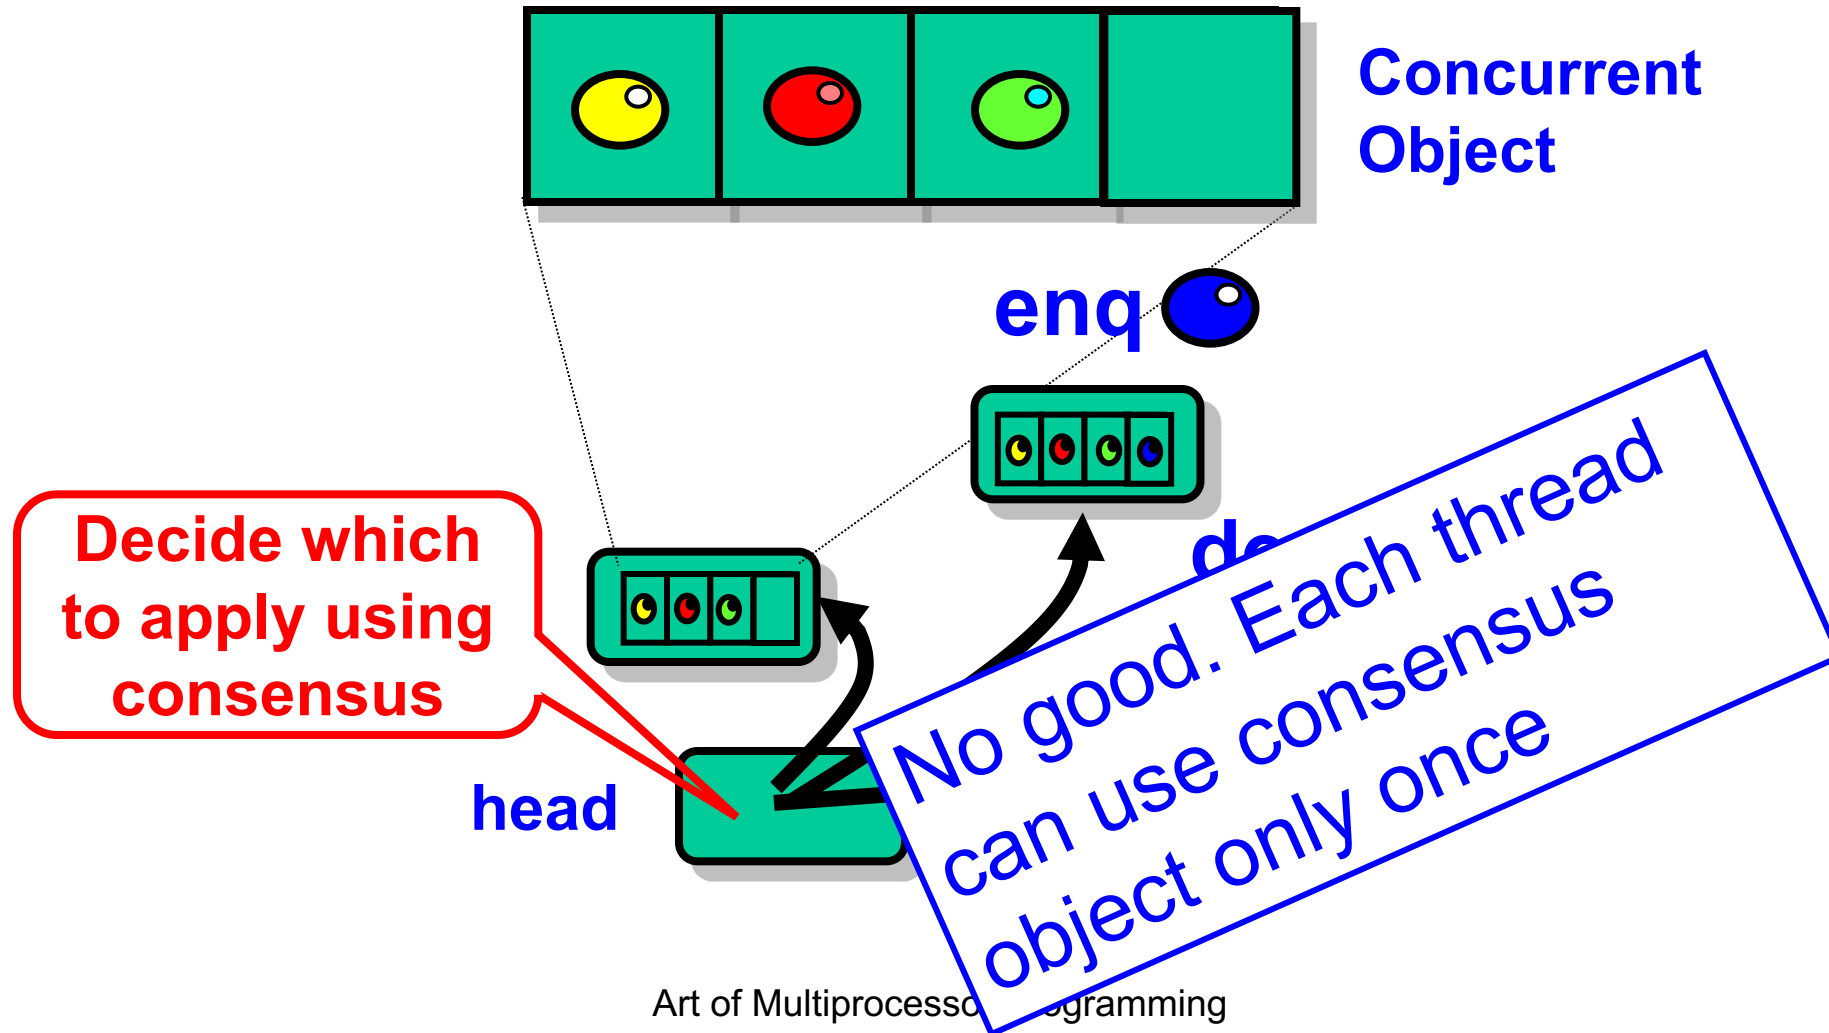
- Consensus object stores reference to cell with current state
- Each thread creates new cell
 - computes outcome,
 - tries to switch pointer to its outcome
- Sadly, no ...
 - consensus objects can be used once only



Naïve Idea



Naïve Idea




Once is not Enough?

**Queue based
consensus**


```
public void propose(int value) {  
    propose(value);  
    Ball ball = queue.deq();  
    if (ball == Ball.RED)  
        return proposed[i];  
    else  
        return proposed[1-i];  
}
```

**Solved one-shot 2-consensus.
Not clear how to reuse or reread ...**



Improved Idea: Linked-List Representation

Each node contains a fresh consensus object used to decide on next operation



Universal Construction

- Object represented as
 - Initial Object State
 - A Log: a linked list of the method calls



Scaling a file system to many cores using an operation log


Srivatsa S. Bhat,[†] Rasha Eqbal,[‡] Austin T. Clements,[§]
M. Frans Kaashoek, Nickolai Zeldovich
MIT CSAIL

SOSP'17, October 28–31, 2017, Shanghai, China.

© 2017 Copyright is held by the owner/author(s).

ACM ISBN 978-1-4503-5085-3/17/10.

<https://doi.org/10.1145/3132747.3132779>



Basic Idea

- Use one-time consensus object to decide next pointer



Basic Idea

- Use one-time consensus object to decide next pointer
- All threads update actual next pointer based on decision
 - OK because they all write the same value



Basic Idea

- Threads use one-time consensus object to decide which node goes next
- Threads update actual **next** field to reflect consensus outcome
 - OK because they all write the same value
- Challenges
 - Lock-free means we need to worry what happens if a thread stops in the middle



Basic Data Structures

```
public class Node implements
java.lang.Comparable {
    public Invoc invoc;
    public Consensus<Node> decideNext;
    public Node next;
    public int seq;
    public Node(Invoc invoc) {
        invoc = invoc;
        decideNext = new Consensus<Node>()
        seq = 0;
    }
}
```



Basic Data Structures

```
public class Node implements
    java.lang.Comparable {
    public Invoc invoc;
    public Consensus<Node> decideNext;
    public Node next;
    public int seq;

    public Consensus<Node> decideNext() {
        decideNext = new Consensus<Node>();
        seq = 0;
    }
}
```


Standard interface for class whose objects are totally ordered



Basic Data Structures

```
public class Node implements
java.lang.Comparable {
    public Invoc invoc;
    public Consensus<Node> decideNext;
    public Node next;
    public int seq;
    public
        INVOC = INVOC;
        decideNext = new Consensus<Node>()
        seq = 0;
}
```

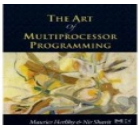
the invocation



Basic Data Structures

```
public class Node implements  
    java.lang.Comparable {  
    public Invoc invoc;  
    public Consensus<Node> decideNext;  
    public Node next;  
    public int seq;  
    public Node(Invoc invoc) {  
        invoc = invoc;  
    }  
}
```

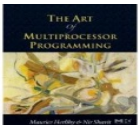
**Decide on next node
(next method applied to object)**



Basic Data Structures

```
public class Node implements  
java.lang.Comparable {  
    public Invoc invoc;  
    public Consensus<Node> decideNext;  
    public Node next;  
    public int seq;  
    public Node(Invoc invoc) {  
        invoc = invoc;  
    }  
}
```

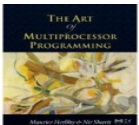
**Traversable pointer to next node
(needed because you cannot
repeatedly read a consensus object)**



Basic Data Structures

```
public class Node implements  
java.lang.Comparable {  
    public Invoc invoc;  
    public Consensus<Node> decideNext;  
    public Node next;  
    public int seq;  
    public Node(Invoc invoc) {  
        invoc = invoc;  
        decideNext = new Consensus<Node>()  
        seq = 0;  
    }  
}
```

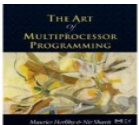
Seq number

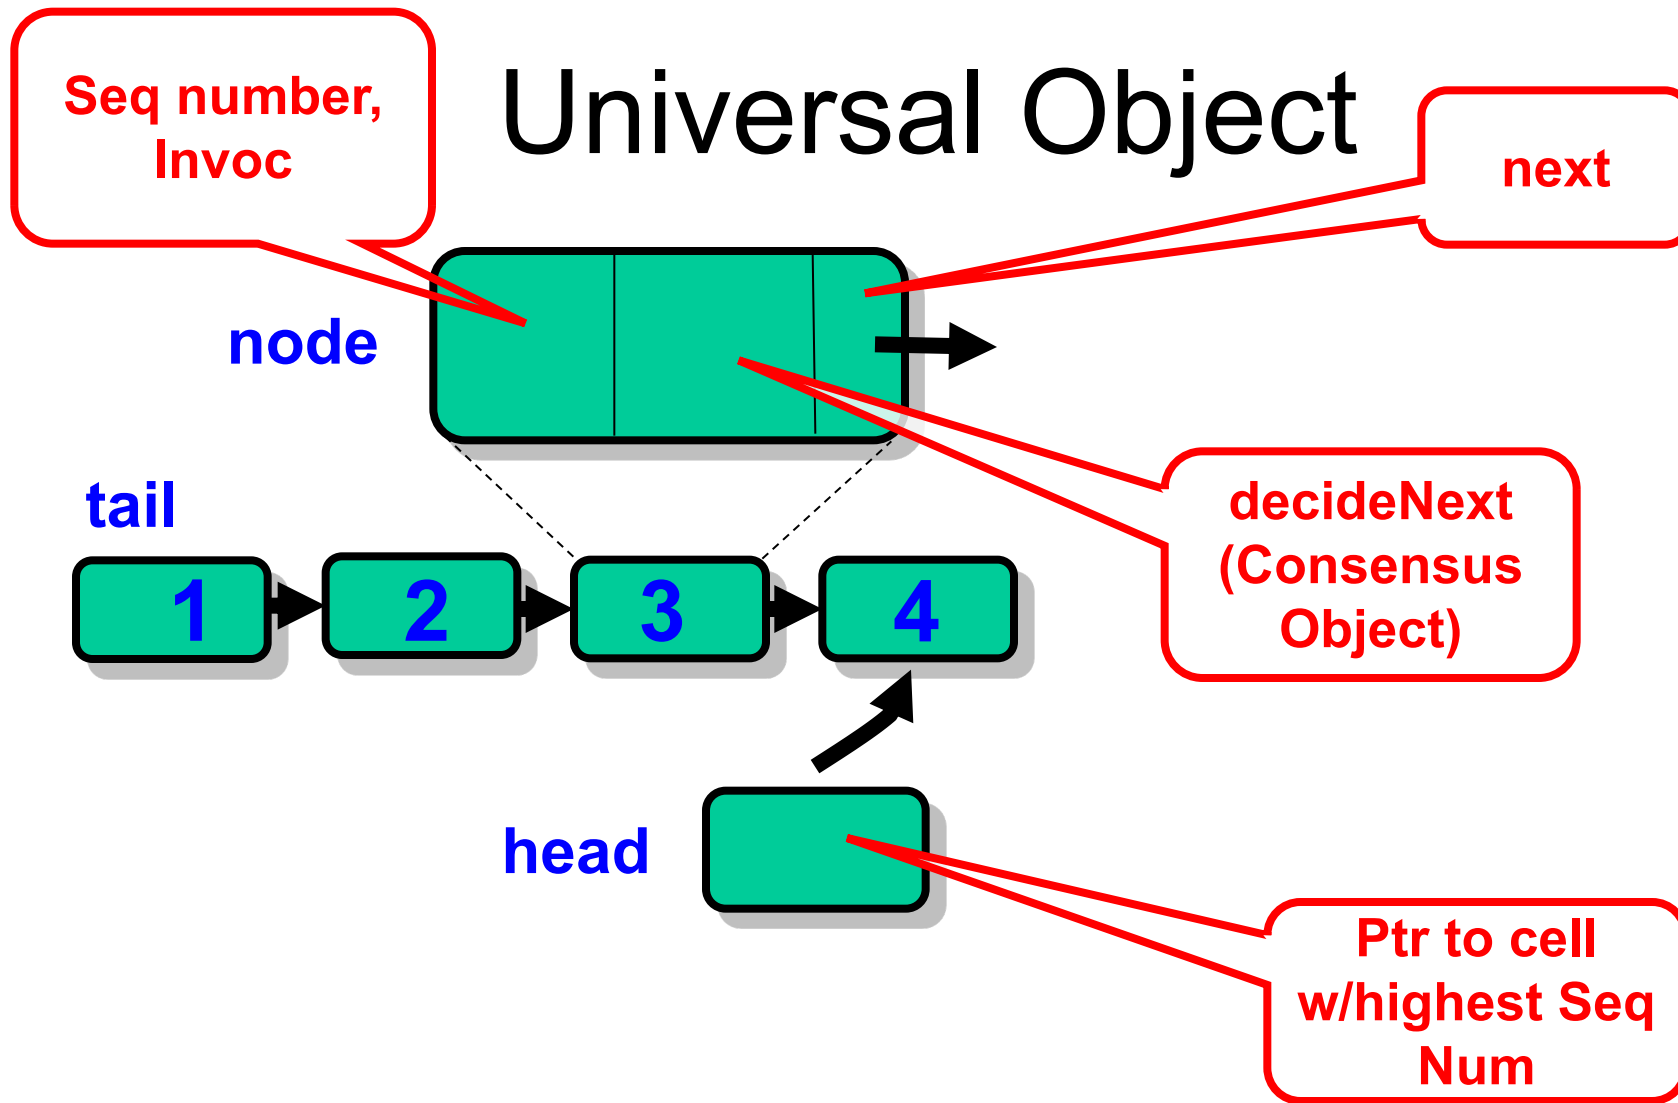


Basic Data Structures

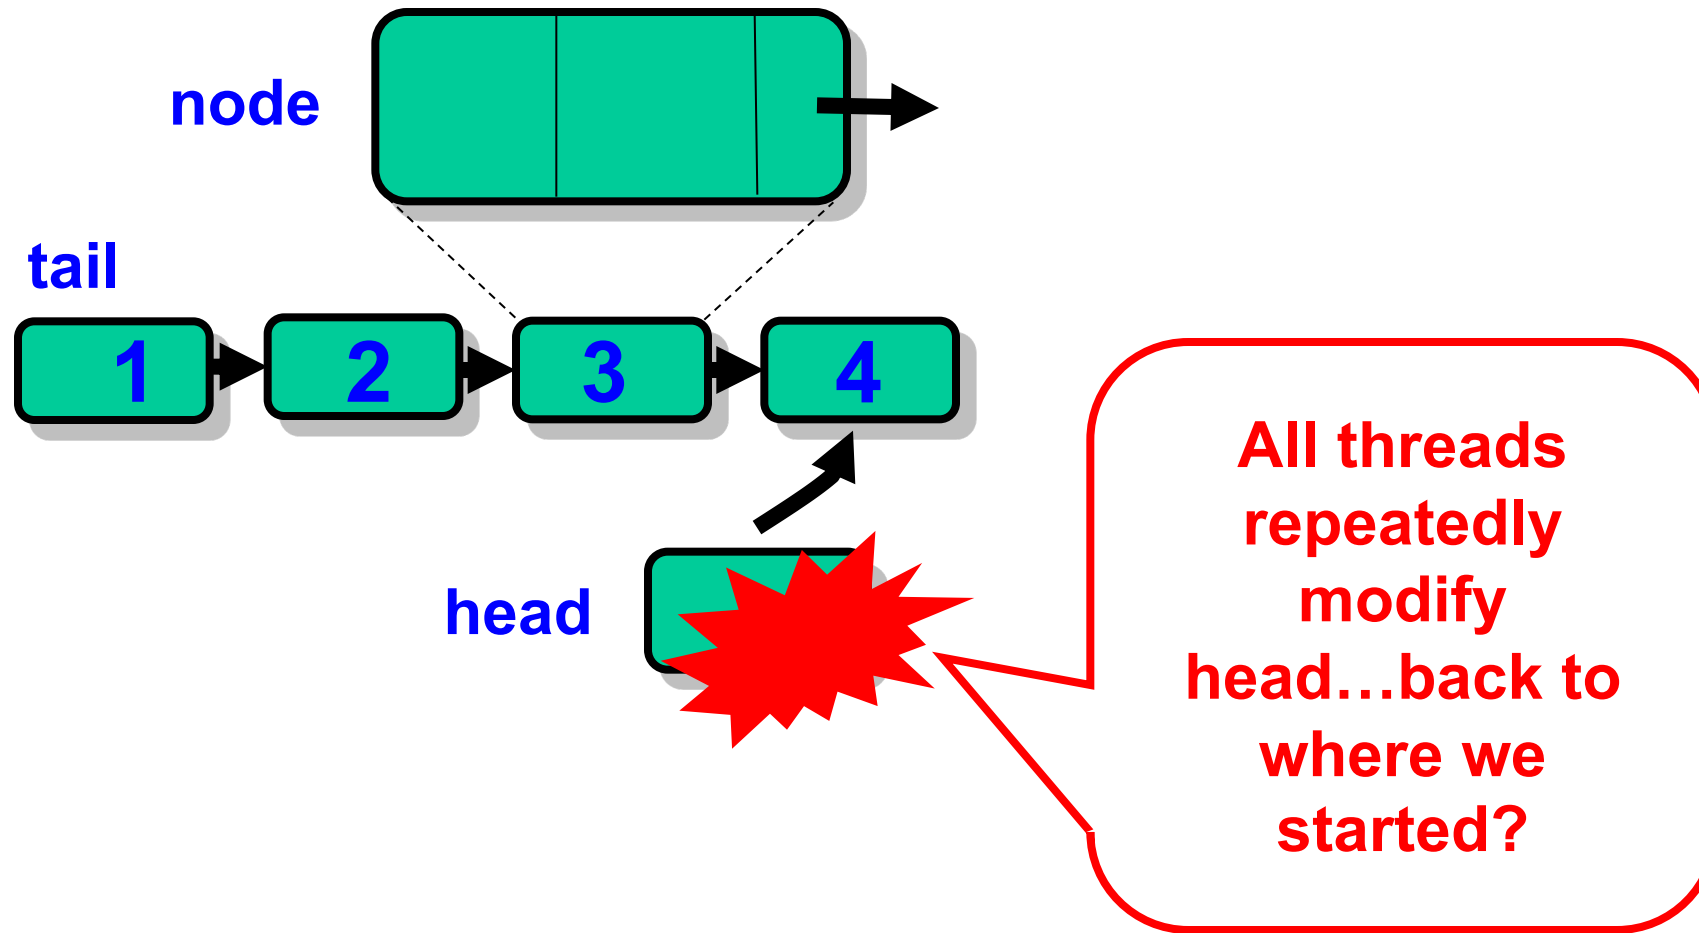
Create new node for this method invocation

```
public Invoc invoc;  
public Consensus<Node> decideNext;  
public Node next;  
public int seq;  
public Node(Invoc invoc) {  
    invoc = invoc;  
    decideNext = new Consensus<Node>()  
    seq = 0;  
}
```

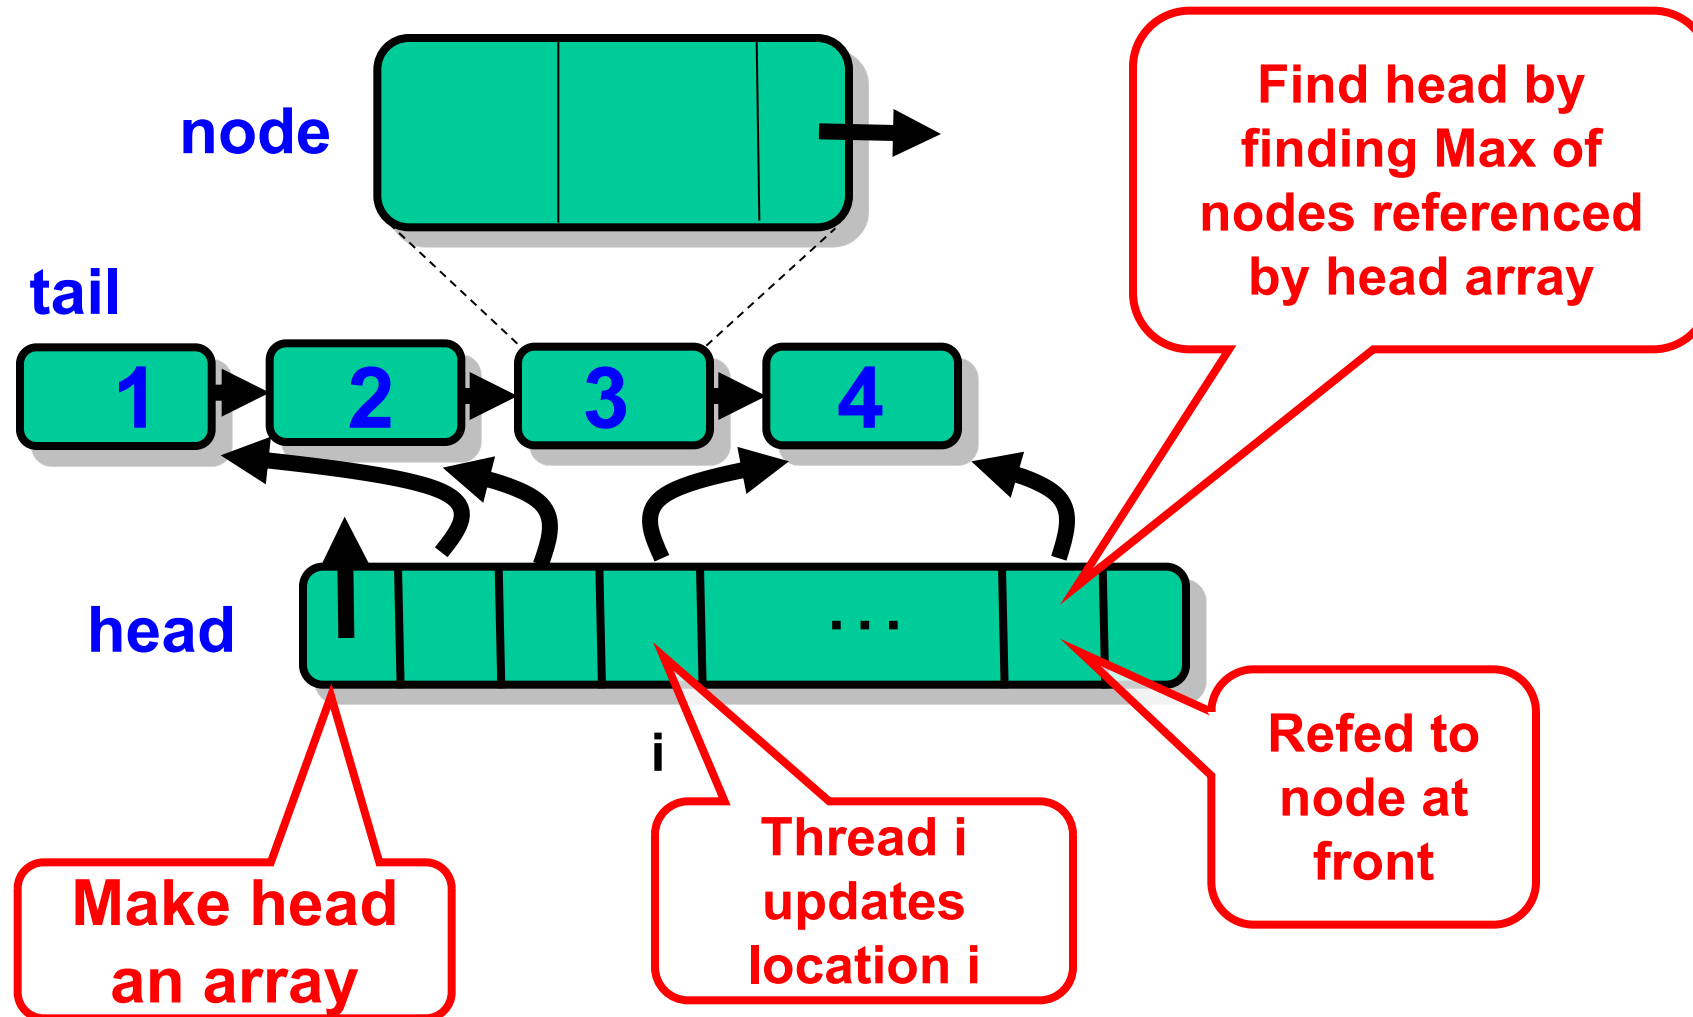




Universal Object



The Solution



Universal Object

```
public class Universal {  
    private Node[] head;  
    private Node tail = new Node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail  
    }  
}
```



Universal Object

```
public class Universal {  
    private Node[] head;  
    private Node tail = new Node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail  
    }  
}
```

Head Pointers Array

Universal Object

```
public class Universal {  
    private Node[] head;  
    private Node tail = new Node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail  
    }  
}
```

**Tail is a sentinel node with
sequence number 1**

Universal Object

```
public class Universal {  
    private Node[] head;  
    private Node tail = new Node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail  
    }
```

Initially head points to tail

Find Max Head Value

```
public static Node max(Node[] array) {  
    Node max = array[0];  
    for (int i = 1; i < array.length; i++)  
        if (max.seq < array[i].seq)  
            max = array[i];  
    return max;  
}
```



Find Max Head Value

```
public static Node max(Node[] array) {  
    Node max = array[0];  
    for (int i = 0; i < array.length; i++)  
        if (max.seq < array[i].seq)  
            max = array[i];  
    return max;  
}
```

**Traverse
the array**

Find Max Head Value

```
public static Node max(Node[] array) {  
    Node max = array[0];  
    for (int i = 0; i < array.length; i++)  
        if (max.seq < array[i].seq)  
            max = array[i];  
    return max;  
}
```

**Compare the seq nums of nodes
pointed to by the array**

Find Max Head Value

```
public static Node max(Node[] array) {  
    Node max = array[0];  
    for (int i = 0; i < array.length; i++)  
        if (max.seq < array[i].seq)  
            max = array[i];  
    return max;  
}
```

Return node with maximal sequence number

Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
            before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
        head[i] = after;  
    }  
    ...  
}
```




Universal Application Part I

```
public Response apply(Invoc invoc) {
```

```
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
            before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
    }
```

**Apply has invocation as input and
returns the appropriate response**

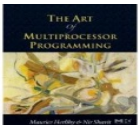
...



Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
            before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
        he  
    }  
    ...  
}
```

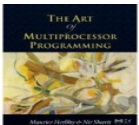
my ID



Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
            before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
        head[i] = after;  
    }  
    ...  
}
```

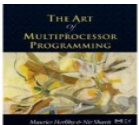
My method call



Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
            before.decideNextNode(prefer);  
        before.  
after.  
head[i] = after;  
    }  
    ...  
}
```


**As long as I have not been
threaded into list**



Universal Application Part I

```
public Response apply(Invoc invoc) {
    int i = ThreadID.get();
    Node prefer = new node(invoc);
    while (prefer.seq == 0) {
        Node before = Node.max(head);
        Node after =
            before.decideNext.decide(prefer);
        before.i
        after.s
        head[i] = after,
    }
    ...
}
```

**Head of list to which we
will try to append**



Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    while (prefer.seq == 0) {  
        Node before = Node.max(head);  
        Node after =  
        before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
        head[i  
    }  
}
```

Propose next node

Universal Application

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new Node(invoc);  
    Set next field to consensus winner  
    Node before = Node.max(head);  
    Node after =  
        before.decideNext.decide(prefer);  
    before.next = after;  
    after.seq = before.seq + 1;  
    head[i] = after;  
}  
}
```


Universal Application Part I

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    Node prefer = new node(invoc);  
    ...
```


**Set seq number
(indicating node was appended)**

```
    before.decideNext decide(prefer);  
    before.next = after;
```

```
    after.seq = before.seq + 1;
```

```
    head[i] = after;  
}
```


...



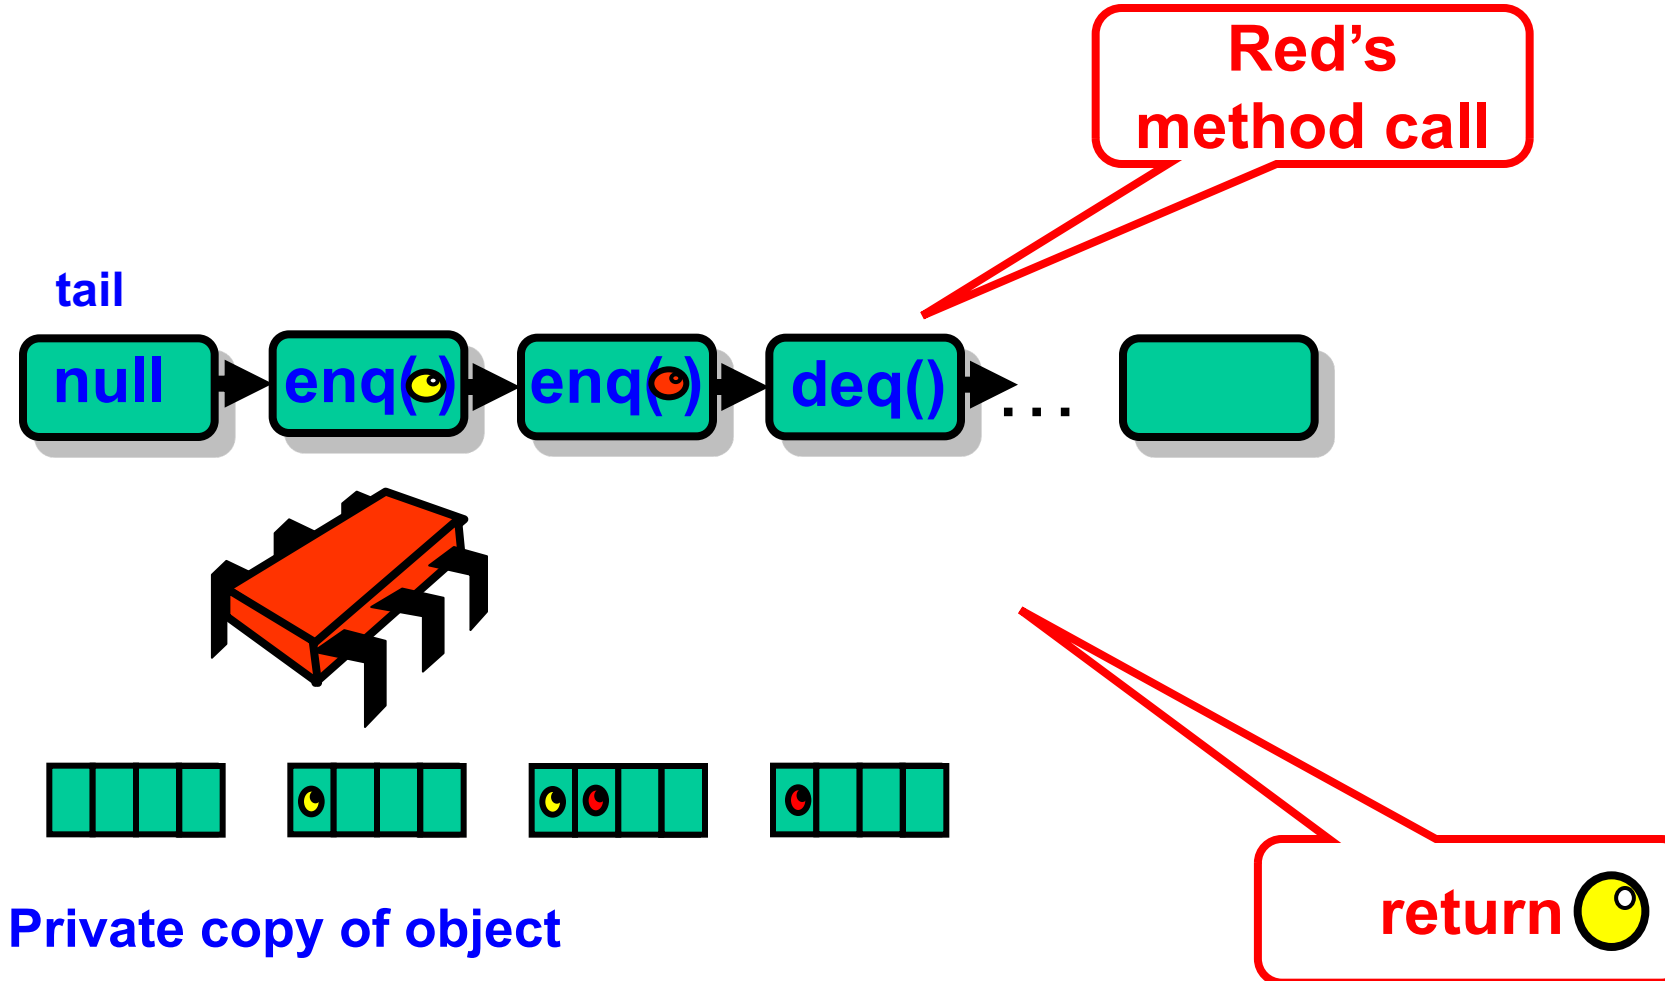
Universal Application Part I

```
public Response apply(Transaction tx) {  
    int i = 0;  
    Node pre = head[0];  
    while (pre.seq == 0) {  
        Node before = Node.max(head);  
        Node after = Node.max(head);  
        before.decideNext.decide(prefer);  
        before.next = after;  
        after.seq = before.seq + 1;  
        head[i] = after;  
    }  
    ...  
}
```

add to head array so new head will be found



Part 2 – Compute Response



Universal Application Part 2

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```



Universal Application Part II


```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}
```

**Compute result by sequentially applying
method calls in list to a private copy**

Universal Application Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

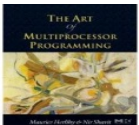
Start with copy of sequential object



Universal Application Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

new method call appended after tail



Universal Application Part II

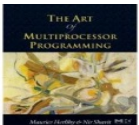
```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

While my method call not linked ...

Universal Application Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

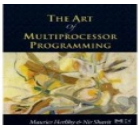
Apply current node's method



Universal Application Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != prefer){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

Return result after my method call applied



Correctness

- List defines linearized sequential history
- Thread returns its response based on list order



Lock-freedom

- Lock-free because
 - A thread moves forward in list
 - Can repeatedly fail to win consensus on “real” head only if another succeeds
 - Consensus winner adds node and completes within a finite number of steps



Wait-free Construction

- Lock-free construction + **announce array**
- Stores (pointer to) node in **announce**
 - If a thread doesn't append its node
 - Another thread will see it in array and *help* append it




Helping

- “Announcing” my intention
 - Guarantees progress
 - Even if the scheduler hates me
 - My method call will complete
- Makes protocol wait-free
- Otherwise starvation possible



Wait-free Construction



The Announce Array

```
public class Universal {  
    private Node[] announce;  
    private Node[] head;  
    private Node tail = new node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail; announce[j] = tail  
    };  
};
```


The Announce Array

```
public class Universal {  
    private Node[] announce;  
    private Node[] head;  
    private Node tail = new node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail; announce[j] = tail  
    };  
};
```

New field: announce array

The Announce Array

```
public class Universal {  
    private Node[] announce;  
    private Node[] head;  
    private Node tail = new node();  
    tail.seq = 1;  
    for (int j=0; j < n; j++){  
        head[j] = tail; announce[j] = tail  
    };  
};
```

All entries initially point to tail

A Cry For Help

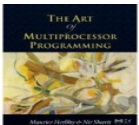
```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    announce[i] = new Node(invoc);  
    head[i] = Node.max(head);  
    while (announce[i].seq == 0) {  
        ...  
        // while node not appended to list  
        ...  
    }  
}
```



A Cry For Help

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    announce[i] = new Node(invoc);  
    head[i] = Node.max(head);  
    while (announce[i].seq == 0) {  
        ...  
        // while node not appended to list  
        ...  
    }  
}
```

**Announce new method call,
asking help from others**



A Cry For Help


```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    announce[i] = new Node(invoc);  
    head[i] = Node.max(head);  
    while (announce[i].seq == 0) {  
        ...  
        // while node not appended to list  
        ...  
    }  
}
```

Look for end of list

A Cry For Help

```
public Response apply(Invoc invoc) {  
    int i = ThreadID.get();  
    announce[i] = new Node(invoc);  
    head[i] = Node.max(head);  
    while (announce[i].seq == 0) {  
        ...  
        // while node not appended to list  
        ...  
    }  
}
```

**Main loop, while node not appended
(either by me or helper)**



Main Loop

- Non-zero sequence # means success



Main Loop

- Non-zero sequence # means success
- Thread keeps helping append nodes



Main Loop

- Non-zero sequence # means success
- Thread keeps helping append nodes
- Until its own node is appended



Main Loop

```
while (announce[i].seq == 0) {  
    Node before = head[i];  
    Node help = announce[(before.seq + 1) % n];  
    if (help.seq == 0)  
        prefer = help;  
    else  
        prefer = announce[i];  
    ...  
}
```

Main Loop

```
while (announce[i].seq == 0) {
```

```
    Node before = head[i];  
    Node help = announce[(before.seq + 1) % n];  
    if (help.seq == 0)  
        prefer = help;  
    else
```

**Keep trying until my cell gets a
sequence number**

Main Loop

```
while (announce[i].seq == 0) {  
    Node before = head[i];  
    Node help = announce[(before.seq + 1) % n];  
    if (help.seq == 0)  
        prefer = help;  
    else  
        prefer = announce[i].
```

Possible end of list

Main Loop

```
while (announce[i].seq == 0) {  
    Node before = head[i];  
    Node help = announce[(before.seq + 1) % n];  
    if (help.seq == 0)  
        prefer = help;  
    else  
        prefer = announce[i];  
}
```

Whom do I help?

Altruism

- Choose a thread to “help”



Altruism

- Choose a thread to “help”
- If that thread needs help
 - Try to append its node
 - Otherwise append your own



Altruism

- Choose a thread to “help”
- If that thread needs help
 - Try to append its node
 - Otherwise append your own
- Worst case
 - Everyone tries to help same pitiful loser
 - Someone succeeds



Help!

- When last node in list has sequence number k



Help!

- When last node in list has sequence number k
- All threads check ...
 - Whether thread $k+1 \bmod n$ wants help
 - If so, try to append her node first



Help!

- First time after thread $k+1$ announces
 - No guarantees



Help!

- First time after thread $k+1$ announces
 - No guarantees
- After n more nodes appended
 - Everyone sees that thread $k+1$ wants help
 - Everyone tries to append that node
 - Someone succeeds

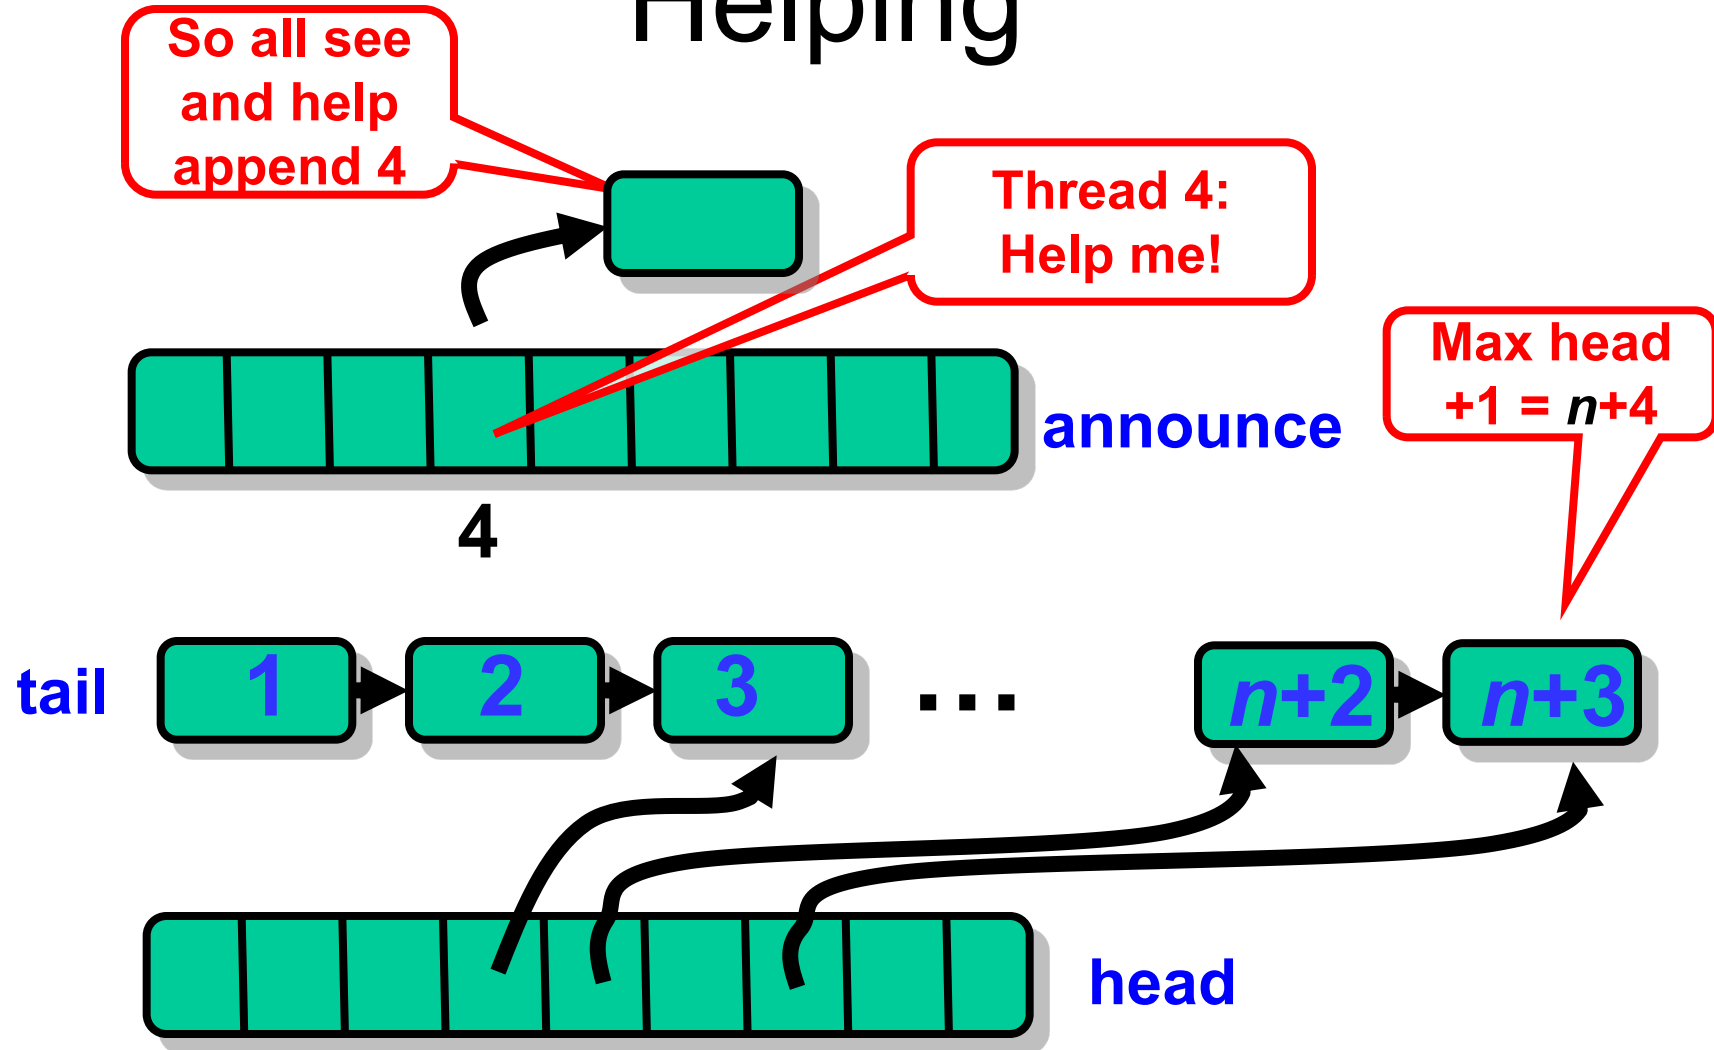


Sliding Window Lemma


- After thread *A* announces its node
- No more than *n* other calls
 - Can start and finish
 - Without appending *A*'s node



Helping



The Sliding Help Window



Sliding Help Window

```
while (announce[i].seq == 0) {  
    Node before = head[i];  
    Node help = announce[(before.seq + 1) % n];  
    if (help.seq == 0)  
        prefer = help;  
    else  
        prefer = announce[i];  
}
```

**In each main loop iteration pick
another thread to help**

Sliding Help Window

**Help if help required, but
otherwise it's all about me!**

```
while (announced) {
    Node before = head[i];
    Node help = announce[(before.seq + 1) % n];
    if (help.seq == 0)
        prefer = help;
    else
        prefer = announce[i];
    ...
}
```

Rest is Same as Lock-free

```
while (announce[i].seq == 0) {  
    ...  
    Node after =  
        before.decideNext.decide(prefer);  
    before.next = after;  
    after.seq = before.seq + 1;  
    head[i] = after;  
}
```

Rest is Same as Lock-free

```
while (announce[i].seq == 0) {
```

```
...
```

```
Node after =  
    before.decideNext.decide(prefer);
```

```
before.next = after;  
after.seq = before.seq + 1;  
head[i] = after;
```

```
}
```

Call consensus to attempt to append

Rest is Same as Lock-free

```
while (announce[i].seq == 0) {  
    . cache consensus result for later use  
    Node after =  
        before.decideNext.decide(prefer);  
    before.next = after;  
    after.seq = before.seq + 1;  
    head[i] = after;  
}
```

Rest is Same as Lock-free

```
while (announce[i].seq == 0) { ...  
    Tell world that node is appended  
    before.decideNext.decide(prefer);  
    before.next = after;  
    after.seq = before.seq + 1;  
    head[i] = after;  
}
```

Finishing the Job

- Once thread's node is linked ...



Finishing the Job

- Once thread's node is linked...
- The rest same as lock-free algorithm



Finishing the Job

- Once thread's node is linked ...
- The rest same as lock-free algorithm
- Compute result by
 - sequentially applying list's method calls
 - to a private copy of the object
 - starting from the initial state




Then Same Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
while (current != announce[i]){  
    MyObject.apply(current.invoc);  
    current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```


Universal Application Part II

```
...  
//compute my response  
SeqObject MyObject = new SeqObject();  
current = tail.next;  
Return result after applying my method  
myobject.apply(current.invoc);  
current = current.next;  
}  
return MyObject.apply(current.invoc);  
}
```

Shared-Memory Computability




Universal
Construction



Wait-free/Lock-free computable

=

Solving n -consensus



Veni, Vidi, Vici

- We saw
 - how to define concurrent objects



Veni, Vidi, Vici

- We saw
 - how to define concurrent objects
- We discussed
 - computational power of machine instructions



Veni, Vidi, Vici

- We saw
 - how to define concurrent objects
- We discussed
 - computational power of machine instructions
- Next
 - use these foundations to understand the real world

